

# Chef Fundamentals

training@chef.io

Copyright (C) 2014 Chef Software, Inc.

v2.1.6

## Introductions

v2.1.6



# Introduce Yourselves

- Name
- Current job role
- Previous job roles/background
- Experience with Chef and/or config management
- Favorite Text Editor

## Course Objectives and Style

## Course Objectives

- Upon completion of this course you will be able to
  - Automate common infrastructure tasks with Chef
  - Describe Chef's architecture
  - Describe Chef's various tools
  - Apply Chef's primitives to solve your problems

## How to learn Chef

- You bring the domain expertise about your business and problems
- Chef provides a framework for solving those problems
- Our job is to work together to teach you how to express solutions to your problems with Chef

## Chef is a Language

- Learning Chef is like learning the basics of a language
- 80% fluency will be reached very quickly
- The remaining 20% just takes practice
- The best way to **learn** Chef is to **use** Chef

## Training is really a discussion

- We will be doing things the **hard way**
- We're going to do **a lot** of typing
- You can't be:
  - Absent
  - Late
  - Left Behind
- We will troubleshoot and fix bugs on the spot
- The result is you reaching fluency fast

## Training is really a discussion

- I'll post objectives at the beginning of a section
- Ask questions when they come to you
- Ask for help when you need it
- You'll get the slides **after** class

## Fundamentals

- This course covers the fundamentals of Chef
- We likely will not have time to discuss the latest trends in the Chef ecosystem
- Any discussion items that are too far off topic will be captured
  - We'll return to these items as time permits

# Stages of Learning - Shuhari

- Shuhari - first learn, then detach, and finally transcend
  - shu - "obey" - traditional wisdom
  - ha - "detach" - break with tradition
  - ri - "separate" - transcend

守破離

## Agenda

## Topics

- Overview of Chef
- Workstation Setup
- Organization Setup
- Test Node Setup
- Writing your first cookbook
- Dissecting your first Chef run
- Introducing the Node object
- Attributes, Templates, and Cookbook Dependencies

## Topics

- Template Variables, Notifications, and Controlling Idempotency
- Data bags, search
- Roles
- Environments
- Using community cookbooks
- Further Resources

# Breaks!

- We'll take a break between each section, or every hour, whichever comes first
- We'll obviously break for lunch :)

## Overview of Chef

# Lesson Objectives

- After completing the lesson, you will be able to
  - Describe how Chef thinks about Infrastructure Automation
  - Define the following terms:
    - Node
    - Resource
    - Recipe
    - Cookbook
    - Run List
    - Roles
    - Search

## Complexity



## Items of Manipulation (Resources)

- Networking
- Files
- Directories
- Symlinks
- Mounts
- Registry Keys
- Powershell Scripts
- Users
- Groups
- Packages
- Services
- Filesystems

A tale of growth...

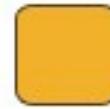


Application

## Add a database



Application



Application Database

## Make database redundant



Application

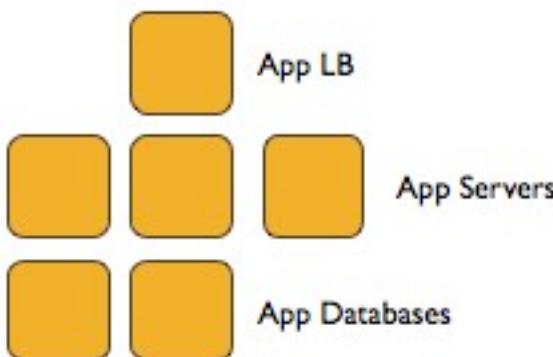


App Databases

# Application server redundancy



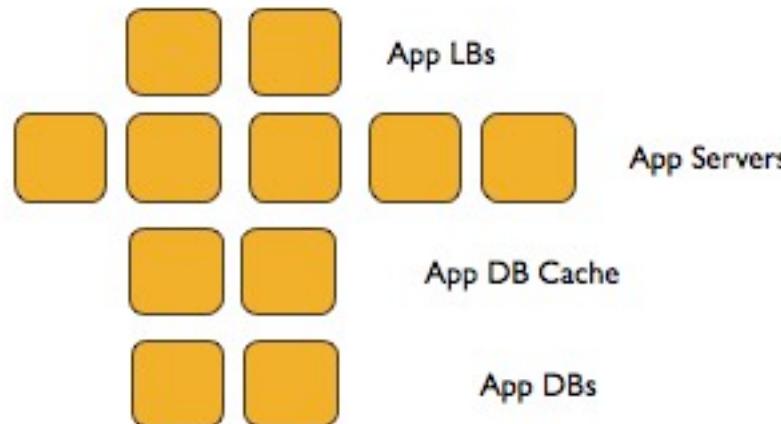
## Add a load balancer



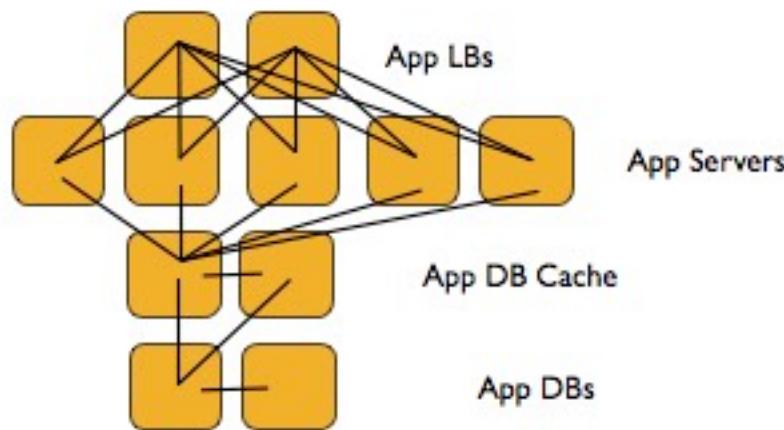
# Webscale!



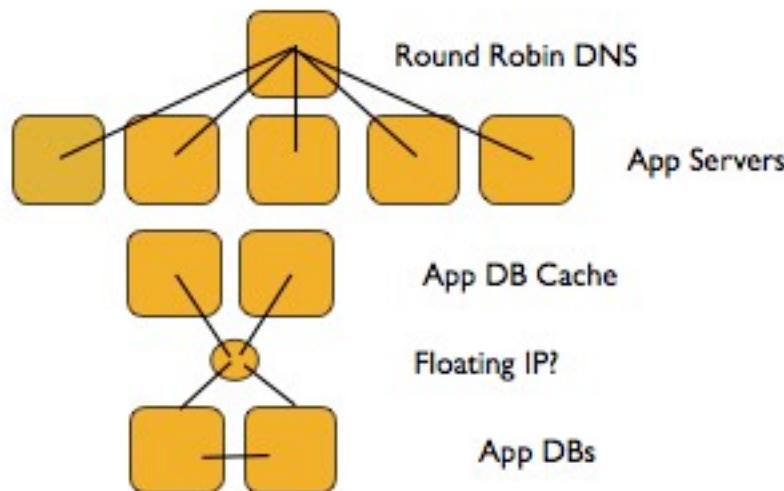
Now we need a caching layer



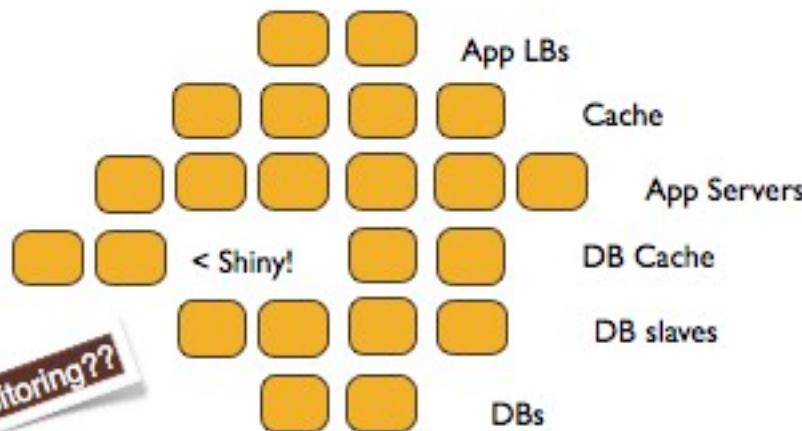
# Infrastructure has a Topology



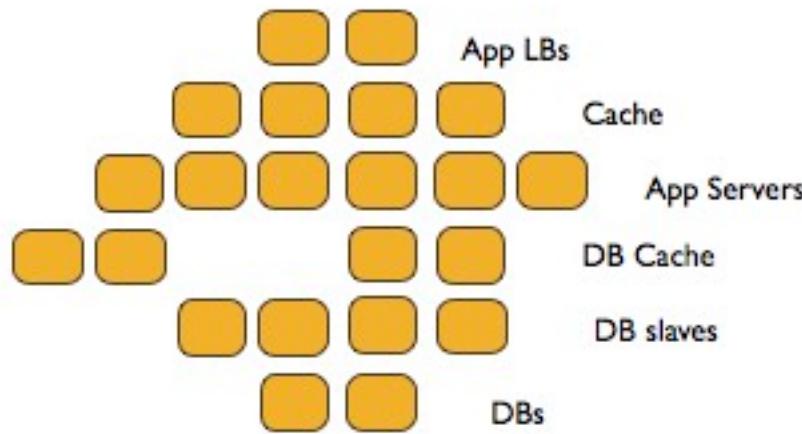
# Your Infrastructure is a Snowflake



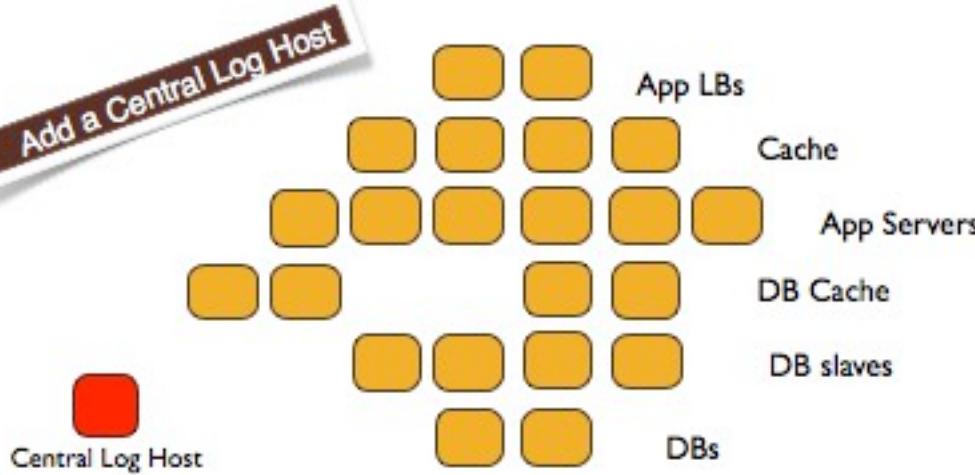
# Complexity Increases Quickly



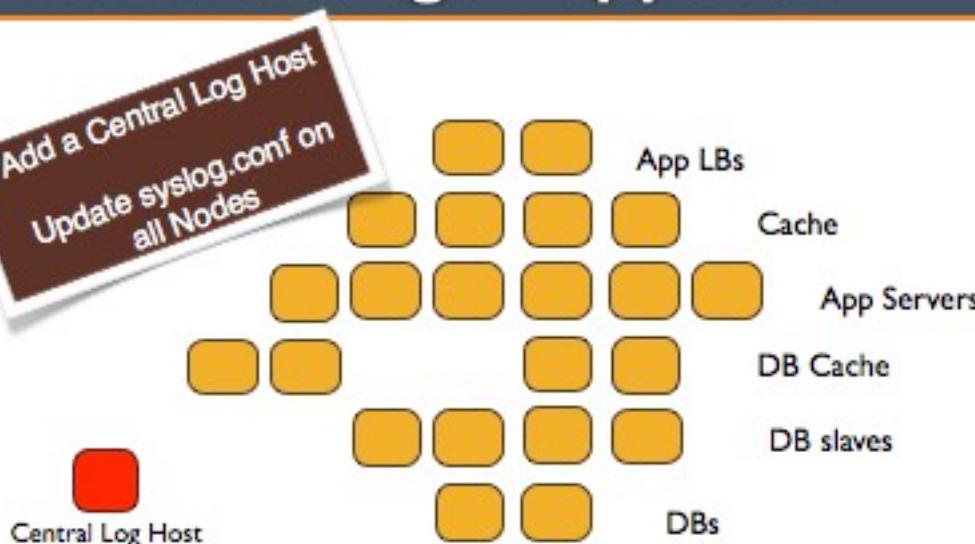
...and change happens!



# ...and change happens!



# ...and change happens!



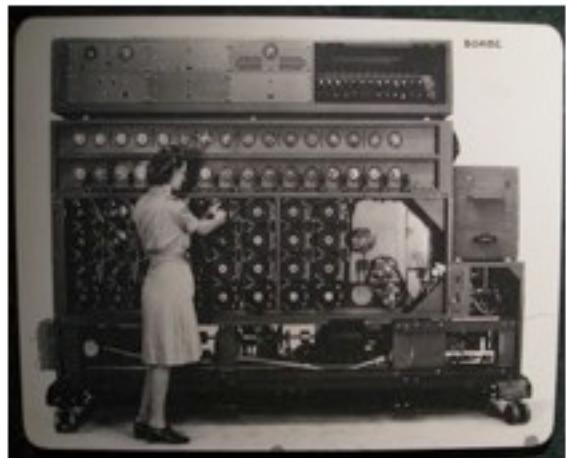
## Chef Solves This Problem



- But you already guessed that, didn't you?

**CHEF™**  
GETCHEF.COM

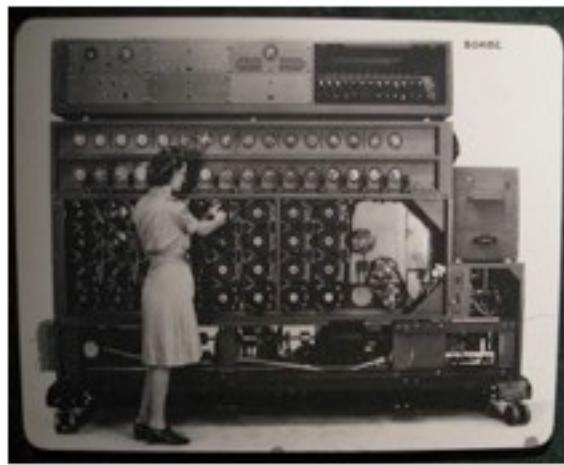
## Chef is Infrastructure as Code



- Programmatically provision and configure components

# Chef is Infrastructure as Code

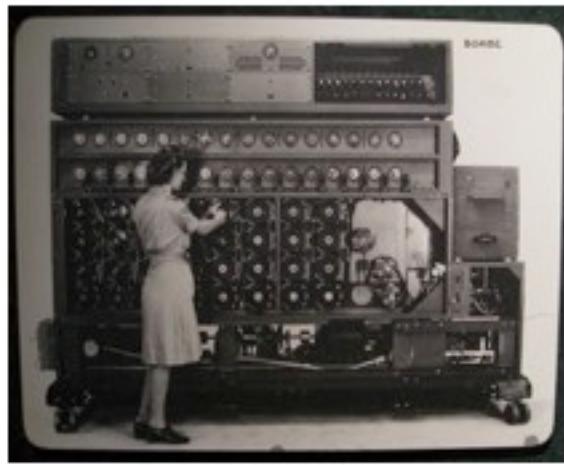
- Treat like any other code base



<http://www.flickr.com/photos/boobie/151295180/>

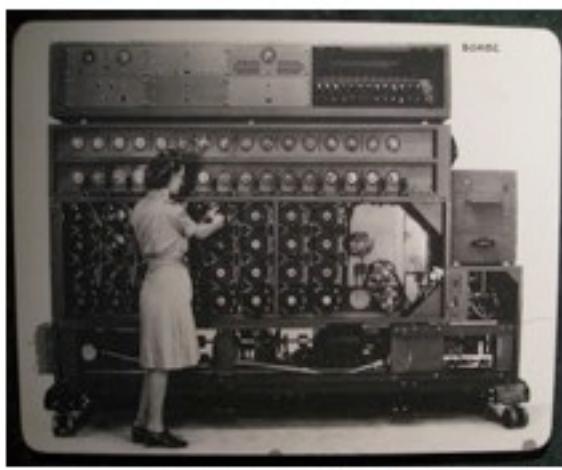
# Chef is Infrastructure as Code

- Reconstruct business from **code repository**, **data backup**, and **compute resources**



<http://www.flickr.com/photos/boobie/151295180/>

# Chef is Infrastructure as Code



<http://www.flickr.com/photos/soothi/151295180/>

- Programmatically provision and configure components
- Treat like any other code base
- Reconstruct business from **code repository**, **data backup**, and **compute resources**

## Configuration Code

- Chef ensures each Node complies with the policy
- Policy is determined by the configurations in each Node's run list
- Reduce management complexity through abstraction
- Store the configuration of your infrastructure in version control

## Declarative Interface to Resources

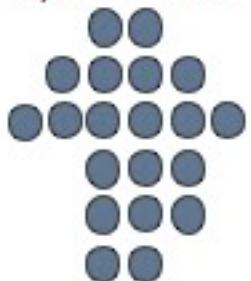
- You define the policy in your Chef configuration
- Your policy states what state each resource should be in, but not how to get there
- Chef-client will pull the policy from the Chef Server and enforce the policy on the Node

## Managing Complexity

- Organizations
- Environments
- Roles
- Nodes
- Recipes
- Cookbooks
- Search

# Organizations

My Infrastructure



Your Infrastructure



Their Infrastructure

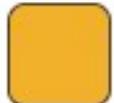


## Organizations

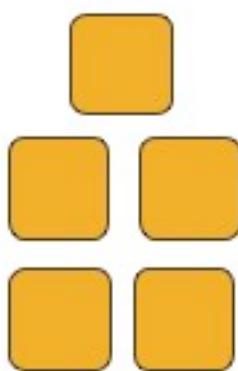
- Completely independent tenants of Enterprise Chef
- Share nothing with other organizations
- May represent different
  - Companies
  - Business Units
  - Departments

# Environments

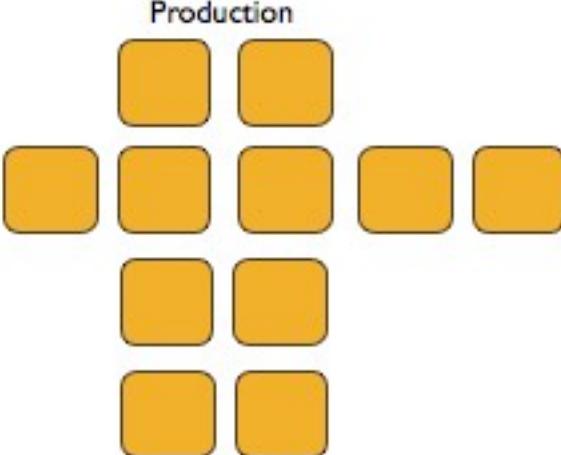
Development



Staging



Production



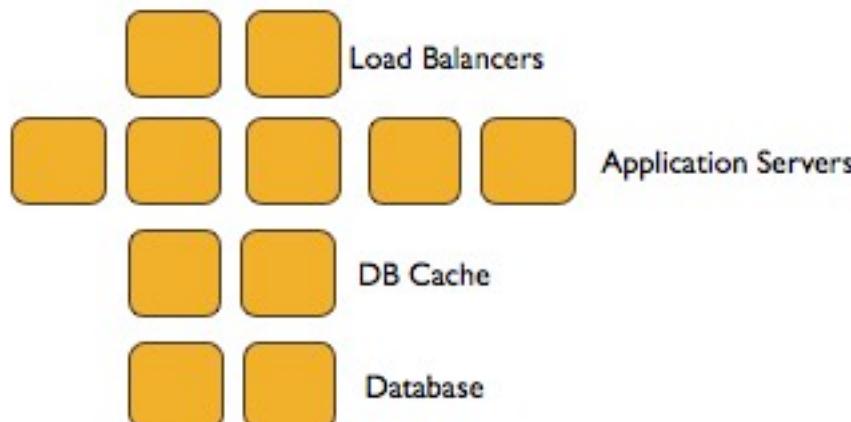
## Environments

- Environments reflect your patterns and workflow, and can be used to model the life-stages of your applications
  - Development
  - Test
  - Staging
  - Production
  - etc.
- Every Organization starts with a single environment!

## Environments Define Policy

- Environments may include data attributes necessary for configuring your infrastructure, e.g.
  - The URL of your payment service's API
  - The location of your package repository
  - The version of the Chef configuration files that should be used

## Roles



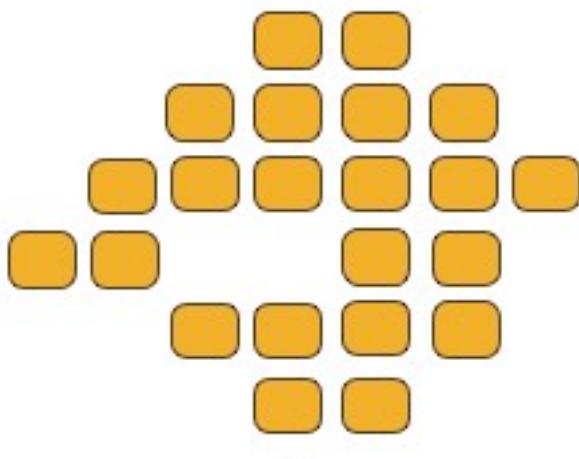
## Roles

- Roles represent the types of servers in your infrastructure
  - Load Balancer
  - Application Server
  - Database Cache
  - Database
  - Monitoring

## Roles Define Policy

- Roles may include an ordered list of Chef configuration files that should be applied
  - This list is called a Run List
  - Order is always important in the Run List
- Roles may include data attributes necessary for configuring your infrastructure, for example:
  - The port that the application server listens on
  - A list of applications that should be deployed

# Nodes



## Nodes

- Nodes represent the servers in your infrastructure
  - Could be physical servers or virtual servers
  - May represent hardware that you own or compute instances in a public or private cloud
- Could also be network hardware - switches, routers, etc

## Node

- Each Node will
  - Belong to one Organization
  - Belong to one Environment
  - Have zero or more Roles

## Nodes Adhere to Policy

- The chef-client application runs on each node, which
  - Gathers the current system configuration of the node
  - Downloads the desired system configuration policies from the Chef server for that node
  - Configures the node such that it adheres to those policies

## Resources

- A Resource represents a piece of the system and its desired state
  - A package that should be installed
  - A service that should be running
  - A file that should be generated
  - A cron job that should be configured
  - A user that should be managed
  - and more

## Resources in Recipes

- Resources are the fundamental building blocks of Chef configuration
- Resources are gathered into Recipes
- Recipes ensure the system is in the desired state

# Recipes

- Configuration files that describe resources and their desired state
- Recipes can:
  - Install and configure software components
  - Manage files
  - Deploy applications
  - Execute other recipes
  - and more

## Example Recipe

```
package "apache2"
```

```
template "/etc/apache2/apache2.conf" do
  source "apache2.conf.erb"
  owner "root"
  group "root"
  mode "0644"
  variables(:allow_override => "All")
  notifies :reload, "service[apache2]"
end
```

```
service "apache2" do
  action [:enable,:start]
  supports :reload => true
end
```

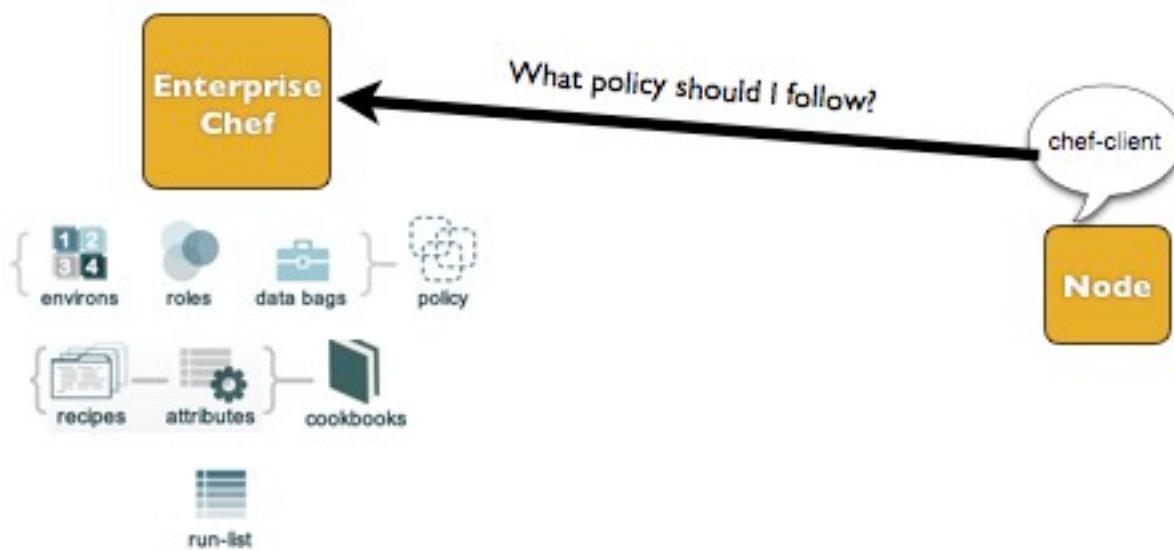
# Cookbooks

- Recipes are stored in Cookbooks
- Cookbooks contain recipes, templates, files, custom resources, etc
- Code re-use and modularity

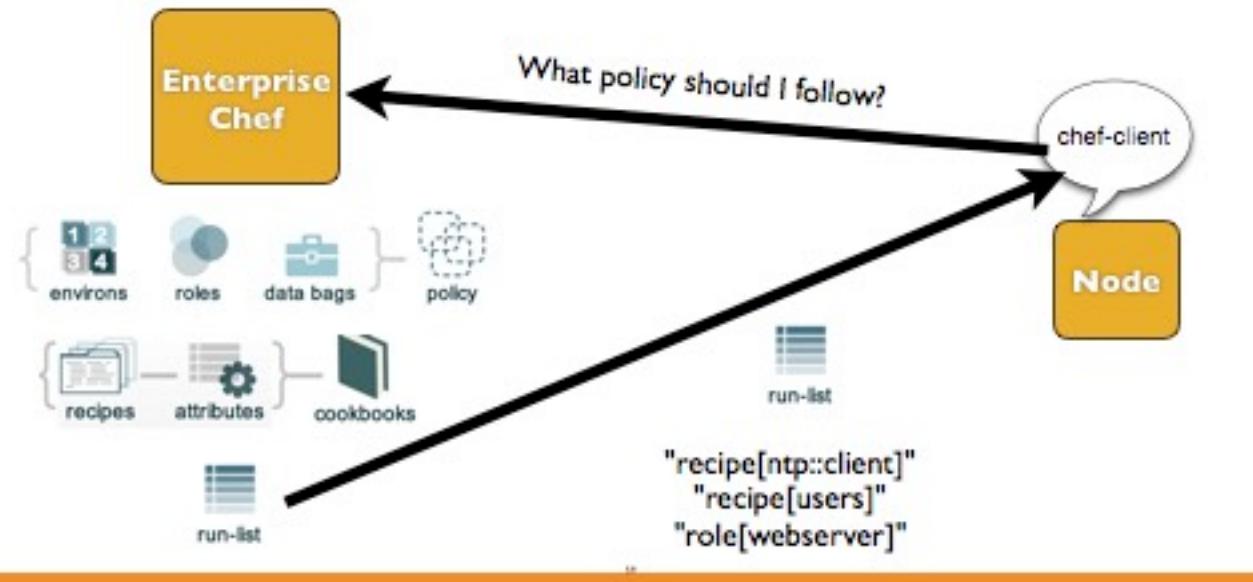


<http://www.flickr.com/photos/shutterhacks/4474421855/>

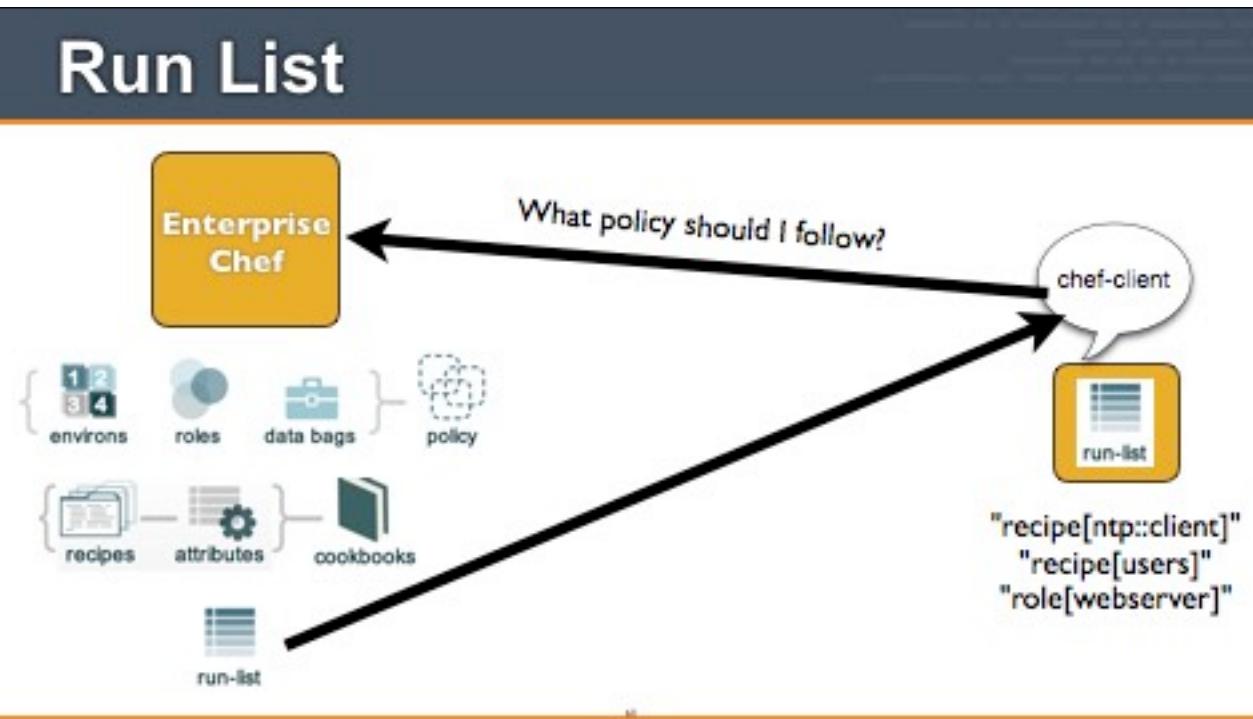
## Run List



# Run List



# Run List



## Run List Specifies Policy

- The Run List is an ordered collection of policies that the Node should follow
- Chef-client obtains the Run List from the Chef Server
- Chef-client ensures the Node complies with the policy in the Run List

## Search

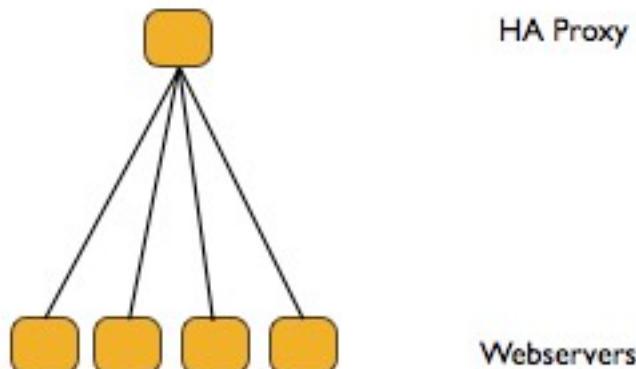
- Search for nodes with Roles
- Find Topology Data
- IP addresses
- Hostnames
- FQDNs



## Search for Nodes

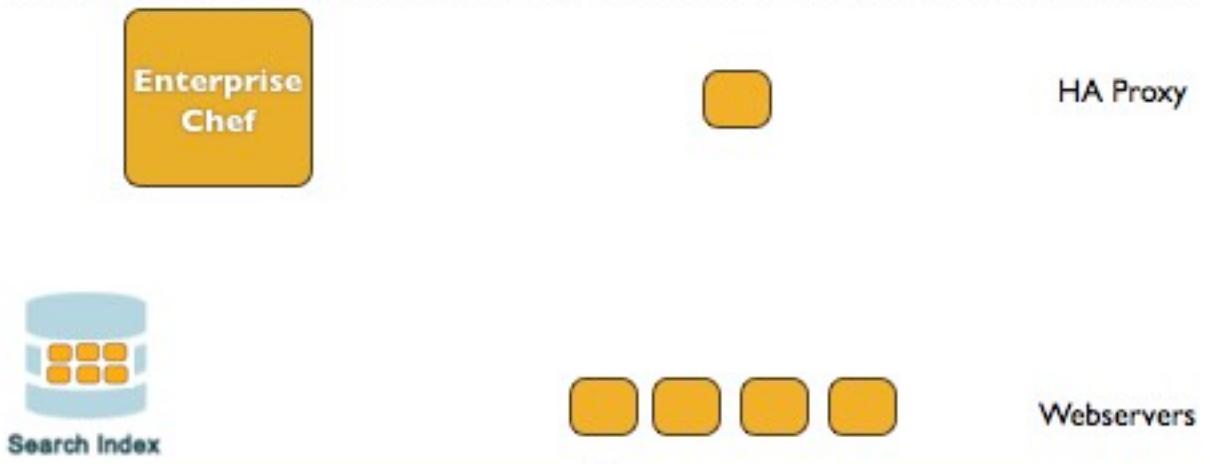
```
pool_members = search("node", "role:webserver")  
  
template "/etc/haproxy/haproxy.cfg" do  
  source "haproxy-app_lb.cfg.erb"  
  owner "root"  
  group "root"  
  mode 0644  
  variables :pool_members => pool_members.uniq  
  notifies :restart, "service[haproxy]"  
end
```

## HAProxy Configuration



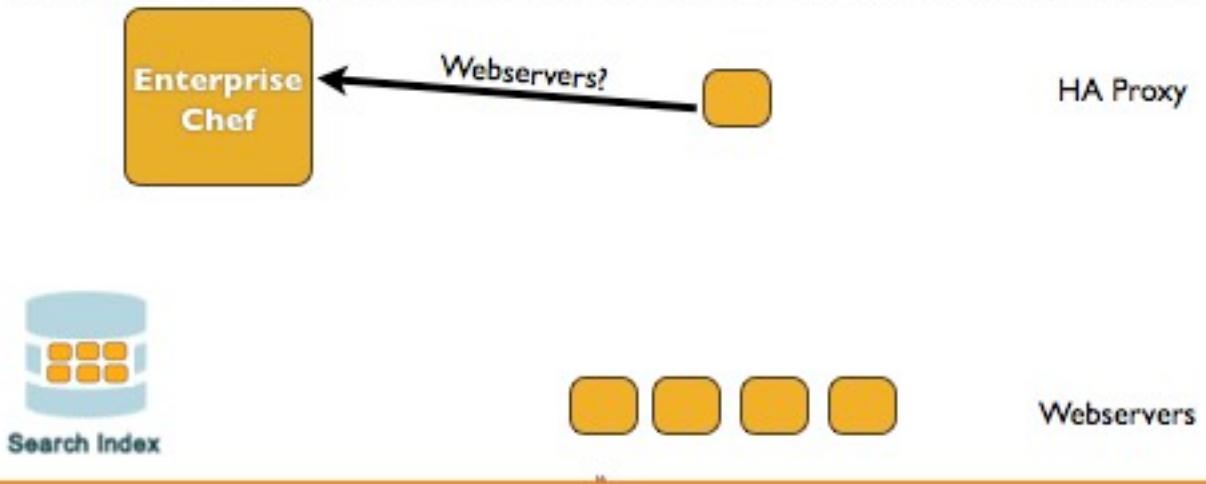
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



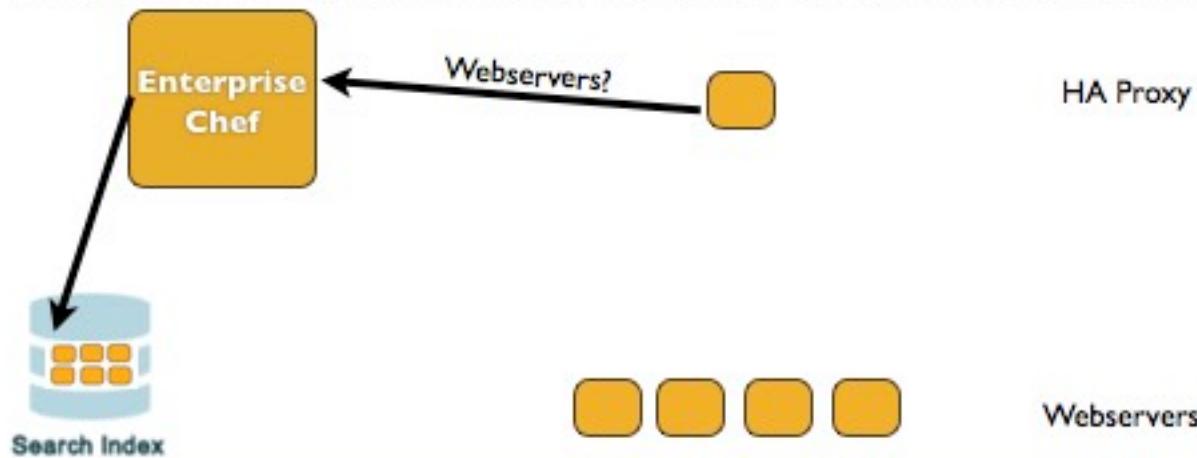
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



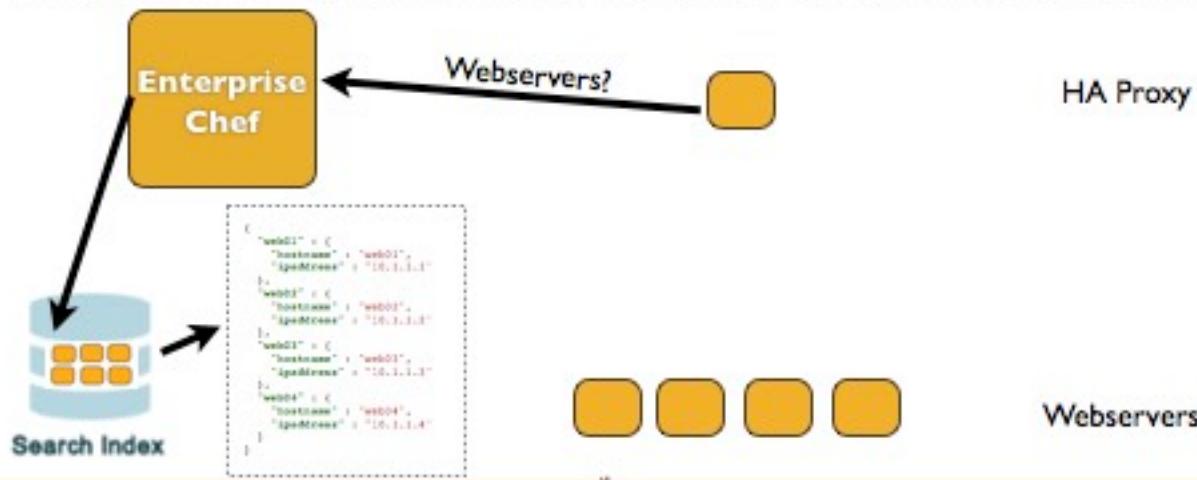
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



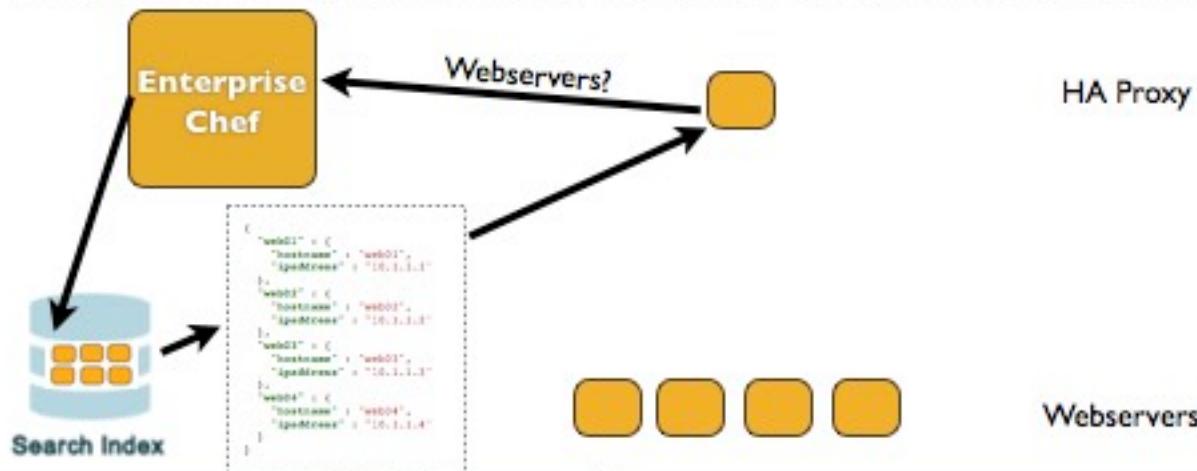
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



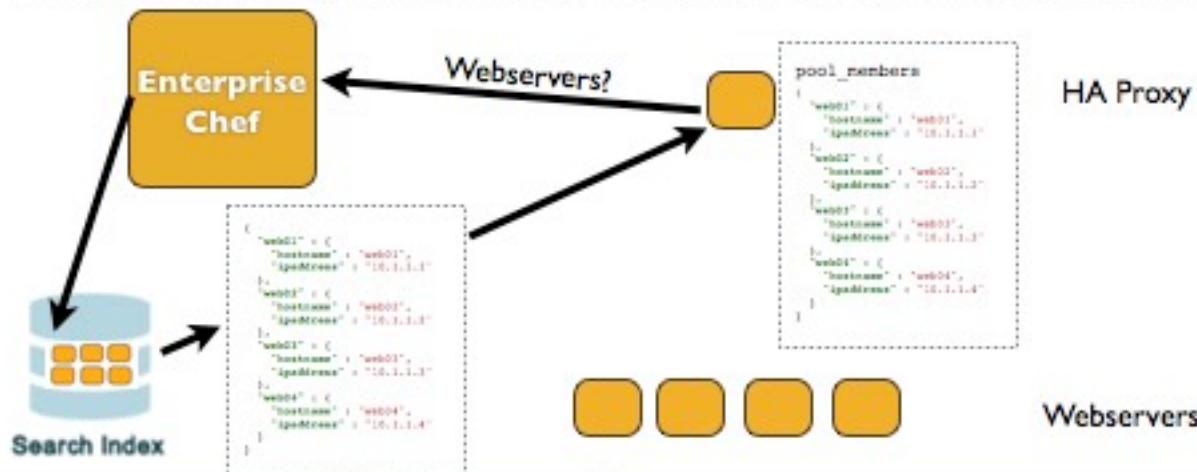
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



# Search for Nodes

```
pool_members = search("node", "role:webserver")\n\n  template "/etc/haproxy/haproxy.cfg" do\n    source "haproxy-app_lb.cfg.erb"\n    owner "root"\n    group "root"\n    mode 0644\n    variables :pool_members => pool_members.uniq\n    notifies :restart, "service[haproxy]"\n  end
```

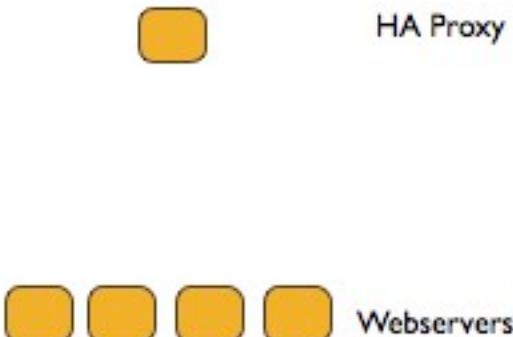
# Pass results into Templates

```
# Set up application listeners here.\n\nlisten application 0.0.0.0:80\n  balance roundrobin\n  <@ pool_members.each do |member| -@>\n    server <@= member[:hostname] @> <@= member[:ipaddress] @>; weight 1 maxconn 1 check\n  <@ end -@>\n<@ if node["haproxy"]["enable_admin"] -@>\nlisten admin 0.0.0.0:22002\n  mode http\n  stats uri /\n<@ end -@>
```

# HAProxy Configuration

```
<@ pool_members.each do |member| -t>
server <@= member(:hostname) ><@= member(:ipaddress) >i> weight 1 maxconn 1 check
<@ end -t>
```

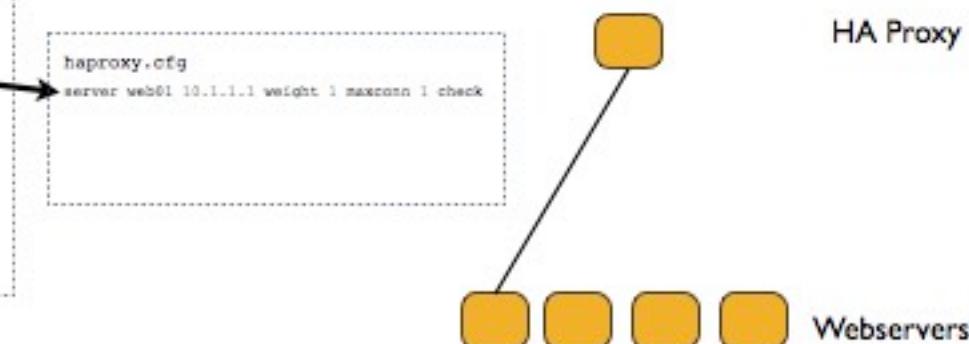
```
pool_members
{
  "web01" : {
    "hostname" : "web01",
    "ipaddress" : "10.1.1.1"
  },
  "web02" : {
    "hostname" : "web02",
    "ipaddress" : "10.1.1.2"
  },
  "web03" : {
    "hostname" : "web03",
    "ipaddress" : "10.1.1.3"
  },
  "web04" : {
    "hostname" : "web04",
    "ipaddress" : "10.1.1.4"
  }
}
```



# HAProxy Configuration

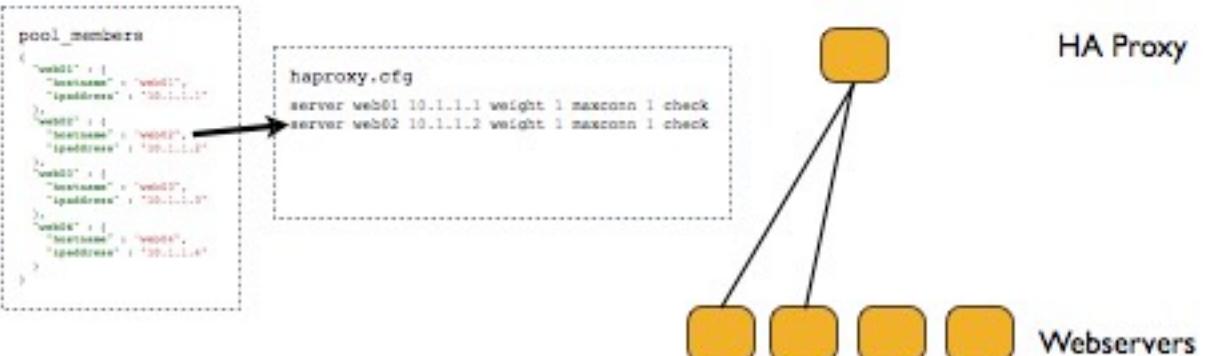
```
<@ pool_members.each do |member| -t>
server <@= member(:hostname) ><@= member(:ipaddress) >i> weight 1 maxconn 1 check
<@ end -t>
```

```
pool_members
{
  "web01" : {
    "hostname" : "web01",
    "ipaddress" : "10.1.1.1"
  },
  "web02" : {
    "hostname" : "web02",
    "ipaddress" : "10.1.1.2"
  },
  "web03" : {
    "hostname" : "web03",
    "ipaddress" : "10.1.1.3"
  },
  "web04" : {
    "hostname" : "web04",
    "ipaddress" : "10.1.1.4"
  }
}
```



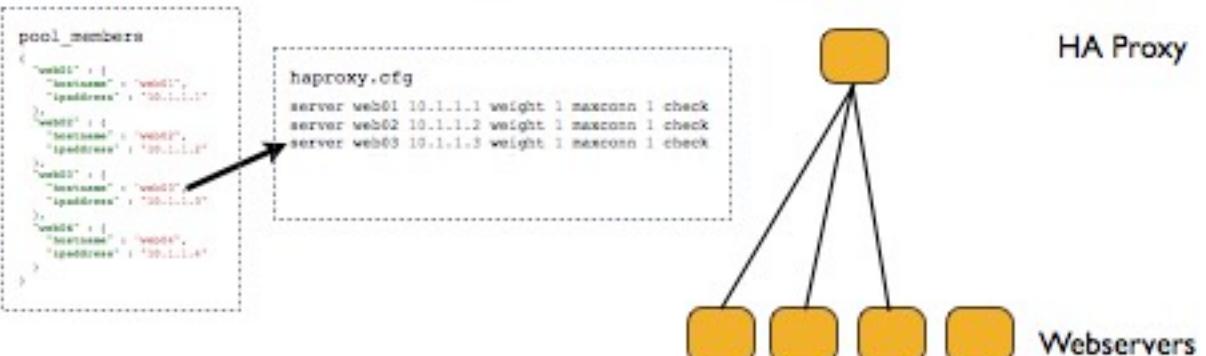
# HAProxy Configuration

```
<% @pool_members.each do |member| -%>
server <%= member(:hostname) %> <%= member(:ipaddress) %>; weight 1 maxconn 1 check
<% end -%>
```



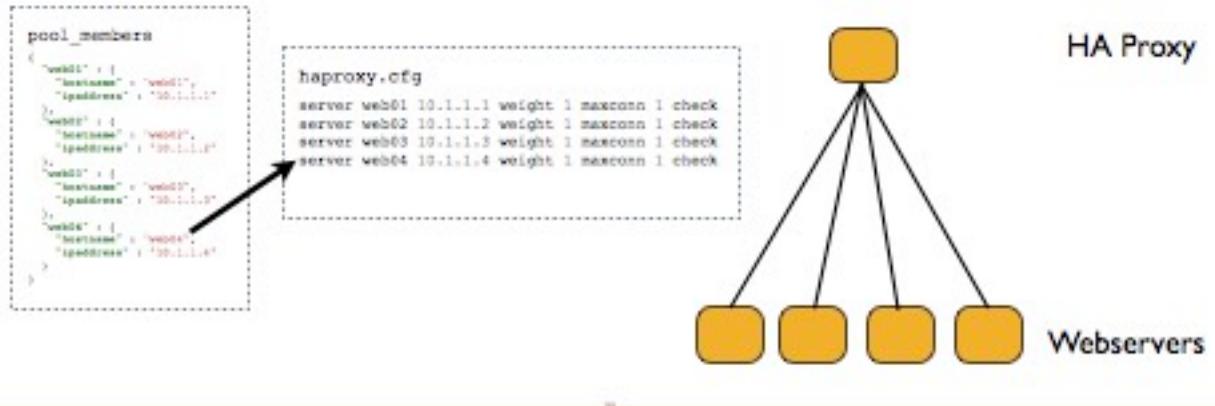
# HAProxy Configuration

```
<% @pool_members.each do |member| -%>
server <%= member(:hostname) %> <%= member(:ipaddress) %>; weight 1 maxconn 1 check
<% end -%>
```

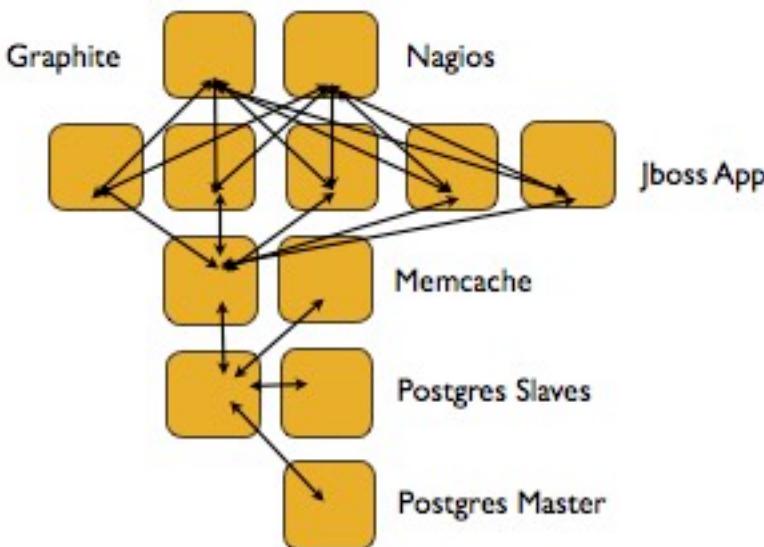


# HAProxy Configuration

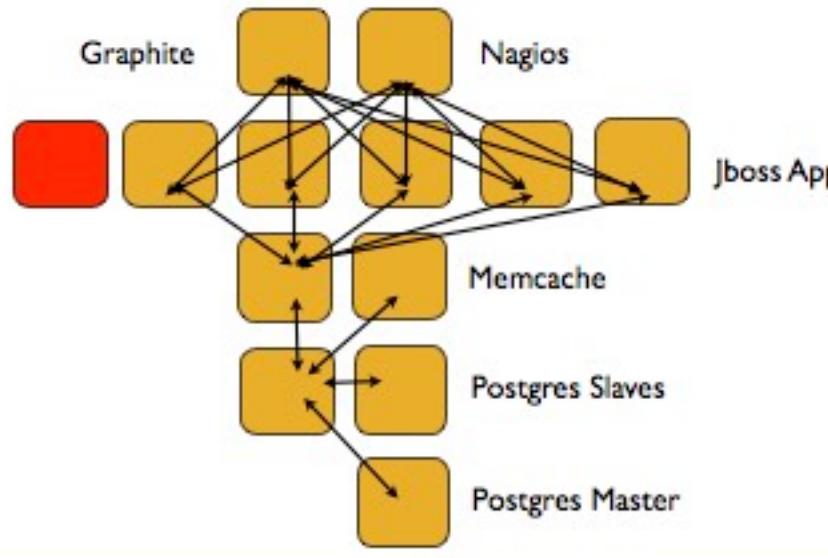
```
<% @pool_members.each do |member| -%>
server <%= member(:hostname) %> <%= member(:ipaddress) %>; weight 1 maxconn 1 check
<% end -%>
```



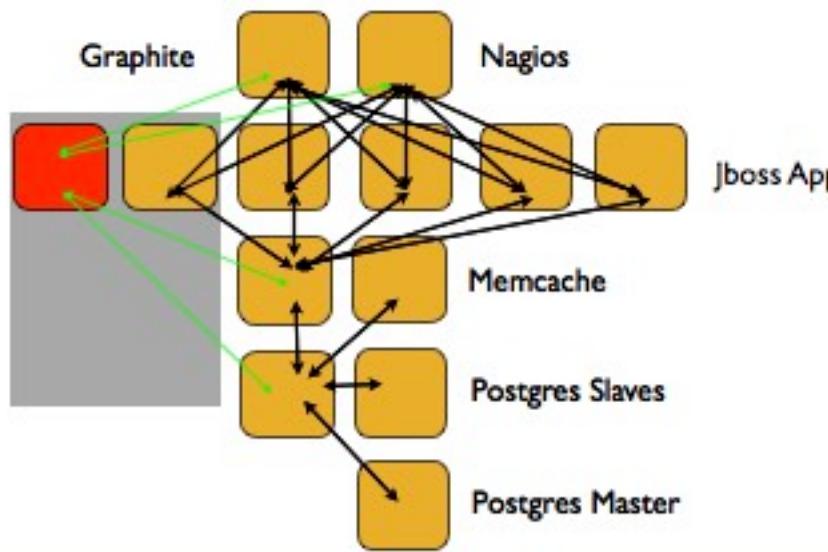
So when this...



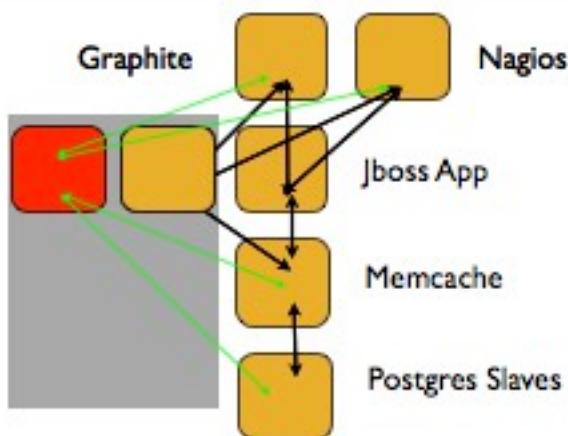
# ...becomes this



# ...this can happen automatically



# Count the Resources



- 12+ resource changes for 1 node addition

- Load balancer config
- Nagios host ping
- Nagios host ssh
- Nagios host HTTP
- Nagios host app health
- Graphite CPU
- Graphite Memory
- Graphite Disk
- Graphite SNMP
- Memcache firewall
- Postgres firewall
- Postgres authZ config

# Manage Complexity

- Determine the desired state of your infrastructure
- Identify the Resources required to meet that state
- Gather the Resources into Recipes
- Compose a Run List from Recipes and Roles
- Apply a Run List to each Node in your Environment
- Your infrastructure adheres to the policy modeled in Chef

## Configuration Drift

- Configuration Drift happens when:
  - Your infrastructure requirements change
  - The configuration of a server falls out of policy
- Chef makes it easy to manage
  - Model the new requirements in your Chef configuration files
  - Run the chef-client to enforce your policies

## Design Tenets of Chef

- Whipuptitude - aptitude for whipping things up
- Reasonability
- Sane defaults
- Flexibility
- Manipulexity - manipulation of complex things

# Review Questions

- What is a Node?
- What is a Resource?
- What is a Recipe? How is it different from a Cookbook?
- What is a Run List?
- What is a Role?

## Workstation Setup

Getting started

# Lesson Objectives

- After completing the lesson, you will be able to
- Login to Enterprise Chef
- View your Organization in Enterprise Chef
- Describe Knife, the Chef command line utility
- Use Knife on your Workstation

## Legend

## Legend: Do I run that command on my workstation?

This is an example of a command to run on your workstation

```
$ whoami  
i-am-a-workstation
```

This is an example of a command to run on your target node via SSH.

```
user@hostname:~$ whoami  
i-am-a-chef-node
```

## Legend: Example Terminal Command and Output

```
$ ifconfig
```

```
lo0: flags=8049<UP,LOOPBACK,RUNNING,MULTICAST> mtu 16384  
      options=3<RXCSUM,TXCSUM>  
      inet6 fe80::1%lo0 prefixlen 64 scopeid 0x1  
        inet 127.0.0.1 netmask 0xff000000  
          inet6 ::1 prefixlen 128  
            gif0: flags=8010<POINTOPOINT,MULTICAST> mtu 1280  
stf0: flags=0<> mtu 1280  
en0: flags=8843<UP,BROADCAST,SMART,RUNNING,SIMPLEX,MULTICAST> mtu 1500  
      ether 28:c9:e9:1f:79:a3  
      inet6 fe80::2ac9:e9ff:fe1f:79a3%en0 prefixlen 64 scopeid 0x4  
        inet 10.100.0.84 netmask 0xffffffff broadcast 10.100.0.255  
          media: autoselect  
            status: active  
p2p0: flags=8843<UP,BROADCAST,RUNNING,SIMPLEX,MULTICAST> mtu 2304  
      ether 6a:c9:e9:1f:79:a3  
      media: autoselect  
        status: inactive
```

## Legend: Example of editing a file on your workstation

 OPEN IN EDITOR: ~/hello\_world

Hi!

I am a friendly file.

**SAVE FILE!**

## Landscape of a Chef-managed Infrastructure



## Landscape of a Chef-managed Infrastructure



## Install Chef

- Install Chef (if not already installed)
- <http://www.chef.io/chef/install>

# Install Chef

## Download options

Chef Client Chef Server

### Installing the Chef Client

Select the kind of system you would like to install the Chef Client on.  
The versions listed have been tested and are supported.

[Select an Operating System]

[Select a Version]

[Select an Architecture]

16

# Install on Mac OSX

## Download options

Chef Client Chef Server

### Installing the Chef Client

Select the kind of system you would like to install the Chef Client on.  
The versions listed have been tested and are supported.

#### Quick Install Instructions

Open a shell on the target system and run the following command to download and install the latest version of the Chef client:

```
curl -L https://www.opscode.com/chef/install.sh | sudo bash
```

#### Downloads

You can install manually by downloading the package below after you have selected a Chef version. For more information about manual installation, [please read the documentation](#).

OS X

10.7

x86\_64

11.8.2-1

chef-11.8.2-1.macosx.10.7.2.sh

# Install on Enterprise Linux

## Download options

Chef Client

Chef Server

### Installing the Chef Client

Select the kind of system you would like to install the Chef Client on. The versions listed have been tested and are supported.

### Quick Installation Instructions

Open a root shell on the target system and run the following command to download and install the latest version of the Chef client:

```
curl -L https://www.opscode.com/chef/install.sh | bash
```

### Downloads

You can install manually by downloading the package below after you have selected a Chef version. For more information about manual installation, [please read the documentation](#).

Enterprise Linux

6

x86\_64

11.8.2-1

chef-11.8.2-1.el6.x86\_64.rpm

# Workstation Setup - Mac OS X / Linux

```
$ curl -L http://www.chef.io/chef/install.sh | sudo bash
```

```
* Total    * Received * Xferd  Average Speed   Time     Time      Time  Current
          *       *        *      Download Upload  Total   Spent   Left  Speed
100 14100  100 14100    0      0  8814      0  0:00:01  0:00:01  ---:---  8812
Downloading Chef for mac_os_x...
downloading http://www.getchef.com/chef/metadata?v=aprerelease=false&p=mac_os_x&pv=10.7&m=x86_64
to file /tmp/install.sh.79770/metadata.txt
trying curl...
url http://opscode-omnibus-packages.s3.amazonaws.com/mac_os_x/10.7/x86_64/chef-11.8.2_1.mac_os_x.
10.7.2.sh
md5 af157c6eff941e52ff69a9dd6a3b57f597
sha256c003c0951d80245b1f02c4588f9157a55e7b94a00dd6ac163aed8ef2e854619a
downloaded metadata file looks valid...
downloading http://opscode-omnibus-packages.s3.amazonaws.com/mac_os_x/10.7/x86_64/
chef-11.8.2_1.mac_os_x.10.7.2.sh
to file /tmp/install.sh.79770/chef--mac_os_x-10.7-x86_64.sh
trying curl...
-----
Thank you for installing Chef!
```

# Workstation Setup - Windows

- Windows
- 2008 (Windows 7) or 2012 (Windows 8)
- i686 (32-bit) or x86\_64 (64-bit)
- 11.8.2

Download options

Chef Client   Chef Server

Installing the Chef Client

Select the kind of system you would like to install the Chef Client on. The versions listed have been tested and are supported.

Downloads

You can install manually by downloading the package. For more information about manual installation, please refer to the documentation.

11.8.2-1

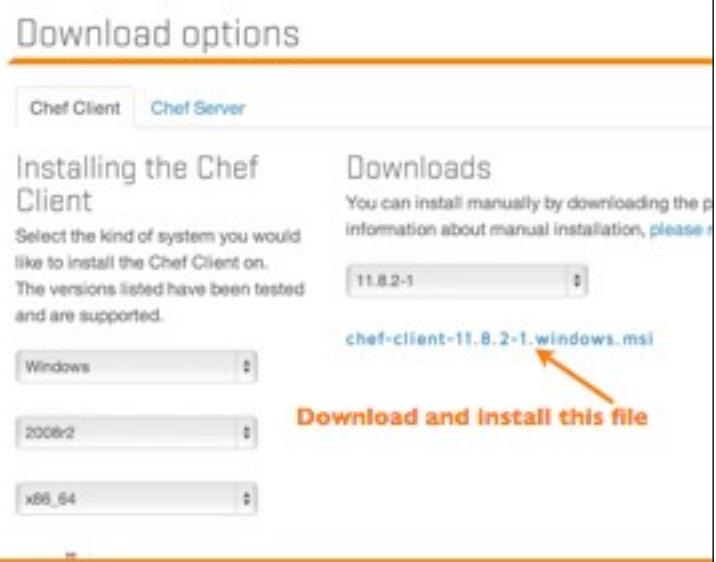
chef-client-11.8.2-1.windows.msi

Windows

2008x2

x86\_64

Download and Install this file



## Install on Windows



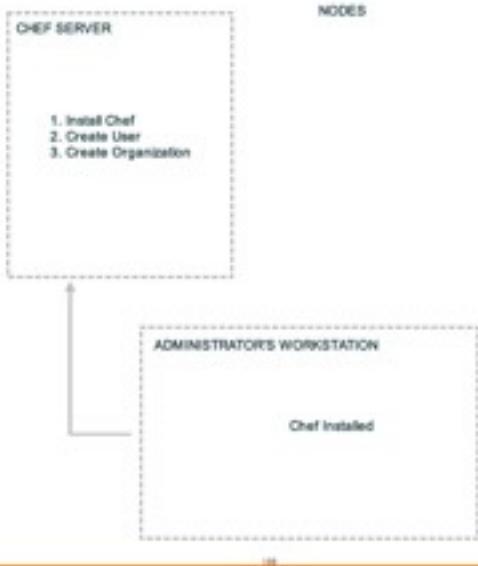
# What just happened?

- Chef and all of its dependencies installed via an operating system-specific package ("omnibus installer")
- Installation includes
  - The Ruby language - used by Chef
  - knife - Command line tool for administrators
  - chef-client - Client application
  - ohai - System profiler
  - ...and more

## Workstation or Node?



# Landscape of a Chef-managed Infrastructure



## Your Chef Server for this class...

- Hosted Enterprise Chef <http://www.chef.io>

WHAT IS CHEF LEARN CHEF RESOURCES **GET CHEF**

Automation for Web-Scale IT  
Chef delivers fast, scalable, flexible IT automation.

Choose your installation  
Looking for the Chef Development Kit, Chef Client or Open-Source Chef Server? [GET CHEF](#)

**HOSTED CHEF**  
Let us manage  
Chef Server for you  
[GET HOSTED CHEF](#)

**ON PREMIS CHEF**  
Install & manage  
your own Chef Server  
[GET ON PREMIS CHEF](#)

Hosted Enterprise Chef is the  
quickest and easiest way to  
get started with Chef!  
You don't have to worry about  
updates or maintenance. Just

If you need to run Chef in  
a firewall, or just want to  
set up and manage your own  
server, choose to install  
Enterprise Chef on-prem.

# Create new account

- Sign up for a new account
- Chef Organization
  - provides multi-tenancy
  - name must be globally unique

## Start your free trial of Enterprise Chef

You're one step away from access to all the power and flexibility of Chef, hosted and supported by Opscode. Get ready to automate your infrastructure to accelerate your time-to-market, manage complexity, and safeguard your systems. Just complete the form to get started.

Full Name

Username

Email

Password

Company  (Optional)

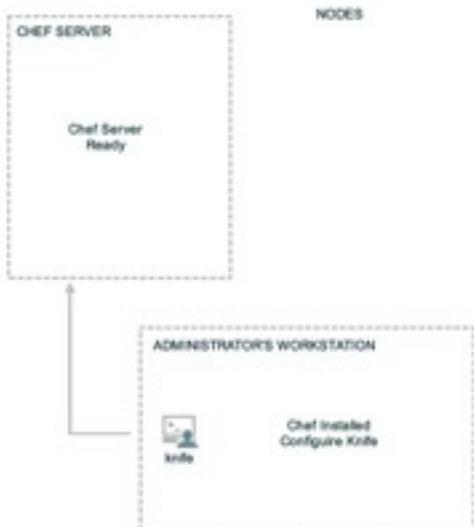
Chef Organization

Organization is the name of your instance of Enterprise Chef.

I agree to the [Terms of Service](#) and the [Master License and Services Agreement](#).

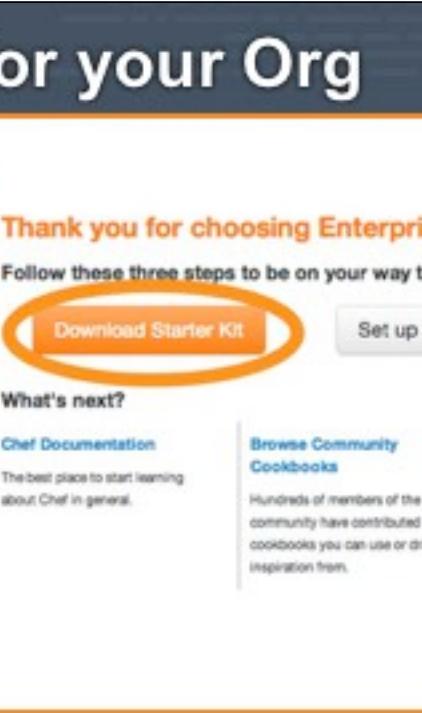
[Get Started](#)

# Landscape of a Chef-managed Infrastructure



## Download "Starter Kit" for your Org

- You get a .zip file from clicking this
- Unzip the zipfile - you'll get a "chef-repo"
- Put the "chef-repo" somewhere, e.g.:
  - C:\Users\you\chef-repo (Win)
  - /Users/you/chef-repo (Mac)
  - /home/you/chef-repo (Linux)



## What is the Starter Kit?

- The 'Starter Kit' is a archive file (e.g. chef-starter.zip) that contains
  - A sample chef repository with a sample 'starter' cookbook
  - Configuration files allowing the workstation to talk to the chef server using knife
- We'll look at this in detail later...

## Knife is the command-line tool for Chef

- Knife provides an API interface between a local Chef repository and the Chef Server, and lets you manage:
  - Nodes
  - Cookbooks and recipes
  - Roles
  - Stores of JSON data (data bags), including encrypted data
  - Environments
  - Cloud resources, including provisioning
  - The installation of Chef on management workstations
  - Searching of indexed data on the Chef Server

## A quick tour of the chef-repo

- Every infrastructure managed with Chef has a Chef Repository ("chef-repo")
- Type all commands in this class from the chef-repo directory
- Let's see what's inside the chef-repo...

## Tour the chef-repo

```
$ cd chef-repo
```

```
[~/chef-repo]$
```

## Tour the chef-repo

```
$ ls -al
```

```
total 40
drwxr-xr-x@ 11 opscode  opscode   374 Dec 15 09:42 .
drwxr-xr-x@ 92 opscode  opscode  3128 Dec 15 09:43 ..
drwxr-xr-x@  3 opscode  opscode   102 Dec 15 2013 .berkshelf
drwxr-xr-x@  5 opscode  opscode   170 Dec 15 2013 .chef
-rw-r--r--@  1 opscode  opscode   495 Dec 15 2013 .gitignore
-rw-r--r--@  1 opscode  opscode  1433 Dec 15 2013 Berksfile
-rw-r--r--@  1 opscode  opscode  2416 Dec 15 2013 README.md
-rw-r--r--@  1 opscode  opscode  3567 Dec 15 2013 Vagrantfile
-rw-r--r--@  1 opscode  opscode   588 Dec 15 2013 chefignore
drwxr-xr-x@  3 opscode  opscode   102 Dec 15 2013 cookbooks
drwxr-xr-x@  3 opscode  opscode   102 Dec 15 2013 roles
```

## What's inside the .chef directory?

```
$ ls .chef
```

```
ORGNAME-validator.pem  
USERNAME.pem  
knife.rb
```

## What's inside the .chef directory?

- `knife.rb` is the configuration file for Knife.
- The other two files are certificates for authentication with the Chef Server
  - We'll talk more about that later.

# knife.rb

- Default location
  - `~/.chef/knife.rb`
  - `c:\Users\You\.chef` (Windows)
- Use a project specific configuration
  - `.chef/knife.rb` of the current directory
  - `chef-repo/.chef/knife.rb`
- [http://docs.chef.io/config\\_rb\\_knife.html](http://docs.chef.io/config_rb_knife.html)

# knife.rb



OPEN IN EDITOR: `chef-repo/.chef/knife.rb`

```
current_dir = File.dirname(__FILE__)
log_level           :info
log_location        STDOUT
node_name           "USERNAME"
client_key          "#{current_dir}/USERNAME.pem"
validation_client_name "ORGNAME-validator"
validation_key       "#{current_dir}/ORGNAME-validator.pem"
chef_server_url     "https://api.opscode.com/organizations/ORGNAME"
cache_type          'BasicFile'
cache_options( :path => "#{ENV['HOME']}/.chef/checksums" )
cookbook_path        ["#{current_dir}/../cookbooks"]
```

# Verify Knife

```
$ knife --version  
Chef: 11.8.2
```

- Your version may be different, that's ok!



```
$ knife client list  
ORGNAME-validator
```

## knife client list

1. Reads the `chef_server_url` from `knife.rb`
2. Invokes HTTP GET to `#{{chef_server_url}}/clients`
3. Displays the result



## Exercise: Run ‘knife help list’

```
$ knife help list
```

```
Available help topics are:
```

```
bootstrap
chef-shell
client
configure
cookbook
cookbook-site
data-bag
delete
deps
diff
download
edit
environment
exec
list
```

## A few knife tips

- Commands are always structured as follows
  - knife
  - NOUN (client)
  - VERB (list)
- You can get more help with
  - knife NOUN help
  - knife --help just shows options

## Best Practice: Use a text editor with a project drawer, or equivalent

- Chef is about Infrastructure as Code
- People who code for a living use text editors that are designed for the task
  - Vim, Emacs, Sublime Text, Notepad++, etc.

## Benefits of a good text editor

- A good editor will
  - Show line numbers
  - Highlight syntax
  - Autocomplete commands
  - Allow you to open multiple files

If you do not have a preferred text editor on your workstation already...

- Download Sublime Text
- Free trial, not time bound
- Works on every platform
- [sublimetext.com](http://sublimetext.com)



## Checkpoint



# What's Next?



# Source Code Repository



## Review Questions

- What is the chef-repo?
- What is knife?
- What is name of the knife configuration file?
- Where is the knife configuration file located?
- What is your favorite text editor? :)

## Organization Setup

Setup an Organization

## Lesson Objectives

- After completing the lesson, you will be able to
- Explain the purpose of Organizations
- Manage your Chef Organization
- Invite users into your organization
- Accept invitations into other organizations

## Organizations

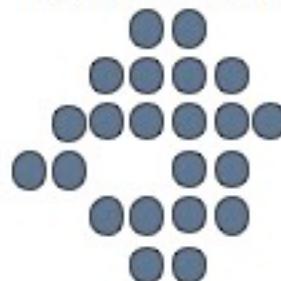
My Infrastructure



Your Infrastructure



Their Infrastructure



# Organizations

- Provide multi-tenancy in Enterprise Chef
- Nothing is shared between Organizations - they're completely independent
- May represent different
  - Companies
  - Business Units
  - Departments

## Manage Organizations

- Login to your Hosted Enterprise Chef



# Organizations

The screenshot shows the Chef Manage interface with the title "Organizations". The top navigation bar includes "Nodes", "Policies", and "Administrative" tabs, with "Administrative" being the active tab. On the left, a sidebar menu lists "Organizations", "Create", "Reset Validation Key", "Generate Knife Config", "Leave Organization", and "Starter Kit". Below this are sections for "Users", "Groups", and "Global Permissions". The main content area displays the message "Showing All Organizations" above a table header "Organization". A single row in the table is shown, labeled "nhcompany".

## Manage Organizations

- Reset Validation Key
- Generate Knife Config
- Leave Organization
- Download the 'Starter Kit'

This screenshot is similar to the one above, showing the "Organizations" page in Chef Manage. The "Administrative" tab is active. The sidebar menu includes "Organizations", "Create", "Reset Validation Key", "Generate Knife Config", "Leave Organization", and "Starter Kit". The main content shows "Showing All Organizations" and a table with a single row "nhcompany". A context menu is open over the "nhcompany" row, listing options: "Reset Validation Key", "Generate Knife Config", "Leave Organization", and "Starter Kit". The top right corner of the screen shows the user "nathanharvey" and the status "Logged in".

## Manage Organizations

- Each Organization may have multiple Users
- Manage an Organization's Users via the Enterprise Server interface
  - <http://manage.opscode.com>

## Exercise: Invite users into your org

- **The Problem:** You need assistance from a colleague on one of your chef projects
- **Proposed Solution:** Invite the colleague into your organization so they can log into the chef server with **their own** credentials, but can see **your** organization

## Exercise: Invite Users into your Org

- For the Lab Exercise you can team up with one other person in the class, then you will
  1. Invite your classmate into your organization
  2. Accept your classmate's invite into their org
  3. Look around your classmate's organization while they look around yours
  4. Finally remove your classmate's access from your account

## Exercise: Invite a classmate into your org

- Navigate to <http://manage.opscode.com>
- Click the "Administration" tab,
- Select the appropriate Organization
- Click "Invite User" from the left menu
- Enter your classmate's 'Chef Username' and click Invite

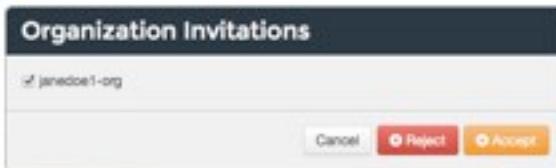


## Exercise: Accept an invite

- You will get a notification once your classmate has invited you into their account



- Click the notification, select the Organization and click 'Accept'



## Exercise: View classmate's org

- Select your classmate's organization from the drop down list and peruse their org

A screenshot of the CHEF interface showing the "Organizations" page. On the left is a sidebar with options like "Create", "Revert Validation-Key", "Generate Knife-Config", "Invite User", "Leave Organization", "Start Kit", "Users", "Groups", and "Global Permissions". The main area shows a table with two organizations: "johndoe1-org" (highlighted with a yellow oval) and "johndoe1-org". The "johndoe1-org" row has a "Details" button. Below the table, there's a section for "Organization: johndoe1-org" with a "Members" table containing one member: "johndoe1". A "Search Members..." input field and "Add Member..." and "Remove" buttons are also present.

# Exercise: Revoke access

- Now either 'Leave Organization' you've been invited into, or remove your classmate from your organization

The screenshot shows the CHEF web interface for managing organizations. On the left, a sidebar lists options like 'Create', 'Reset Validation Key', 'Generate Knife Config', 'Invite User', 'Leave Organization' (which is circled in red), and 'Starter Kit'. The main area shows a list of organizations: 'johndoe1-org' and 'johndoe2-org' (also circled in red). Below this, a specific organization's members are listed: 'johndoe1' (with an 'Actions' button circled in red). The bottom of the screen includes standard navigation links like 'Feedback', 'What Next?', and 'About OpenSUSE Manager'.

## Review Questions

- What is an Organization?
- How do you regenerate the Starter Kit for your Organization?

# Node Setup

Setup a Node to manage

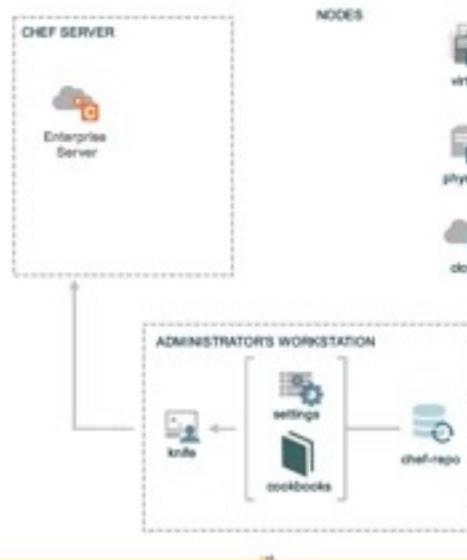
92.1.8



## Lesson Objectives

- After completing the lesson, you will be able to
  - Install Chef nodes using "knife bootstrap"
  - Explain how knife bootstrap configures a node to use the Organization created in the previous section
  - Explain the basic configuration needed to run chef-client

# Nodes



# Nodes

- Nodes represent the servers in your infrastructure  
these may be
  - Physical or virtual servers
  - Hardware that you own
  - Compute instances in a public or private cloud
- Could also be network hardware - switches, routers, etc

# We Have No Nodes Yet

The screenshot shows the Chef Manage web application. At the top, there's a dark header bar with the Chef logo and the word "CHEF" in white. Below it is a navigation bar with three tabs: "Nodes" (which is active and highlighted in blue), "Policies", and "Administrative". On the left side, there's a sidebar with a title "Nodes" and a list of actions: "Delete", "Manage Tags", "Reset Key", "Edit Run List", and "Edit Attributes". The main content area has a heading "Showing All Nodes" and a message "There are no items to display.".

## Training Node

- The labs require a node to be managed
- We allow for two different options
  - Bring your own Node
  - Use the CloudShare training environment

# Bring Your Own Node

- Use your own Virtual Machine (VM) or Server
- Required for the labs:
  - CentOS 6+
  - 512 MB RAM
  - 15 GB Disk
  - sudo or root level permissions

## CloudShare Node

- Register and login to CloudShare (see invite)
- Start Using This Environment

### Chef Fundamentals Webinar

Your dedicated hands-on environment is just a click away.

We believe you shouldn't have to waste time copying gigabytes of software, shipping machines, or traveling, just to get your IT into people's hands.

When you click the 'Start Using' button to your right, you'll have instant access worldwide to a full, enterprise-grade IT environment, available through your browser, mobile device, or dedicated client and VPN connection. You can connect your existing datacenter servers, change configurations, load new software, and provide this functionality to others. It's just like an on-premises data center, but with more flexibility and none of the headaches of backup, access, replication, logging, SLAs, and other annoyances - we take care of all of that for you.

CloudShare's SaaS platform is a quick and easy way to share copies of your complex IT environments, online. So you can collaborate with customers, partners, and colleagues - for Demos, Proofs-of-Concept, Training, or other enterprise applications.

Questions? Click [here](#) to email this environment's author or click [here](#) for CloudShare support.

Start Using  
This Environment



# CloudShare Node

Chef Fundamentals 2.x

CentOS 6

**CentOS 6.3 Server**

Description: OS: CentOS 6.3 x64  
Spec: 16 GB HD / 1 GB RAM  
OS: Linux  
State: Running

[More details ▶](#)

 View VM

 Reboot VM  Revert  Delete

# CloudShare Node

Chef Fundamentals 2.x

CentOS 6

**CentOS 6.3 Server**

Description: OS: CentOS 6.3 x64  
Spec: 16 GB HD / 1 GB RAM  
OS: Linux  
State: Running

[Hide details ▾](#)

**VM Details**

External Address: [uv01khywgnej7omyxst.vm.clld.nr](#)

Internal IP: 10.180.201.90

Total Memory: 1024 MB

Disk Size: 16 GB

CPU: 1

The machine was prepared in 21 seconds

Credentials ([show password](#))

Auto-login: root (local user)

Username: root

Password: \*\*\*\*\*

 Reboot VM  Revert  Delete

## Lab - Login

```
$ ssh chef@<EXTERNAL_ADDRESS>
```

```
The authenticity of host 'uvolqrwls0jdgs3blvt.vm.cld.sr
(69.195.232.110)' can't be established.
RSA key fingerprint is d9:95:a3:b9:02:27:e9:cd:
74:e4:a2:34:23:f5:a6:8b.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'uvolqrwls0jdgs3blvt.vm.cld.sr,
69.195.232.110' (RSA) to the list of known hosts.
chef@uvolqrwls0jdgs3blvt.vm.cld.sr's password:
Last login: Mon Jan  6 16:26:24 2014 from
host86-145-117-53.range86-145.btcentralplus.com
[chef@CentOS63 ~]$
```

## Checkpoint

- At this point you should have
  - One virtual machine (VM) or server that you'll use for the lab exercises
  - The IP address or public hostname
  - An application for establishing an ssh connection
  - 'sudo' or 'root' permissions on the VM

# Checkpoint



## "Bootstrap" the Target Instance

```
$ knife bootstrap <EXTERNAL_ADDRESS> --sudo -x chef -P chef -N "node1"
```

```
Bootstrapping Chef on uvo1qrwls0jdgs3blvt.vm.cld.sr
uvo1qrwls0jdgs3blvt.vm.cld.sr knife sudo password:
Enter your password:
...
...
uvo1qrwls0jdgs3blvt.vm.cld.sr Creating a new client identity for node1 using the
validator key.
uvo1qrwls0jdgs3blvt.vm.cld.sr resolving cookbooks for run list: []
uvo1qrwls0jdgs3blvt.vm.cld.sr Synchronizing Cookbooks:
uvo1qrwls0jdgs3blvt.vm.cld.sr Compiling Cookbooks...
uvo1qrwls0jdgs3blvt.vm.cld.sr [2014-01-28T11:03:14-05:00] WARN: Node node2 has an
empty run list.
uvo1qrwls0jdgs3blvt.vm.cld.sr Converging 0 resources
uvo1qrwls0jdgs3blvt.vm.cld.sr Chef Client finished, 0 resources updated
```

# knife bootstrap



Workstation



Node

# knife bootstrap

```
knife bootstrap HOSTNAME --sudo -x chef -P
```



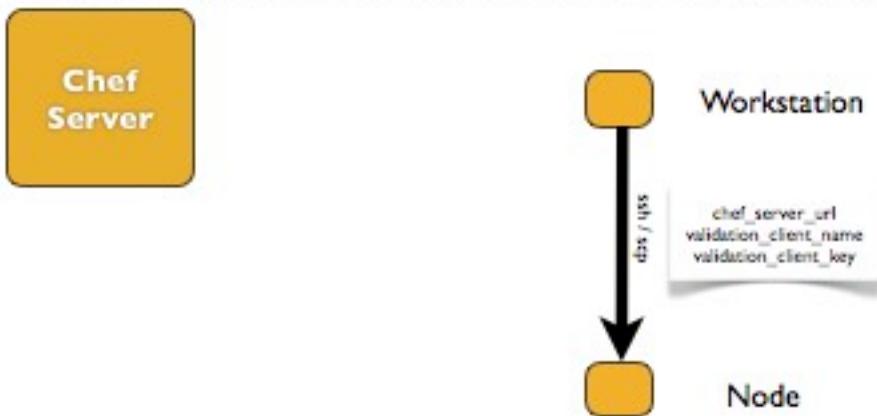
Workstation



Node

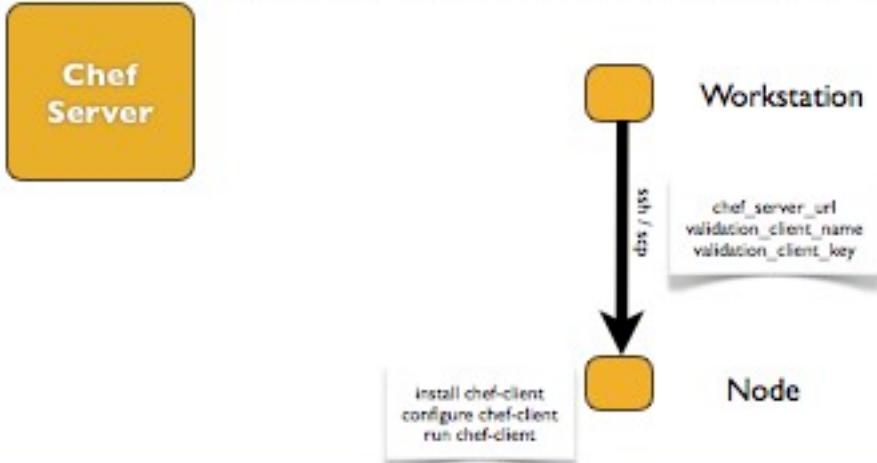
# knife bootstrap

```
knife bootstrap HOSTNAME --sudo -x chef -P
```



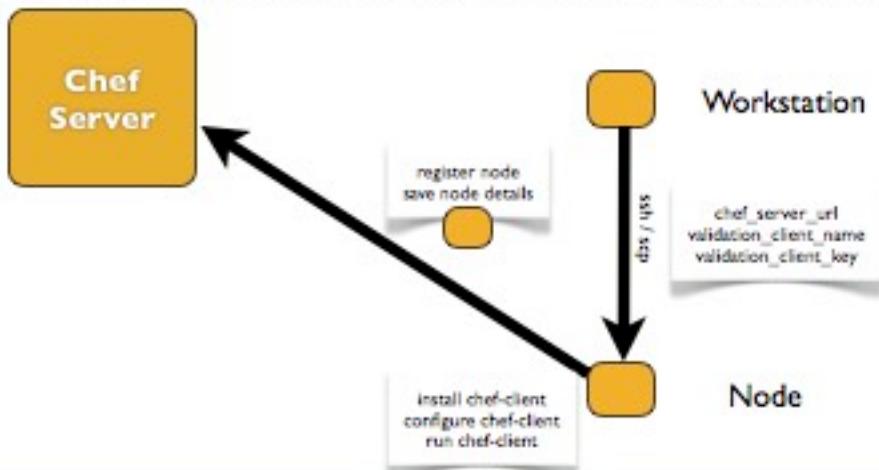
# knife bootstrap

```
knife bootstrap HOSTNAME --sudo -x chef -P
```



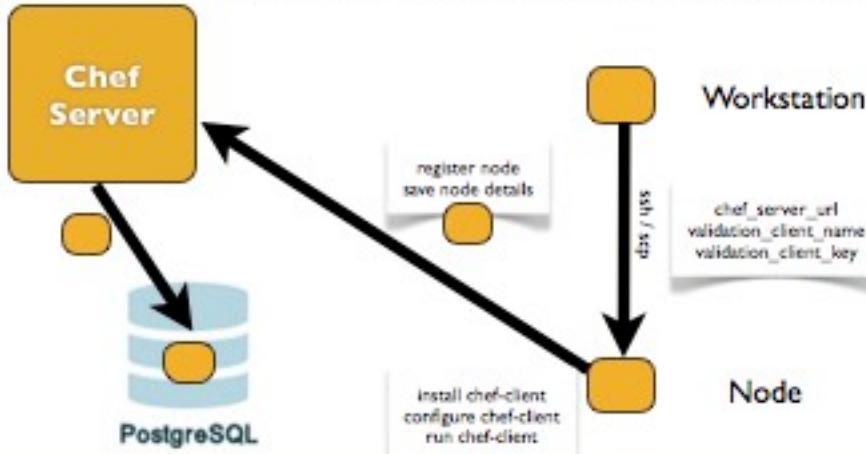
# knife bootstrap

```
knife bootstrap HOSTNAME --sudo -x chef -P
```



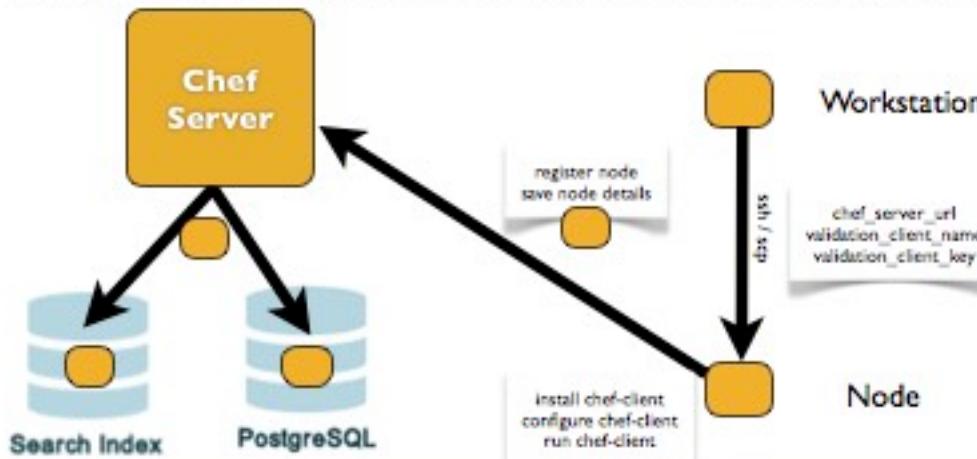
# knife bootstrap

```
knife bootstrap HOSTNAME --sudo -x chef -P
```



# knife bootstrap

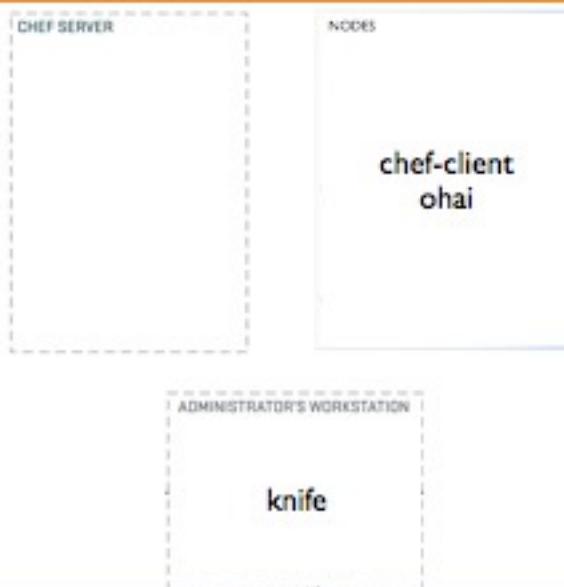
```
knife bootstrap HOSTNAME --sudo -x chef -P
```



## What just happened?

- Chef and all of its dependencies installed via an operating system-specific package ("omnibus installer")
- Installation includes
  - The Ruby language - used by Chef
  - knife - Command line tool for administrators
  - chef-client - Client application
  - ohai - System profiler
  - ...and more

# Workstation or Node?



## Verify Your Target Instance's Chef-Client is Configured Properly

```
$ ssh chef@<EXTERNAL_ADDRESS>

chef@node1:~$ ls /etc/chef
client.pem  client.rb  first-boot.json validation.pem

chef@node1:~$ which chef-client
/usr/bin/chef-client
```

## Examine /etc/chef/client.rb

```
chef@node1:~$ cat /etc/chef/client.rb
```

```
log_location      STDOUT
chef_server_url   "https://api.opscode.com/organizations/ORGNAME"
validation_client_name "ORGNAME-validator"
node_name         "node1"
```

## Exercise: Change the log level on your test node

```
chef@node1:~$ sudo vim /etc/chef/client.rb
```

```
log_level      :info
log_location    STDOUT
chef_server_url "https://api.opscode.com/organizations/ORGNAME"
validation_client_name "ORGNAME-validator"
node_name       "node1"
```

- Set the default log level for chef-client to **:info**
- More configuration options can be found on the docs site:  
[http://docs.chef.io/config\\_rb\\_client.html](http://docs.chef.io/config_rb_client.html)

# View Node on Chef Server

- Login to your Hosted Enterprise Chef



# View Node on Chef Server

- Click the 'Details' tab

The image shows the 'Nodes' section of the Chef Server interface. The 'Details' tab is highlighted with a red circle. The table lists one node: 'node1' (Platform: centos, FQDN: certos01.example.com, IP Address: 10.100.201.90). Below the table, there are tabs for 'Details', 'Attributes', and 'Permissions'. Under 'Details', it shows 'Last Check In: 3 Minutes Ago' (2014-01-01 11:38:31 UTC) and 'Uptime: 3 Hours' (Since 2014-01-01 08:07:08 UTC). The 'Attributes' tab shows no items, and the 'Permissions' tab also shows no items.

# View Node on Chef Server

- Click the 'Attributes' tab

The screenshot shows the Chef Server web interface. At the top, there's a navigation bar with tabs for Nodes, Reports, Policy, and Administration. Below that, a sidebar on the left lists '3 Nodes' and links for Delete, Manage Tags, Reset Key, Edit Run List, and Edit Attributes. The main content area shows a table of nodes with columns for Node Name, Platform, FQDN, and IP Address. One row is highlighted for 'node1'. Below this, a modal window is open for 'Node: node1'. It has three tabs: Details, Attributes (which is selected and highlighted with an orange oval), and Permissions. Under the Attributes tab, there's a section titled 'Attributes' with 'Expand All' and 'Collapse All' buttons. A list of attributes is shown, each preceded by a plus sign, indicating they can be expanded to view their values. The expanded attributes include:

- tags
- + languages
- + kernel
- os:linux
- os:version: 2.6.32-258.23.2.46.v86\_64
- hostname: centos03
- fqdn: centos03.example.com
- domain: example.com
- + network
- + counters
- ipaddress: 10.180.201.90
- macaddress: 00:50:56:00:00:90

## Node

- The node is registered with Chef Server
- The Chef Server displays information about the node
- This information comes from Ohai - we'll see Ohai later.....

# Checkpoint



## Review Questions

- Where is the chef-client configuration file?
- What is the command to run chef?
- What does a knife bootstrap do?

# Chef Resources and Recipes

Writing an Apache cookbook

92.1.8



## Lesson Objectives

- After completing the lesson, you will be able to
  - Describe in detail what a **cookbook** is
  - Create a new cookbook
  - Explain what a **recipe** is
  - Describe how to use the **package**, **service**, and **cookbook\_file** resources
  - **Upload a cookbook** to the Chef Server
  - Explain what a **run list** is, and how to set it for a node via **knife**
  - Explain the output of a chef-client run

## What is a cookbook?

- A cookbook is like a “package” for Chef recipes.
- It contains all the recipes, files, templates, libraries, etc. required to configure a portion of your infrastructure
- Typically they map 1:1 to a piece of software or functionality.

## The Problem and the Success Criteria

- **The Problem:** We need a web server configured to serve up our home page.
- **Success Criteria:** We can see the homepage in a web browser.

## Required steps

- Install Apache
  - Start the service, and make sure it will start when the machine boots
  - Write out the home page
- 
- Please note in this course we're teaching Chef primitives, not web server management
  - This is probably not the Apache HTTP server configuration you would use in production

## Exercise: Create a new Cookbook

```
$ knife cookbook create apache
```

```
** Creating cookbook apache
** Creating README for cookbook: apache
** Creating CHANGELOG for cookbook: apache
** Creating metadata for cookbook: apache
```

## Exercise: Explore the cookbook

```
$ ls -la cookbooks/apache
```

```
total 24
drwxr-xr-x 13 opscode opscode 442 Jan 24 21:25 .
drwxr-xr-x  5 opscode opscode 170 Jan 24 21:25 ..
-rw-r--r--  1 opscode opscode 412 Jan 24 21:25 CHANGELOG.md
-rw-r--r--  1 opscode opscode 1447 Jan 24 21:25 README.md
drwxr-xr-x  2 opscode opscode  68 Jan 24 21:25 attributes
drwxr-xr-x  2 opscode opscode  68 Jan 24 21:25 definitions
drwxr-xr-x  3 opscode opscode 102 Jan 24 21:25 files
drwxr-xr-x  2 opscode opscode  68 Jan 24 21:25 libraries
-rw-r--r--  1 opscode opscode 276 Jan 24 21:25 metadata.rb
drwxr-xr-x  2 opscode opscode  68 Jan 24 21:25 providers
drwxr-xr-x  3 opscode opscode 102 Jan 24 21:25 recipes
drwxr-xr-x  2 opscode opscode  68 Jan 24 21:25 resources
drwxr-xr-x  3 opscode opscode 102 Jan 24 21:25 templates
```

## Exercise: Edit the default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
#
# Cookbook Name:: apache
# Recipe:: default
#
# Copyright 2013, YOUR_COMPANY_NAME
#
# All rights reserved - Do Not Redistribute
#
```

# Chef Resources

```
package "haproxy" do
  action :install
end

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy.cfg.erb"
  owner "root"
  group "root"
  mode "0644"
  notifies :restart, "service[haproxy]"
end

service "haproxy" do
  supports :restart => true
  action [:enable, :start]
end
```

# Chef Resources

- Have a type

```
package "haproxy" do
  action :install
end

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy.cfg.erb"
  owner "root"
  group "root"
  mode "0644"
  notifies :restart, "service[haproxy]"
end

service "haproxy" do
  supports :restart => true
  action [:enable, :start]
end
```

# Chef Resources

- Have a type
- Have a name

```
package "haproxy" do
  action :install
end

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy.cfg.erb"
  owner "root"
  group "root"
  mode "0644"
  notifies :restart, "service[haproxy]"
end

service "haproxy" do
  supports :restart => true
  action [:enable, :start]
end
```

# Chef Resources

- Have a type
- Have a name
- Have **parameters**

```
package "haproxy" do
  action :install
end

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy.cfg.erb"
  owner "root"
  group "root"
  mode "0644"
  notifies :restart, "service[haproxy]"
end

service "haproxy" do
  supports :restart => true
  action [:enable, :start]
end
```

# Chef Resources

- Have a type
- Have a name
- Have parameters
- Take **action** to put the resource into the desired state

```
package "haproxy" do
  action :install
end

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy.cfg.erb"
  owner "root"
  group "root"
  mode "0644"
  notifies :restart, "service[haproxy]"
end

service "haproxy" do
  supports :restart => true
  action [:enable, :start]
end
```

# Chef Resources

- Have a type
- Have a name
- Have parameters
- Take **action** to put the resource into the desired state
- Can send **notifications** to other resources

```
package "haproxy" do
  action :install
end

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy.cfg.erb"
  owner "root"
  group "root"
  mode "0644"
  notifies :restart, "service[haproxy]"
end

service "haproxy" do
  supports :restart => true
  action [:enable, :start]
end
```

## Exercise: Add a package resource to install Apache to the default recipe

 OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
#  
# Cookbook Name:: apache  
# Recipe:: default  
#  
# Copyright 2013, YOUR_COMPANY_NAME  
#  
# All rights reserved - Do Not Redistribute  
#  
  
package "httpd" do  
  action :install  
end
```

SAVE FILE!

## So the resource we just wrote...

```
package "httpd" do  
  action :install  
end
```

## So the resource we just wrote...

- Is a **package** resource

```
package "httpd" do
  action :install
end
```

## So the resource we just wrote...

- Is a package resource
- Whose **name** is *httpd*

```
package "httpd" do
  action :install
end
```

## So the resource we just wrote...

- Is a package resource
- Whose name is *httpd*
- With an install **action**

```
package "httpd" do
  action :install
end
```

## Notice we didn't say how to install the package

- Resources are **declarative** - that means we say *what* we want to have happen, rather than *how*
- Resources take action through **Providers** - providers perform the *how*
- Chef uses the **platform** the node is running to determine the correct **provider** for a resource

# Package Resource

```
package "git"
```



```
  yum install git
```

```
  apt-get install git
```

```
  pacman sync git
```

```
  pkg_add -r git
```

**Providers are determined by node's platform**

**Exercise:** Add a service resource to ensure the service is started and enabled at boot



**OPEN IN EDITOR:** cookbooks/apache/recipes/default.rb

```
***  
# All rights reserved - Do Not Redistribute  
  
package "httpd" do  
  action :install  
end  
  
service "httpd" do  
  action [ :enable, :start ]  
end
```

**SAVE FILE!**

## So the resource we just wrote...

```
service "httpd" do
  action [ :enable, :start ]
end
```

## So the resource we just wrote...

- Is a **service** resource

```
service "httpd" do
  action [ :enable, :start ]
end
```

## So the resource we just wrote...

- Is a service resource
- Whose **name** is *httpd*

```
service "httpd" do
  action [ :enable, :start ]
end
```

## So the resource we just wrote...

- Is a service resource
- Whose **name** is *httpd*
- With two **actions**:
  - enable
  - start

```
service "httpd" do
  action [ :enable, :start ]
end
```

# Order Matters

- Resources are executed in order

1st

2nd

3rd

```
package "haproxy" do
  action :install
end

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy.cfg.erb"
  owner "root"
  group "root"
  mode "0644"
  notifies :restart, "service[haproxy]"
end

service "haproxy" do
  supports :restart => true
  action [:enable, :start]
end
```

Exercise: Add a `cookbook_file` resource to copy the home page in place



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

...

```
service "httpd" do
  action [ :enable, :start ]
end
```

```
cookbook_file "/var/www/html/index.html" do
  source "index.html"
  mode "0644"
end
```

SAVE FILE!

## So the resource we just wrote...

```
cookbook_file "/var/www/html/index.html" do
  source "index.html"
  mode "0644"
end
```

## So the resource we just wrote...

- Is a **cookbook\_file** resource

```
cookbook_file "/var/www/html/index.html" do
  source "index.html"
  mode "0644"
end
```

## So the resource we just wrote...

- Is a `cookbook_file` resource
- Whose **name** is:  
`/var/www/html/index.html`

```
cookbook_file "/var/www/html/index.html" do
  source "index.html"
  mode "0644"
end
```

## So the resource we just wrote...

- Is a `cookbook_file` resource
- Whose **name** is:  
`/var/www/html/index.html`
- With two **parameters**:
  - **source** of `index.html`
  - **mode** of “`0644`”

```
cookbook_file "/var/www/html/index.html" do
  source "index.html"
  mode "0644"
end
```

# Full contents of the apache recipe

```
#  
# Cookbook Name:: apache  
# Recipe:: default  
#  
# Copyright 2013, YOUR_COMPANY_NAME  
#  
# All rights reserved - Do Not Redistribute  
  
package "httpd" do  
  action :install  
end  
  
service "httpd" do  
  action [ :enable, :start ]  
end  
  
cookbook_file "/var/www/html/index.html" do  
  source "index.html"  
  mode "0644"  
end
```

## Exercise: Add index.html to your cookbook's files/default directory



OPEN IN EDITOR: cookbooks/apache/files/default/index.html

```
<html>  
<body>  
  <h1>Hello, world!</h1>  
</body>  
</html>
```

SAVE FILE!

## Exercise: Upload the cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache      [0.1.0]
Uploaded 1 cookbook.
```

## The Run List

- The Run List is the ordered set of recipes and roles that the Chef Client will execute on a node
  - Recipes are specified by “**recipe[name]**”
  - Roles are specified by “**role[name]**”

## Exercise: Add apache recipe to test node's run list

```
$ knife node run_list add node1 "recipe[apache]"
```

```
node1:
  run_list: recipe[apache]
```

## Exercise: Run Chef Client

```
chef@node1:~$ sudo chef-client
```

```
[2014-01-21T12:22:48-05:00] INFO: Picking chef instance to converge...
Starting Chef Client, version 11.8.2
[2014-01-21T12:22:48-05:00] INFO: *** chef 11.8.2 ***
[2014-01-21T12:22:48-05:00] INFO: Chef-client pid: 34346
[2014-01-21T12:22:48-05:00] INFO: Run list is [recipe[apache]]
[2014-01-21T12:22:48-05:00] INFO: Run list expands to [apache]
[2014-01-21T12:22:48-05:00] INFO: Starting chef run for node1
[2014-01-21T12:22:48-05:00] INFO: Receiving start handle...
[2014-01-21T12:22:48-05:00] INFO: Start handle complete.
[2014-01-21T12:22:48-05:00] INFO: HTTP Request measured 404 objects not found.
searching cookbooks for run
  list: ["apache"]
[2014-01-21T12:22:48-05:00] INFO: Loading cookbooks (apache)
synchronizing cookbooks...
[2014-01-21T12:22:48-05:00] INFO: matching updated
cookbooks/apache/recipes/default.rb in the cache.
[2014-01-21T12:22:48-05:00] INFO: Synchronizing updated
cookbooks/apache/recipes/default.rb in the cache.
[2014-01-21T12:22:48-05:00] INFO: matching updated
cookbooks/apache/webdata.rb in the cache.
[2014-01-21T12:22:48-05:00] INFO: Synchronizing updated
cookbooks/apache/webdata.rb in the cache.
  - apache
Compiling cookbooks...
Converging 3 resources
  recipes/apache/default
  httpd-2.2.15-0.18.20140121122248
httpd-2.2.15-0.18.20140121122248 from base repository
...
```

## Exercise: Verify that the home page works

- Open a web browser
- Type in the URL for your test node



## Congratulate yourself!

- You have just written your first Chef cookbook!
- (clap!)

## Reading the output of a chef-client run

```
Starting Chef Client, version 11.8.2
[2014-01-06T07:06:00-05:00] INFO: *** Chef 11.8.2 ***
[2014-01-06T07:06:00-05:00] INFO: Chef-client pid: 10781
[2014-01-06T07:06:01-05:00] INFO: Run List is [recipe[apache]]
[2014-01-06T07:06:01-05:00] INFO: Run List expands to [apache]
[2014-01-06T07:06:01-05:00] INFO: Starting Chef Run for node1
[2014-01-06T07:06:01-05:00] INFO: Running start handlers
[2014-01-06T07:06:01-05:00] INFO: Start handlers complete.
```

- The run list is shown
- The expanded Run List is the complete list, after nested roles are expanded

## Reading the output of a chef-client run

```
resolving cookbooks for run list: ["apache"]
[2014-01-06T07:06:02-05:00] INFO: Loading cookbooks [apache]
Synchronizing Cookbooks:
[2014-01-06T07:06:02-05:00] INFO: Storing updated cookbooks/apache/recipes/default.rb in the
cache.
[2014-01-06T07:06:02-05:00] INFO: Storing updated cookbooks/apache/metadata.rb in the cache.
  - apache
Compiling Cookbooks...
```

- Loads the cookbooks in the order specified by the run list
- Downloads any files that are missing from the server

## Reading the output of a chef-client run

```
Converging 3 resources
Recipe: apache::default
 * package[httpd] action install[2014-01-06T07:51:48-05:00] INFO: Processing
package[httpd] action install (apache::default line 9)
[2014-01-06T07:51:55-05:00] INFO: package[httpd] installing
httpd-2.2.15-29.el6.centos from base repository
 - install version 2.2.15-29.el6.centos of package httpd
```

- Checks to see if the package httpd is installed

## Reading the output of a chef-client run

```
* service[httpd] action enable[2014-01-06T07:52:06-05:00] INFO: Processing
service[httpd] action enable (apache::default line 13)
[2014-01-06T07:52:06-05:00] INFO: service[httpd] enabled
 - enable service service[httpd]

* service[httpd] action start[2014-01-06T07:52:06-05:00] INFO: Processing
service[httpd] action start (apache::default line 13)
[2014-01-06T07:52:07-05:00] INFO: service[httpd] started
 - start service service[httpd]
```

- Checks to see if httpd is already enabled to run at boot - it is, take no further action
- Checks to see if httpd is already started - it is, take no further action

# Idempotence

- Actions on resources in Chef are designed to be ***idempotent***
  - I.e. they can be applied multiple times but the end result is still the same - like multiplying by 1 in math!
  - Or in mathematical terms,  $F(F(x))=F(x) \forall x$
- Chef is a "desired state configuration" system - if a resource is already configured, no action is taken.  
This is called **convergence**

## Reading the output of a chef-client run

```
* cookbook_file[/var/www/html/index.html] action create[2014-01-06T07:52:07-05:00] INFO: Processing cookbook_file[/var/www/html/index.html] action create (apache::default line 17)
[2014-01-06T07:52:07-05:00] INFO: cookbook_file[/var/www/html/index.html] created file /var/www/html/index.html

- create new file /var/www/html/index.html[2014-01-06T07:52:07-05:00] INFO: cookbook_file[/var/www/html/index.html] updated file contents /var/www/html/index.html

- update content in file /var/www/html/index.html from none to 03fbfd
  --- /var/www/html/index.html      2014-01-06 07:52:07.285214202 -0500
  +++ /tmp/.index.html20140106-10796-1lxoknbq 2014-01-06 07:52:07.868365963 -0500
@@ -1,+1,6 @@
+<html>
+<body>
+ <h1>Hello, world!</h1>
+</body>
+</html>[2014-01-06T07:52:07-05:00] INFO: cookbook_file[/var/www/html/index.html] mode changed to 644

- change mode from '' to '644'
- restore selinux security context
```

- Checks for an index.html file
- There is already one in place, backup the file
- Set permissions on the file
- A diff of the written file is shown with the modified lines called out

## Reading the output of a chef-client run

```
[2014-01-06T07:52:08-05:00] INFO: Chef Run complete in 23.477837576 seconds
[2014-01-06T07:52:08-05:00] INFO: Running report handlers
[2014-01-06T07:52:08-05:00] INFO: Report handlers complete
Chef Client finished, 4 resources updated
[2014-01-06T07:52:08-05:00] INFO: Sending resource update report (run-id: 952a8431-0994-468e-836c-0f7de7aa656e)
```

- Notice that a complete Chef-Run displays:
  - The time the client took to complete convergence
  - Status of report and exception handlers

## Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
...
resolving cookbooks for run list: ['apache']
[2014-01-06T07:57:13-05:00] INFO: Loading cookbooks [apache]
Synchronizing Cookbooks
 - apache
Compiling Cookbooks...
Converging 3 resources
Recipe: apache::default
 * package[httpd] action install[2014-01-06T07:57:13-05:00] INFO: Processing package[httpd] action install (apache::default line 8)
 (up to date)
 * service[httpd] action enable[2014-01-06T07:57:17-05:00] INFO: Processing service[httpd] action enable (apache::default line 13)
 (up to date)
 * service[httpd] action start[2014-01-06T07:57:17-05:00] INFO: Processing service[httpd] action start (apache::default line 13)
 (up to date)
 * cookbook_file[/var/www/html/index.html] action create[2014-01-06T07:57:17-05:00] INFO: Processing cookbook_file[/var/www/html/index.html] action create (apache::default line 17)
 (up to date)
[2014-01-06T07:57:17-05:00] INFO: Chef Run complete in 4.797129533 seconds
[2014-01-06T07:57:17-05:00] INFO: Removing cookbooks/apache/Ellen/default/index.html from the cache; it is no longer needed by chef-client.
[2014-01-06T07:57:17-05:00] INFO: Running report handlers
[2014-01-06T07:57:17-05:00] INFO: Report handlers complete
Chef Client finished, 0 resources updated
[2014-01-06T07:57:17-05:00] INFO: Sending resource update report (run-id: 1c205d2c-7971-43ab-a994-1ba323b9aeb8)
```

## Recap - So resources are grouped into recipes

```
recipe[apache::default] do
  package "httpd" do
    action :install
  end

  service "httpd" do
    action [ :enable, :start ]
  end

  cookbook_file "/var/www/html/index.html" do
    source "index.html"
    mode "0644"
  end
```

## Cookbooks contain recipes and supporting files

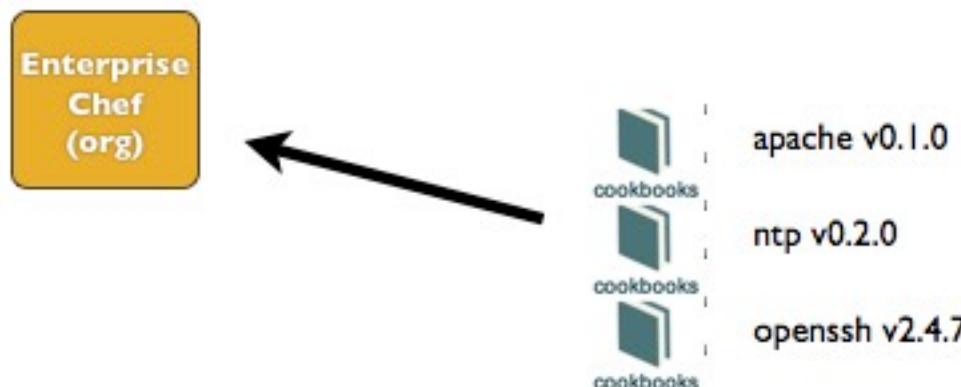
```
& pwd
/Users/you/chef-repo
$ tree
...
├── cookbooks
│   └── apache
│       ├── recipes
│       │   ├── default.rb
│       │   └── server.rb
│       └── files
│           └── default
│               └── index.html
└── metadata.rb
└── attributes
    └── default
...
...
  
```



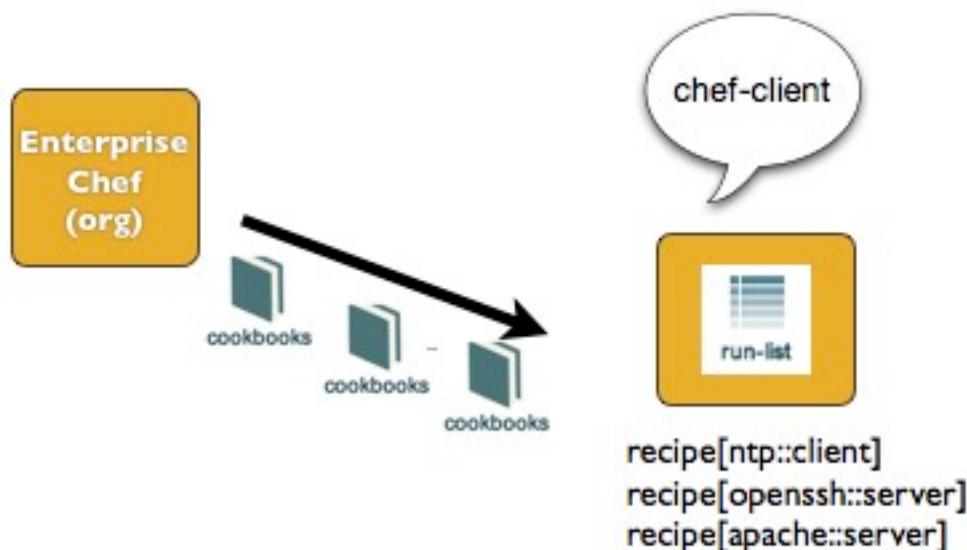
Recipes

Supporting File

## Cookbooks are installed as artifacts on Chef Server



## Nodes have run\_lists made of recipes



# Questions

- What makes up a cookbook?
- How do you create a new cookbook?
- What is a recipe?
- What is a resource?
- How do you upload a cookbook to the Chef Server?
- What is a run list?

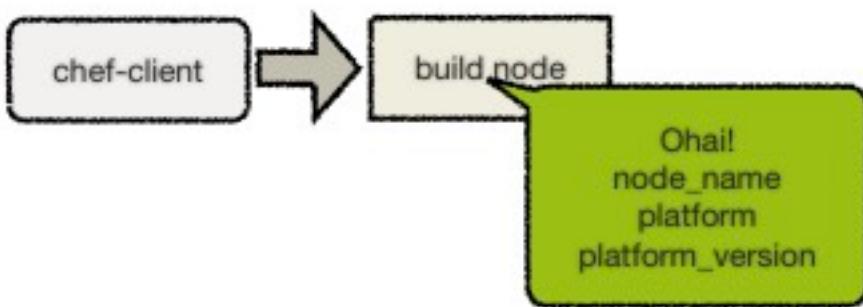
## Dissecting Your First chef-client Run

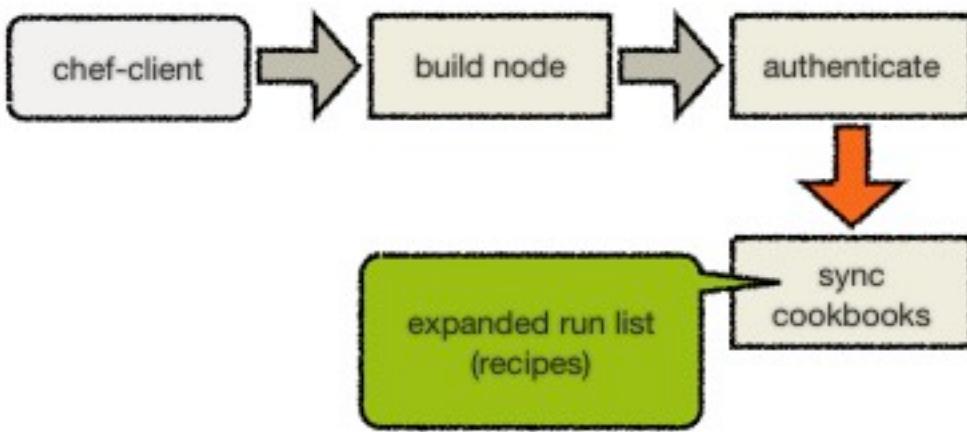
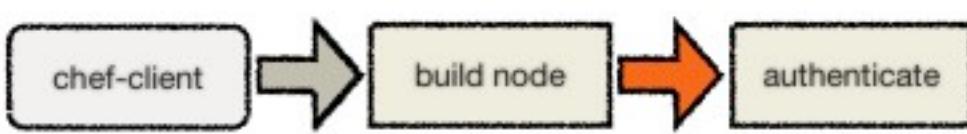
The Anatomy of a Chef run

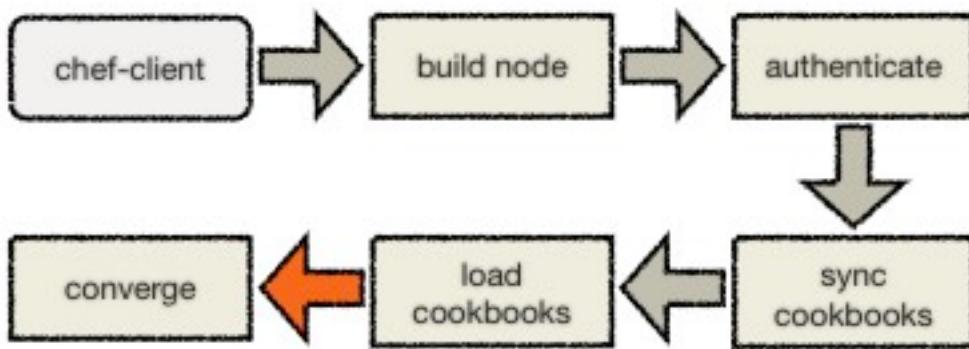
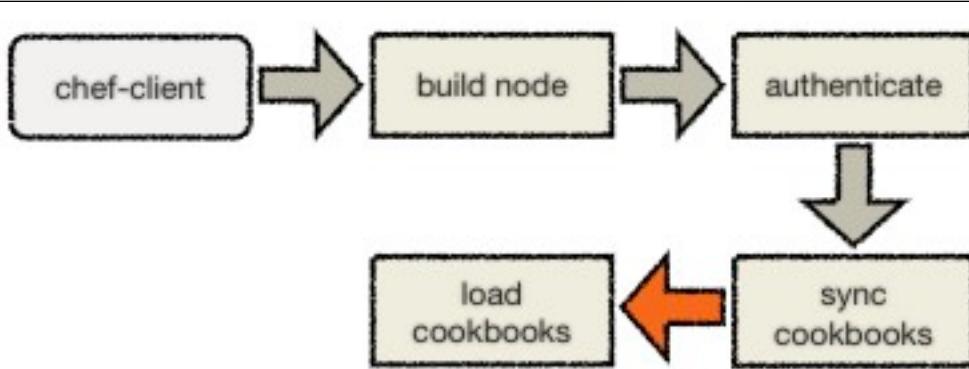
# Lesson Objectives

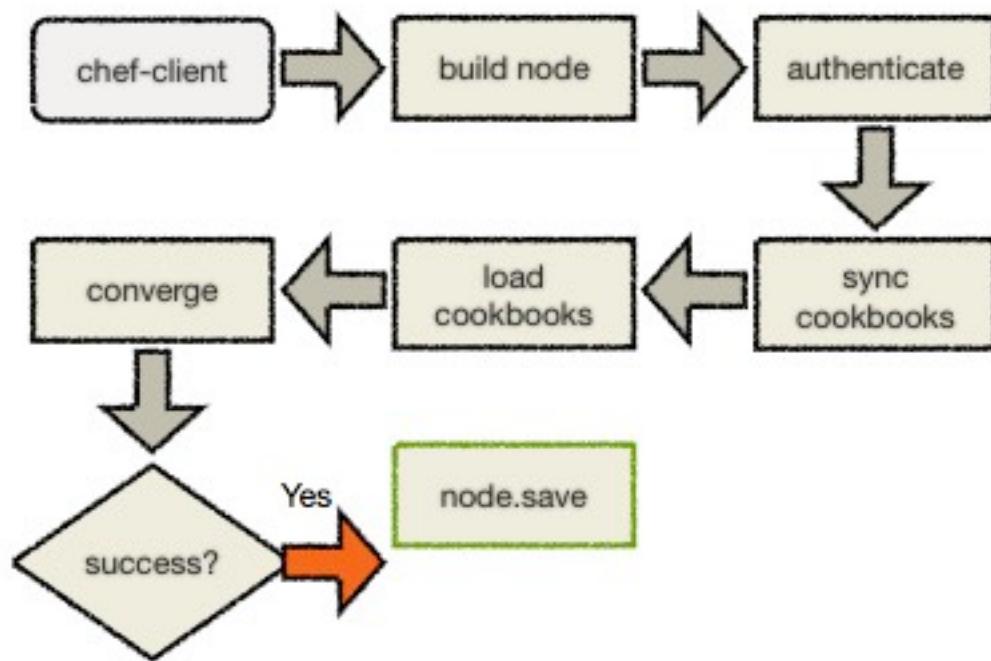
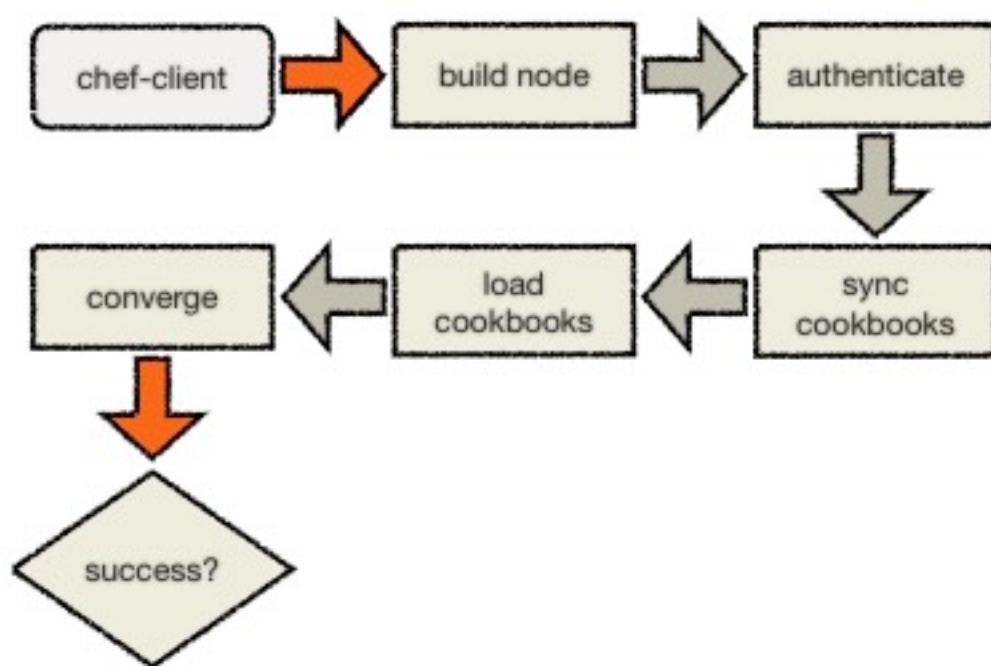
- After completing the lesson, you will be able to
- List all the steps taken by a chef-client during a run
- Explain the basic security model of Chef
- Explain the concepts of the Resource Collection

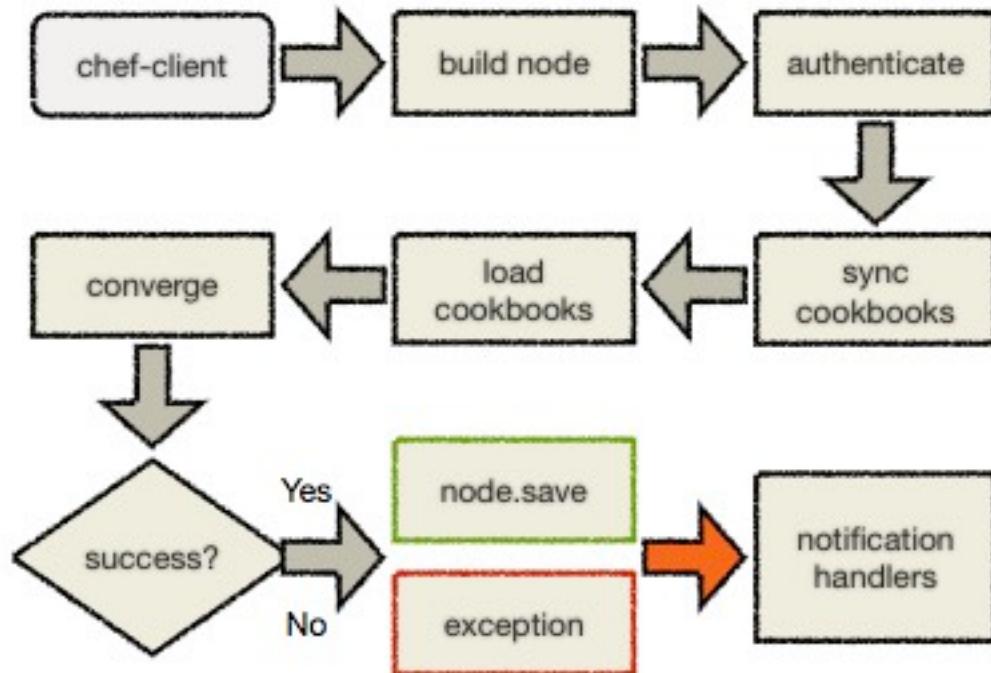
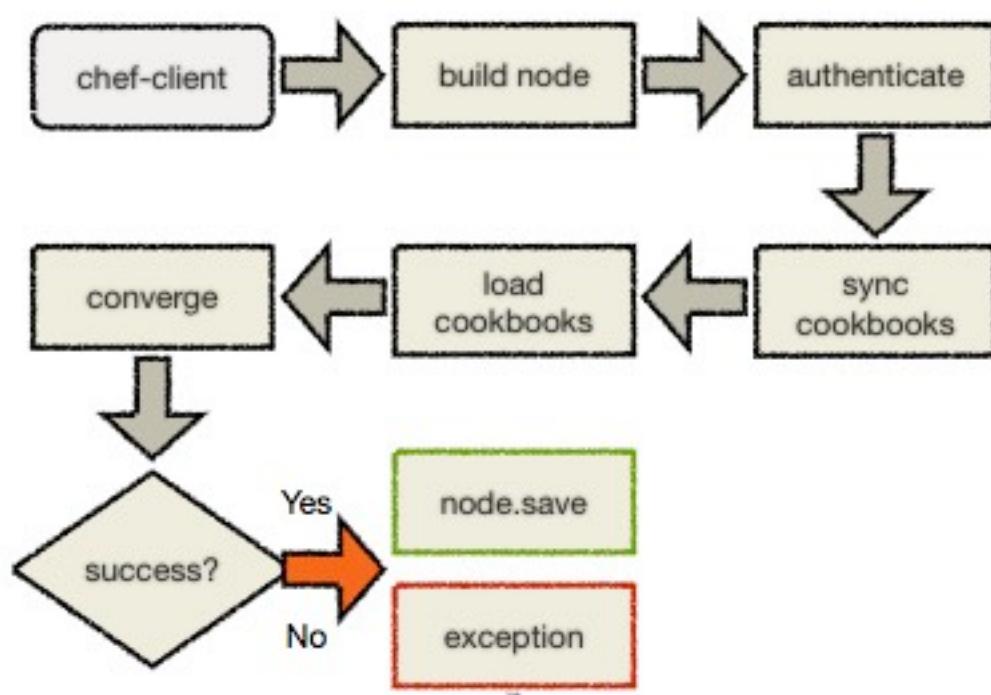
chef-client





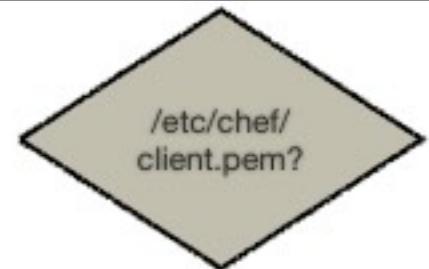


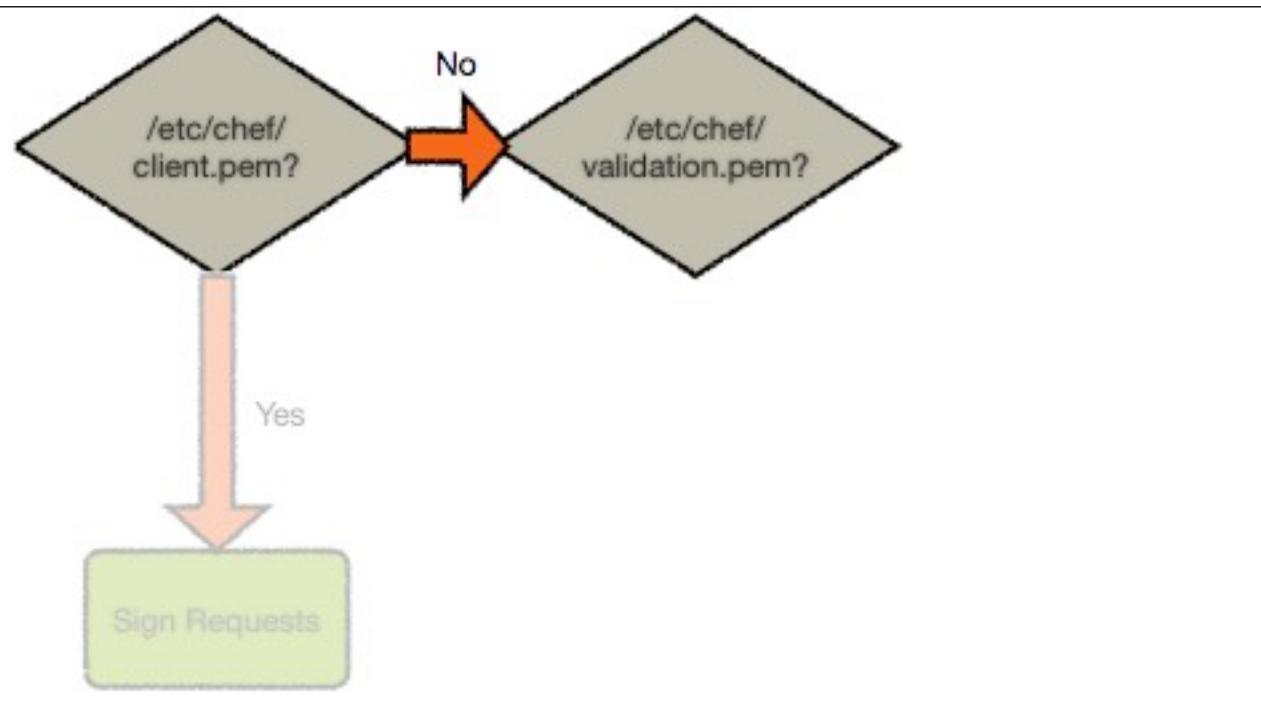
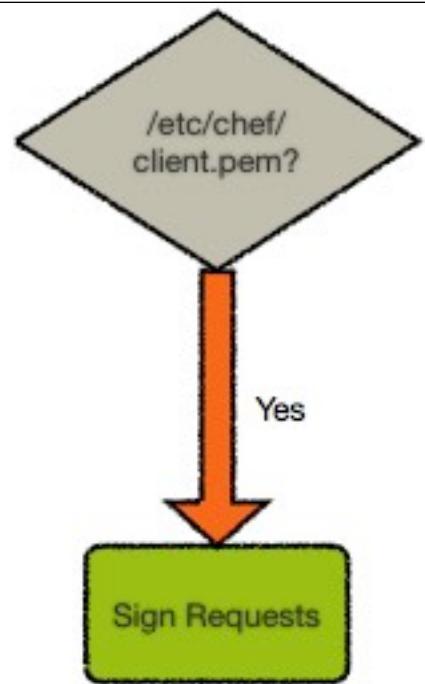


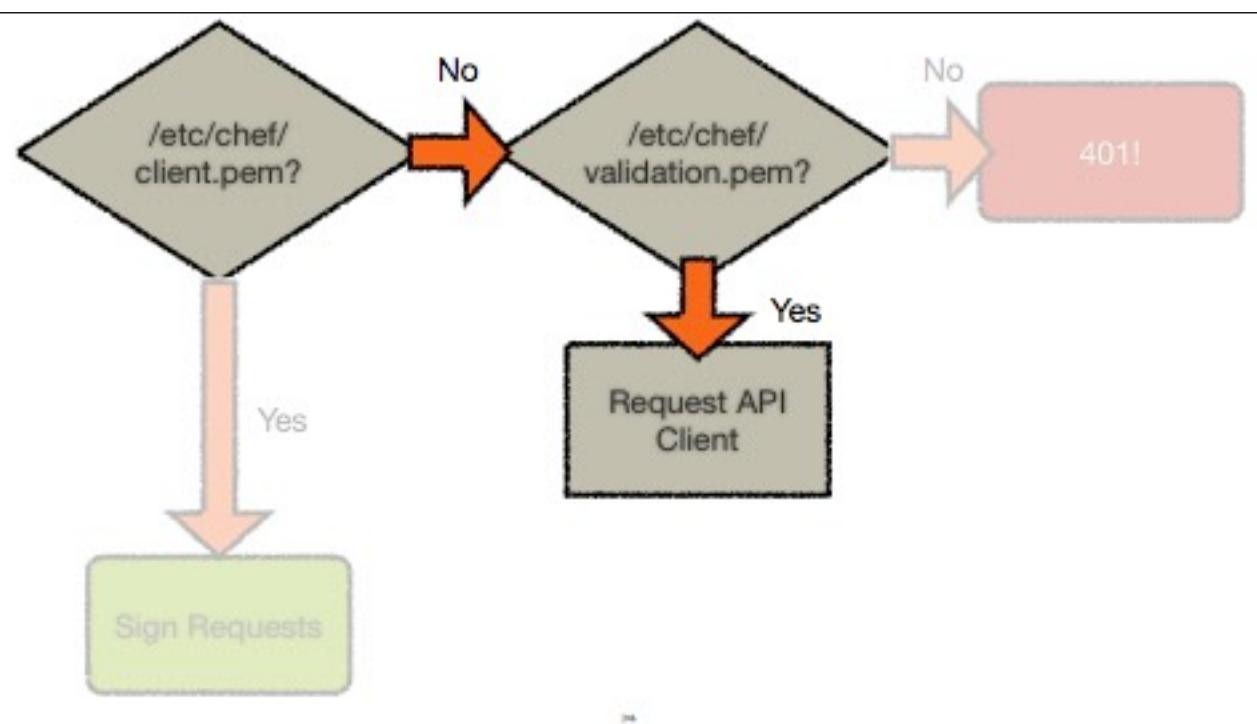
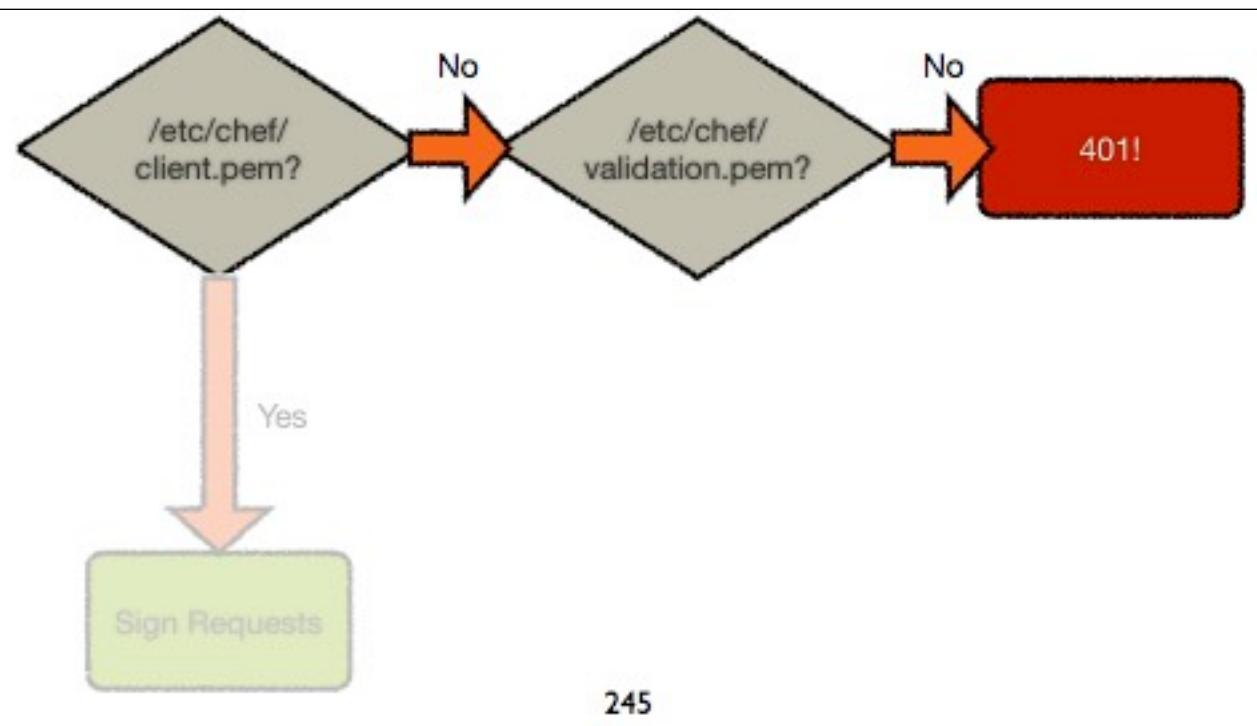


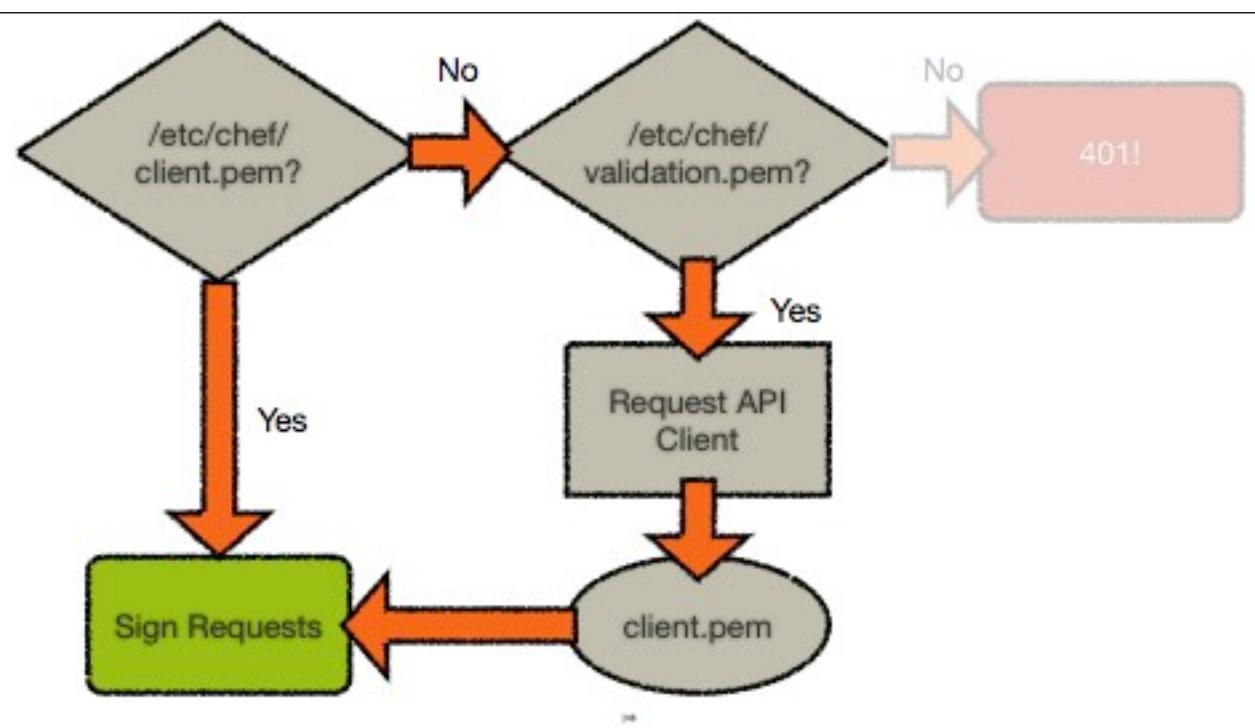
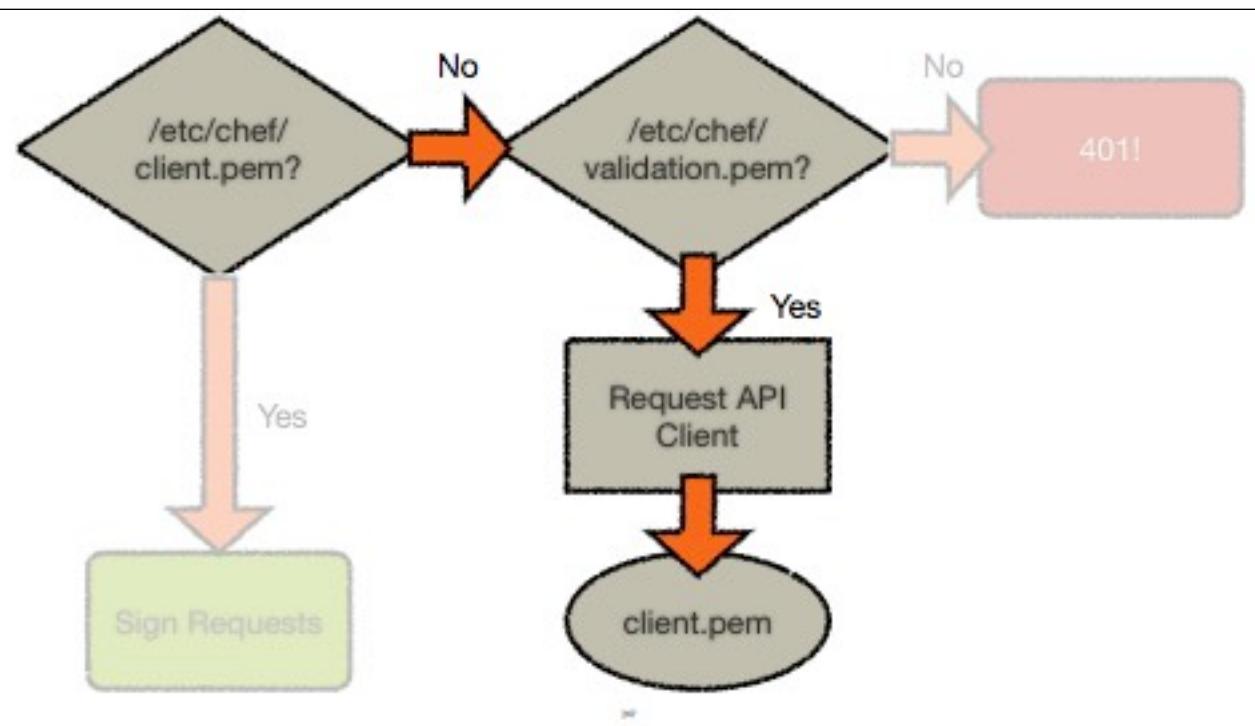
# Private Keys

- Chef Server requires keys to authenticate.
  - `client.pem` - private key for API client
  - `validation.pem` - private key for ORGNAME-validator
- Next, let's see how those are used...









# About the resource collection

v2.1.6

## Multiphase Execution - Compile Phase

- During the compile phase, Chef
  1. Loads all cookbooks from the run list
  2. Reads each recipe to build the resource collection

## Multiphase Execution - Execute Phase

- During the execute phase chef takes the resource collection and for each resource it will
  - Check if the resource is in the required state
    - If 'yes' - do nothing
    - If 'no' - bring resource in line with required state
  - Move on to next resource

## Resource Collection Phase 1 - Compile Phase

### Recipe

```
package "httpd" do
  action :install
end

cookbook_file "/var/www/html/index.html" do
  source "index.html"
  mode "0644"
end

service "httpd" do
  action [ :enable, :start ]
end
```

### Resource Collection

```
resource_collection = [
  package["httpd"],
  cookbook_file["/var/www/html/index.html"],
  service ["httpd"]
]
```

## Resource Collection Phase 2 - Execute Phase

### Recipe

```
package "httpd" do
  action :install
end

cookbook_file "/var/www/html/index.html" do
  source "index.html"
  mode "0644"
end

service "httpd" do
  action [ :enable, :start ]
end
```

### Resource Collection

```
resource_collection = [
  package["httpd"],
  cookbook_file["/var/www/html/index.html"],
  service ["httpd"]
]
```

### Execution



## Recipe order is important!

- Recipes are executed in the order they appear in the run list

Run List: recipe[ntp::client], recipe[openssh::server], recipe[apache::server]

- These recipes are invoked in the following order

1. recipe[ntp::client]
2. recipe[openssh::server]
3. recipe[apache::server]

## Resource Collection - Multiple Recipes

### 1.recipe[ntp::client]

```
package "ntp" do
  action :install
end

template "/etc/ntp.conf" do
  source "ntp.conf.erb"
  owner "root"
  mode "0644"
end

service "ntp" do
  action :start
end
```

### Resource Collection

```
resource_collection [
  package[ntp],
  template[/etc/ntp.conf],
  service[ntp]
```

## Resource Collection - Multiple Recipes

### 2.recipe[openssh::client]

```
package "openssh" do
  action :install
end

template "/etc/sshd/sshd_config" do
  source "sshd_config.erb"
  owner "root"
  mode "0644"
end

service "openssh" do
  action :start
end
```

### Resource Collection

```
resource_collection [
  package[ntp],
  template[/etc/ntp.conf],
  service[ntp],
  package[openssh],
  template[/etc/sshd/sshd_config],
  service[openssh]
```

## Resource Collection - Multiple Recipes

```
3.recipe[httpd::server]
```

```
  package "httpd" do
    action :install
  end

  service "httpd" do
    action [ :enable, :start ]
  end

  cookbook_file "/var/www/html/index.html" do
    source "index.html"
    mode "0644"
  end
```

### Resource Collection

```
resource_collection [
  package[ntp],
  template[/etc/ntp.conf],
  service[ntp],
  package[openssh],
  template[/etc/sshd/sshd_config],
  service[openssh],
  package[httpd],
  service[httpd],
  cookbook_file[/var/www/html/index.html]
]
```

## The final resource collection

- So the resources are invoked in the following order during the execute phase

```
package[ntp]
template[/etc/ntp.conf]
service[ntp]
package[openssh]
template[/etc/sshd/sshd_config]
service[openssh]
package[httpd]
service[httpd]
cookbook_file[/var/www/html/index.html]
```

# Multiphase Execution

- Plain ruby is executed in the compile phase
- Chef DSL is executed in the execute phase

Recipe

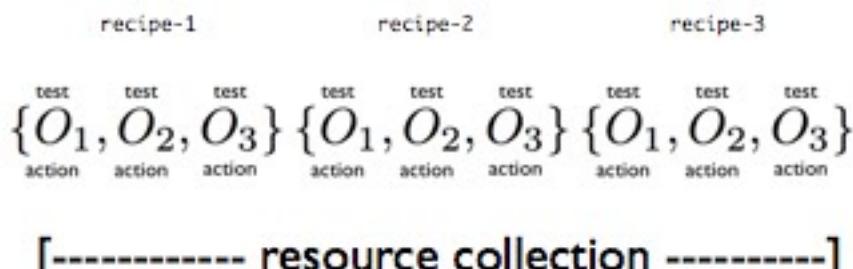
```
#w[sites-available sites-enabled mods-available].each do |dir|
  directory "/var/www/#{dir}" do
    action :create
    mode  '0755'
    owner 'root'
    group node['apache']['root_group']
  end
end
```

Resource Collection

```
resource_collection [
  directory["/var/www/sites-available"],
  directory["/var/www/sites-enabled"],
  directory["/var/www/mods-available"],
]
```

## Remember - Resource order is important!

- Resources are invoked in the order they appear in the recipe



## Review Questions

- What are the steps in a Chef Client run?
- How does a new machine get a private key with which to authenticate requests?
- If you have the right credentials in place, why else might you not be able to authenticate?
- In which phase of a chef-client run do plain Ruby statements get evaluated?

## Introducing the Node object

Attributes & Search

## Lesson Objectives

- After completing the lesson, you will be able to
  - Explain what the Node object represents in Chef
  - List the Nodes in an organization
  - Show details about a Node
  - Describe what Node Attributes are
  - Retrieve a node attribute directly, and via search

## What is the Node object

- A node is any physical, virtual, or cloud machines that is configured to be maintained by a Chef
- The 'node object' is the representation of that physical node within Chef (e.g. in JSON)
- When you are writing Recipes, the Node object is always available to you.

## Node

- The node is registered with Chef Server
- The Chef Server displays information about the node
- This information comes from Ohai

## View Node on Chef Server

The screenshot shows the Chef Server interface for viewing nodes. The top navigation bar includes links for Nodes, Reports, Policy, and Administration. On the left, a sidebar provides options for managing nodes: Delete, Manage Tags, Reset Key, Edit Run List, and Edit Attributes. The main content area displays a table titled "Showing All Nodes" with columns for Node Name, Platform, FQDN, and IP Address. A single row is selected for the node named "node1". Below the table, a modal window titled "Node: node1" is open, showing tabs for Details, Attributes, and Permissions. Under the Details tab, it indicates the last check-in was 3 Minutes Ago (2014-01-28 11:58:31 UTC) and the uptime is 8 Hours (Since 2014-01-28 07:08:18 UTC). At the bottom of the modal, there is a "Tags" section with an "Add" button and a message stating "There are no items to display."

# View Node on Chef Server

The screenshot shows the Chef web interface. In the top navigation bar, there are tabs for Nodes, Reports, Policy, and Administration. The Nodes tab is selected. On the left sidebar, under the 'Nodes' section, there are links for Delete, Manage Tags, Reset Key, Edit Run List, and Edit Attributes. The main content area displays a table titled 'Showing All Nodes' with columns for Node Name, Platform, FQDN, and IP Address. A single row is highlighted for 'node1'. Below the table, a modal window is open for 'Node: node1'. It has tabs for Details, Attributes, and Permissions, with the Attributes tab selected. Under the Attributes heading, there are two buttons: 'Expand All' (highlighted in orange) and 'Collapse All'. The expanded list shows various attributes: tags, languages, name, os, os\_info, os\_version, hostname, fqdn, domain, network, counters, ipaddress, and macaddress.

## Exercise: List nodes

```
$ knife node list
```

```
node1
```

## Exercise: List clients

```
$ knife client list
```

```
ORGNAME-validator  
node1
```

### Each node must have a unique name

- Every node must have a unique name within an organization
- Chef defaults to the *Fully Qualified Domain Name* of the server, i.e. in the format `server.domain.com`
- We overrode it to "node1" to make typing easier

## Exercise: Show node details

```
$ knife node show node1
```

```
Node Name: node1
Environment: _default
FQDN: centos63.example.com
IP: 10.160.201.90
Run List: recipe[apache]
Roles:
Recipes: apache
Platform: centos 6.4
Tags:
```

## What is the Node object

- Nodes are made up of Attributes
- Many are discovered **automatically** (platform, ip address, number of CPUs)
- Many other objects in Chef can also add Node attributes (Cookbooks, Roles and Environments, Recipes, Attribute Files)
- Nodes are stored and indexed on the Chef Server

# Ohai

```
"languages": {  
  "ruby": {  
  },  
  "perl": {  
    "version": "5.14.2",  
    "archname": "x86_64-  
linux-gnu-thread-multi"  
  },  
  "python": {  
    "version": "2.6.6",  
    "builddate": "Jul 10  
2013, 22:48:45"  
  },  
  "perl": {  
    "version": "version",  
    "archname": "x86_64-  
linux-thread-multi"  
  },  
  "lua": {  
    "version": "5.1.4"  
  },  
  "kernel": {  
    "name": "Linux",  
    "release": "3.2.0-32-virtual",  
    "version": "#1 SMP Wed Oct 16  
18:37:12 UTC 2013",  
    "machine": "x86_64",  
    "modules": {  
      "isofs": {  
        "size": "70066",  
        "refcount": "2"  
      },  
      "des_generic": {  
        "size": "16604",  
        "refcount": "0"  
      }  
    },  
    "os": "GNU/Linux"  
  },  
  "os": "linux",  
  "os_version":  
  "2.6.32-358.23.2.el6.x86_64",  
  "ohai_time": 1389105685.7735305,  
  "network": {  
    "interfaces": {  
      "lo": {  
        "mtu": "16436",  
        "flags": [  
          "LOOPBACK", "UP", "LOWER_UP"  
        ],  
        "encapsulation": "Loopback",  
        "addresses": {  
          "127.0.0.1": {  
            "family": "inet",  
            "netmask": "255.0.0.0",  
            "scope": "Node"  
          },  
          "::1": {  
            "family": "inet6",  
            "scope": "Node"  
          }  
        },  
        "eth0": {  
          "type": "eth",  
          "number": "0"  
        }  
      }  
    }  
  }  
}
```

## Exercise: Run Ohai on node

```
chef@node1:~$ ohai | less
```

```
{  
  "languages": {  
    "ruby": {  
    },  
    "java": {  
      "version": "1.6.0_24",  
      "runtime": {  
        "name": "OpenJDK Runtime Environment (IcedTea6 1.11.13)",  
        "build": "rhel-1.65.1.11.13.el6_4-x86_64"  
      },  
      "hotspot": {  
        "name": "OpenJDK 64-Bit Server VM",  
        "build": "20.0-b12, mixed mode"  
      }  
    }  
  }  
}
```

## Exercise: Show all the node attributes

```
$ knife node show node1 -l
```

```
Node Name: node1
Environment: _default
FQDN: centos63.example.com
IP: 10.160.201.90
Run List: recipe[apache]
Roles:
Recipes: apache
Platform: centos 6.4
Tags:
Attributes:
tags:

Default Attributes:

Override Attributes:

Automatic Attributes (Chai Data):
block_device:
  dm-0:
    removable: 0
    size: 28393472
```

## Exercise: Show the raw node object

```
$ knife node show node1 -Fj
```

```
{
  "name": "node1",
  "chef_environment": "_default",
  "run_list": ["recipe[apache]"],
  "normal": {"tags": []}
}
```

## Exercise: Show only the fqdn attribute

```
$ knife node show node1 -a fqdn
```

```
node1:  
  fqdn: centos63.example.com
```

## Exercise: Use search to find the same data

```
$ knife search node "*:*" -a fqdn
```

```
1 items found  
  
node1:  
  fqdn: centos63.example.com
```

## Review Questions

- What is the Node object?
- What is a Node Attribute?
- How do you display all the attributes of a Node?
- Can you search for the **cpu** attribute of your node?

## Setting Node Attributes

Setting attributes in recipes and attribute files

## Lesson Objectives

- After completing the lesson, you will be able to
  - Describe where and how attributes are set
  - Explain the attribute merge order and precedence rules
  - Declare an attribute with a recipe and set its value

## What are Attributes?

- Attributes represent information about your node
- The information can be auto-detected from the node (e.g. # of CPUs, amount of RAM) & populated by Ohai
- You can also set attributes on your node using cookbook recipes & attribute files, roles, environments, etc
- Attributes keep the program code separate from data
- All attributes are set on the "node object", and are indexed for search on the server

# Attribute Sources

- Attributes can be set at various levels (in increasing order of precedence)
- Automatically on the node itself (by Ohai)
- In roles
- In environments
- In cookbook recipes
- In cookbook attribute files

## Ohai - set automatically

```
"languages": {  
  "ruby": {  
  },  
  "perl": {  
    "version": "5.14.2",  
    "archname": "x86_64-  
linux-gnu-thread-multi"  
  },  
  "python": {  
    "version": "2.6.6",  
    "builddate": "Jul 10  
2013, 22:48:45"  
  },  
  "perl": {  
    "version": "5.10.1",  
    "archname": "x86_64-  
linux-thread-multi"  
  }  
},  
  
"kernel": {  
  "name": "Linux",  
  "release": "3.2.0-32-virtual",  
  "version": "#1 SMP Wed Oct 16  
18:37:12 UTC 2013",  
  "machine": "x86_64",  
  "modules": {  
    "isofs": {  
      "size": "70066",  
      "refcount": "2"  
    },  
    "des_generic": {  
      "size": "16604",  
      "refcount": "0"  
    }  
  },  
  "os": "GNU/Linux"  
},  
"os": "linux",  
"os_version":  
"2.6.32-358.23.2.el6.x86_64",  
"ohai_time": 1389105685.7735305,
```

```
"network": {  
  "interfaces": {  
    "lo": {  
      "mtu": "16436",  
      "flags": [  
        "LOOPBACK", "UP", "LOWER_UP"  
      ],  
      "encapsulation": "Loopback",  
      "addresses": {  
        "127.0.0.1": {  
          "family": "inet",  
          "netmask": "255.0.0.0",  
          "scope": "Node"  
        },  
        "::1": {  
          "family": "inet6",  
          "scope": "Node"  
        }  
      },  
      "eth0": {  
        "type": "eth",  
        "number": "0",  
        "mac": "00:0c:29:1d:4f:00",  
        "mtu": "1500",  
        "flags": [  
          "BROADCAST", "MULTICAST", "UP", "LOWER_UP"  
        ],  
        "encapsulation": "Ethernet",  
        "addresses": {  
          "192.168.1.10": {  
            "family": "inet",  
            "netmask": "255.255.255.0",  
            "scope": "Global",  
            "link_layer": "00:0c:29:1d:4f:00"  
          }  
        }  
      }  
    }  
  }  
}
```

## Setting attributes in attribute files

- Attributes can be set in the cookbook's attributes file
  - ./cookbooks/<cookbook>/attributes/default.rb
- Format is

The diagram illustrates the structure of an attribute assignment in an attributes file. It shows three orange boxes with arrows pointing downwards to the corresponding parts of the code:

- The first box is labeled "precedence" and points to the word "default".
- The second box is labeled "attribute name" and points to the string "apache".
- The third box is labeled "attribute value" and points to the string "/etc/apache2".

```
default["apache"]["dir"] = "/etc/apache2"
```

- We'll look at precedence later....

## Setting Attributes in Recipes

- They can also be set directly in Recipes
- Precede attribute name with 'node.' as follows

The diagram illustrates the structure of an attribute assignment in a recipe. It shows three orange boxes with arrows pointing downwards to the corresponding parts of the code:

- The first box is labeled "precedence" and points to the prefix "node.default".
- The second box is labeled "attribute name" and points to the string "apache".
- The third box is labeled "attribute value" and points to the string "/etc/apache2".

A fourth orange box with an upward arrow points to the "node.default" prefix, containing the text "declares it a node attribute".

```
node.default["apache"]["dir"] = "/etc/apache2"
```

## The Problem and the Success Criteria

- **The Problem:** We have defined our 'index.html' homepage in the recipe, but we may want to change that without altering the recipe
- **Success Criteria:** We can change the homepage by changing an attribute value

## Exercise: Set attribute in attributes file

 OPEN IN EDITOR: cookbooks/apache/attributes/default.rb

```
default["apache"]["indexfile"] = "index1.html"
```

SAVE FILE!

- Set an attribute to a specific value in the attributes file

## Exercise: Add index1.html to cookbook's files/default directory

 OPEN IN EDITOR: cookbooks/apache/files/default/index1.html

```
<html>
  <body>
    <h1>Hello, world!</h1>
    <h2>This is index1.html</h2>
    <p>We configured this in the attributes file</p>
  </body>
</html>
```

SAVE FILE!

## Exercise: Set attribute in a recipe

 OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
service "httpd" do
  action [:enable, :start]
end

cookbook_file "/var/www/html/index.html" do
  source node["apache"]["indexfile"]
  mode "0644"
end
```

SAVE FILE!

## Exercise: Upload the cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [0.1.0]
Uploaded 1 cookbook.
```

## Exercise: Re-run Chef Client

```
chef@node1:~$ sudo chef-client
```

```
Starting Chef Client, version 11.8.2
resolving cookbooks for run list: ("apache")
...
Recipe: apache::default
 * package(httdp) action install (up to date)
 * service(httdp) action enable (up to date)
 * service(httdp) action start (up to date)
 * cookbook_file[/var/www/html/index.html] action create
 - update content in file /var/www/html/index.html from 43fbfd to 478a29
   --- /var/www/html/index.html 2014-01-07 04:28:34.994963341 -0500
   +++ /tmp/.index.html20140107-12128-1edca1 2014-01-07 04:31:24.049155232 -0500
   #! <1> +1,8 @@
   <html>
   <body>
   + <body>
   + <h1>Hello, world!</h1>
   + </body>
   +   <h2>This is index.html</h2>
   +   <p>We configured this in the attributes file</p>
   + </body>
   </html>
 - restore selinux security context

Chef Client finished, 1 resources updated
```

## Exercise: Verify new homepage works

- Open a web browser
- The homepage takes the attribute file value



Hello, world!

This is index1.html

We configured this in the attributes file

## Exercise: Set attribute in the recipe

OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
service "apache2" do
  action [:enable, :start]
end

node.default["apache"]["indexfile"] = "index2.html"
cookbook_file "/var/www/html/index.html" do
  source node["apache"]["indexfile"]
  mode "0644"
end
```

SAVE FILE!

## Exercise: Add index2.html to cookbook's files/default directory

 OPEN IN EDITOR: cookbooks/apache/files/default/index2.html

```
<html>
  <body>
    <h1>Hello, world!</h1>
    <h2>This is index2.html</h2>
    <p>We configured this in the recipe</p>
  </body>
</html>
```

SAVE FILE!

## Exercise: Upload the cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache          [0.1.0]
Uploaded 1 cookbook.
```

# Exercise: Re-run Chef Client

```
chef@node1:~$ sudo chef-client
```

```
Starting Chef Client, version 11.8.2
resolving cookbooks for run list: ['apache']
...
Recipe: apache::default
  * package[httpd] action install (up to date)
  * service[httpd] action enable (up to date)
  * service[httpd] action start (up to date)
  * cookbook_file[/var/www/html/index.html] action create
    - update content in file /var/www/html/index.html from 479e29 to 153143
      --- /var/www/html/index.html 2014-01-07 04:31:24.049155232 -0500
      +++ /tmp/.index.html20140107-12296-11u8kv7 2014-01-07 04:32:45.605469431 -0500
      @@ <1,> +1,> @@
      <html>
        <body>
          <h1>Hello, world!</h1>
        <!-->This is index1.html</p>
        + <h2>This is index2.html</h2>
        + <p>We configured this in the attributes file</p>
        </body>
      </html>
    - restore selinux security context

Chef Client finished, 1 resources updated
```

# Exercise: Verify new homepage works

- Open a web browser
- Recipe value has taken precedence



Hello, world!

This is index2.html

We configured this in the recipe

## Exercise: Viewing Attributes via WebUI

The screenshot shows the Chef WebUI interface. In the top navigation bar, the 'Nodes' tab is selected. Below it, a sidebar on the left contains links for 'Delete', 'Manage Tags', 'Reset Key', 'Edit Run List', and 'Edit Attributes'. The main content area displays a table titled 'Showing All Nodes' with one row for 'node2'. The table has columns for 'Node Name', 'Platform', and 'FQDN'. The details for 'node2' are shown in a modal window. The 'Attributes' tab is selected in the modal's navigation bar. The attributes are listed under sections: 'apache' (with 'indexfile: index2.html'), 'tags', 'languages', 'kernel', and 'os: linux' (with 'os\_version: 2.6.32-358.23.2.el6.x86\_64' and 'hostname: CentOS643').

## Exercise: Viewing attributes using knife

```
$ knife node show node1 -l | more
```

```
...
Attributes:
tags:

Default Attributes:
apache:
  indexfile: index2.html

Override Attributes:

Automatic Attributes (Ohai Data):
block_device:
  dm-0:
    removable: 0
    size:      31563776
...

```

## Setting Attributes in Roles and Environments

- Attributes can also be set in Roles and Environments

- These use raw JSON or Ruby format

```
default_attributes({ "apache" => { "port" => 80 } })
```

- More later in Roles and Environments sections...

## Default Attribute Precedence



- Please note this is a simplified diagram, and the precedences shown can be overridden

# Questions

- What is an attribute?
- Where can attributes be set?
- What knife command can you use to view attributes?
- How many levels of attribute precedence are there?
- Which takes precedence, a default attribute set in the attribute file or in a recipe?

## Attributes, Templates, and Cookbook Dependencies

Writing an MOTD Cookbook

## Lesson Objectives

- After completing the lesson, you will be able to
  - Describe Cookbook Attribute files
  - Use ERB Templates in Chef
  - Explain Attribute Precedence
  - Describe Cookbook Metadata
  - Specify cookbook dependencies
  - Perform the cookbook creation, upload, and test loop

## The Problem and the Success Criteria

- **The Problem:** We need to add a message that appears at login that states:
  - "This server is property of COMPANY"
  - "This server is in-scope for PCI compliance" if the server is, in fact, in scope.
- **Success Criteria:** We see the message when we log in to the test node

## Possible Solutions

- We could use an MOTD cookbook to supply the 'Company' name as a node attribute and write out the message
- The line on PCI compliance only gets written if a 'PCI Compliance' attribute is set to 'true' for that Company

## We have a small problem...

- We don't have a node attribute for 'Company', nor if the server is in scope for 'PCI Compliance'
- Also, well factored cookbooks only contain the information relevant to their domain
  - We know we will likely have other things related to PCI (security settings, for example)
- So the best thing to do is create a separate PCI cookbook, and add our attribute there

# Solution

- We'll create an MOTD cookbook to write out the MOTD - this cookbook configures 'Company' name as a node attribute
- We'll also create a PCI cookbook which configures a 'PCI compliance' attribute - this attribute is then read by MOTD cookbook
- We'll create the MOTD cookbook first, then we'll create the PCI cookbook

## Exercise: Create a cookbook named 'motd'

```
$ knife cookbook create motd
```

```
** Creating cookbook motd
** Creating README for cookbook: motd
** Creating CHANGELOG for cookbook: motd
** Creating metadata for cookbook: motd
```

## Exercise: Create a default.rb attribute file

 OPEN IN EDITOR: cookbooks/motd/attributes/default.rb

```
default["motd"]["company"] = "Chef"
```

SAVE FILE!

- Creates a new Node attribute: node[ "motd" ][ "company" ]
- Sets the values to the string "Chef"
- Note: there is no space between 'default' and '['
- The 'motd' in the attribute is just convention so you know the cookbook the attribute was set in. This is not an enforced syntax

## Exercise: Open the default recipe in your editor

 OPEN IN EDITOR: cookbooks/motd/recipes/default.rb

```
#  
# Cookbook Name:: motd  
# Recipe:: default  
#  
# Copyright 2013, YOUR_COMPANY_NAME  
#  
# All rights reserved - Do Not Redistribute  
#
```

## What resource should we use?

- We could try and use a `cookbook_file` here, and rely on the file copy rules
  - Create a file per server, basically
- Obviously, that's dramatically inefficient
- Instead, we will render a **template** - a file that is a mixture of the contents we want, and embedded Ruby code

11

### Exercise: Add template resource for /etc/motd

- Use a **template** resource
- The **name** is "`/etc/motd`"
- The resource has two parameters
  - **source** is "`motd.erb`"
  - **mode** is "`0644`"

11

# The template[/etc/motd] resource



OPEN IN EDITOR: cookbooks/motd/recipes/default.rb

```
#  
# Cookbook Name:: motd  
# Recipe:: default  
#  
# Copyright 2013, YOUR_COMPANY_NAME  
#  
# All rights reserved - Do Not Redistribute  
#  
  
template "/etc/motd" do  
  source "motd.erb"  
  mode "0644"  
end
```

SAVE FILE!

## Exercise: Open motd.erb in your Editor



OPEN IN EDITOR: cookbooks/motd/templates/default/motd.erb

```
This server is property of <%= node["motd"]["company"] %>  
<% if node["pci"]["in_scope"] -%>  
This server is in-scope for PCI compliance  
<% end -%>
```

SAVE FILE!

- "erb" stands for "Embedded Ruby"

## Exercise: Open motd.erb in your Editor



[OPEN IN EDITOR: cookbooks/motd/templates/default/motd.erb](#)

```
This server is property of <%= node["motd"]["company"] %>
<% if node["pci"]["in_scope"] -%>
This server is in-scope for PCI compliance
<% end -%>
```

- To embed a value within an ERB template:
  - Start with <%=
  - Write your Ruby expression - most commonly a node attribute
  - End with %>

## Exercise: Open motd.erb in your Editor



[OPEN IN EDITOR: cookbooks/motd/templates/default/motd.erb](#)

```
This server is property of <%= node["motd"]["company"] %>
<% if node["pci"]["in_scope"] -%>
This server is in-scope for PCI compliance
<% end -%>
```

- You can use any Ruby construct in a template
  - Starting with <% will evaluate the expression, but not insert the result
  - Ending with -%> will not insert a line in the resulting file
  - Note there are no spaces in 'node[ "pci" ][ "in\_scope" ]'

## Templates Are Used For Almost All Configuration Files

- Templates are very flexible ways to create your configuration files
- Coupled with Chef's attribute precedence rules, you can create very effective, data-driven cookbooks

### Best Practice: Recipes contain the pattern, attributes supply the details

- Recipes contain the pattern for how to do something. ("How we deploy Tomcat")
- Attributes contain the details. ("What port do we run tomcat on?")

## Exercise: Upload the motd cookbook

```
$ knife cookbook upload motd
```

```
Uploading motd [0.1.0]
Uploaded 1 cookbook.
```

## Exercise: Create a cookbook named 'pci'

```
$ knife cookbook create pci
```

```
** Creating cookbook pci
** Creating README for cookbook: pci
** Creating CHANGELOG for cookbook: pci
** Creating metadata for cookbook: pci
```

## Exercise: Create a default.rb attribute file



OPEN IN EDITOR: cookbooks/pci/attributes/default.rb

```
default["pci"]["in_scope"] = true
```

SAVE FILE!

- Creates a new Node attribute: node["pci"]["in\_scope"]
- Sets the value to the Ruby 'true' literal

## Best Practice: Always use 'default' attributes in your cookbooks

- When setting an attribute in a cookbook, it should (almost) always be a **default** attribute

## Exercise: Upload the PCI cookbook

```
$ knife cookbook upload pci
```

```
Uploading pci [0.1.0]
Uploaded 1 cookbook.
```

## Exercise: Add the motd recipe to your test node's run list

```
$ knife node run_list add node1 "recipe[motd]"
```

```
node1:
  run_list:
    - recipe[apache]
    - recipe[motd]
```

## Exercise: Add the motd recipe to your test node's run list

```
$ knife node show node1
```

```
Node Name: node1
Environment: _default
FQDN: centos63.example.com
IP: 10.160.201.90
Run List: recipe[apache], recipe[motd]
Roles:
Recipes: apache
Platform: centos 6.4
Tags:
```

## Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
Starting Chef Client, version 11.8.2
...
- mord
Compiling Cookbooks...
Converging 4 resources
Recipes apache[_default]
 * package[httpd] action install [2014-01-06T09:14:33-05:00] INFO: Processing package[httpd] action install (apache::_default line 6)
  (up to date)
 * service[httpd] action enable [2014-01-06T09:14:33-05:00] INFO: Processing service[httpd] action enable (apache::_default line 13)
  (up to date)
 * service[httpd] action start [2014-01-06T09:14:33-05:00] INFO: Processing service[httpd] action start (apache::_default line 13)
  (up to date)
 * cookbook_file[/var/www/html/index.html] action create [2014-01-06T09:14:33-05:00] INFO: Processing cookbook_file[/var/www/html/index.html] action create (motd::_default line 9)
  (up to date)
Recipes motd[_default]
 * template[/etc/motd] action create [2014-01-06T09:14:37-05:00] INFO: Processing template[/etc/motd] action create (motd::_default line 9)
  - create new file /etc/motd
=====
Error executing action "create" on resource "template[/etc/motd]"
```

**FAIL!**

```
Chef::Mixin::Template::TemplateError
-----
undefined method `[]' for nil:NilClass
```

You probably see this at the bottom of your screen...

```
Template Context:  
----  
on line #2  
1: This server is property of <%= node["motd"]["company"] %>  
2: <% if node["pci"]["in_scope"] -%>  
3: This server is in-scope for PCI compliance  
4: <% end -%>  
  
[2014-01-06T09:14:37-05:00] INFO: Running queued delayed notifications before re-raising  
exception  
[2014-01-06T09:14:37-05:00] ERROR: Running exception handlers  
[2014-01-06T09:14:37-05:00] ERROR: Exception handlers complete  
[2014-01-06T09:14:37-05:00] FATAL: Stacktrace dumped to /var/chef/cache/chef-stacktrace.out  
Chef Client failed. 0 resources updated  
[2014-01-06T09:14:37-05:00] INFO: Sending resource update report (run-id: c001b9d5-  
c2f6-4026-9cc6-ff0901aa563c)  
[2014-01-06T09:14:37-05:00] ERROR: "\xE2" from ASCII-8BIT to UTF-8  
[2014-01-06T09:14:37-05:00] FATAL: Chef::Exceptions::ChildConvergeError: Chef run process  
exited unsuccessfully (exit code 1)
```

## Stack Traces

- A stack trace tells you where in a program an error occurred
- They can (obviously) be very detailed
- They can also be intensely useful, as they supply the data you need to find a problem

# Scroll up

- In this case, Chef actually knows exactly what went wrong.
- Scroll up to find out.

```
Error executing action "create" on resource 'template[/etc/motd]'

-----
Chef::Mixin::Template::TemplateError
-----
undefined method `()' for nil:NilClass

Resource Declaration:
-----
# In /var/chef/cache/cookbooks/motd/recipes/default.rb

10: template "/etc/motd" do
11:   source "motd.erb"
12:   mode "0644"
13: end
```

## We do not have the attribute we are using in the conditional

```
INFO: Run List is [recipe[apache], recipe[motd]]
INFO: Run List expands to [apache, motd]
INFO: Starting Chef Run for node1.local
INFO: Running start handlers
INFO: Start handlers complete.
resolving cookbooks for run list: ["apache", "motd"]
INFO: Loading cookbooks [apache, motd]
```

- Can anyone guess why?
- We did not load the PCI cookbook!

## nil:NilClass error

- The bottom line is when you see this error

```
undefined method `[]' for nil:NilClass
```

- It most often means you tried to look up a node attribute that does not exist!

## Best Practice: Make sure cookbooks have default values

- If a cookbook needs an attribute to exist, it should
  1. Define a default value for it in an attribute file, or
  2. Depend on another cookbook that does
- **Never rely on an attribute being created manually**

## Exercise: Add a dependency on the PCI cookbook to the MOTD cookbook



[OPEN IN EDITOR: cookbooks/motd/metadata.rb](#)

```
maintainer      "YOUR_COMPANY_NAME"
maintainer_email "YOUR_EMAIL"
license         "All rights reserved"
description     "Installs/Configures motd"
long_description IO.read(File.join(File.dirname(__FILE__), 'README.md'))
version         "0.1.0"
```

## Exercise: Add a dependency on the PCI cookbook to the MOTD cookbook



[OPEN IN EDITOR: cookbooks/motd/metadata.rb](#)

```
maintainer      "YOUR_COMPANY_NAME"
maintainer_email "YOUR_EMAIL"
license         "All rights reserved"
description     "Installs/Configures motd"
long_description IO.read(File.join(File.dirname(__FILE__), 'README.md'))
version         "0.1.0"
depends "pci"
```

- Cookbooks that depend on other cookbooks will cause the dependent cookbook to be downloaded to the client, and evaluated

# Cookbook Metadata



[OPEN IN EDITOR: cookbooks/motd/metadata.rb](#)

```
maintainer      "YOUR_COMPANY_NAME"
maintainer_email "YOUR_EMAIL"
license         "All rights reserved"
description     "Installs/Configures motd"
long_description IO.read(File.join(File.dirname(__FILE__), 'README.md'))
version         "0.1.0"
depends "pci"
```

## Cookbook Attributes are applied for all downloaded cookbooks!

- Cookbooks downloaded as dependencies will have their attribute files evaluated
- Even if there is no recipe from the cookbook in the run-list

## Exercise: Upload the motd cookbook

```
$ knife cookbook upload motd
```

```
Uploading motd [0.1.0]
Uploaded 1 cookbook.
```

## Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
Forking chef instance to converge...
Starting Chef Client, version 11.8.2
*** Chef 11.8.2 ***
Chef-client pid: 15152
Run List is [recipe[apache], recipe[motd]]
Run List expanded to [apache, motd]
Starting Chef Client or node...
-
- update co   in   /m   com e3   to 62
  (new con   t   a   l   get su   sed)
- restore sel   cit,
Chef Run complete in 00:00:00.035877
Removing cookbooks/   /files/   /built/index2
from   cache; it is no longer needed by chef-client.
Running report handlers
Report handlers complete
Chef Client finished, 1 resources updated
Sending resource update report (run-id: f923449e-cc7b-45f7-a9fb-ac1da3653557)
```



## Exercise: Check your work

```
chef@node1:~$ cat /etc/motd
```

```
This server is property of Chef  
This server is in-scope for PCI compliance
```

## Exercise: Show your test node's pci attribute

```
$ knife search node "pci: *" -a pci
```

```
1 items found
```

```
node1:  
  pci:  
    in_scope: true
```

## Exercise: Set attribute to false



OPEN IN EDITOR: cookbooks/pci/attributes/default.rb

```
default["pci"]["in_scope"] = false
```

SAVE FILE!

- Set the attribute to false

## Exercise: Upload the PCI cookbook

```
$ knife cookbook upload pci
```

```
Uploading pci [0.1.0]
Uploaded 1 cookbook.
```

## Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
Forking chef instance to converge...
Starting Chef Client, version 11.8.2
*** Chef 11.8.2 ***
Chef-client pid: 15350
Run List is [recipe[apache], recipe[motd]]
Run List expands to [apache, motd]
Starting Chef Run for node1
-
template[/etc/motd] updated file contents /etc/motd

  - update content in file /etc/motd from 62e8bb9 to d36elf
    (current file is binary, diff output suppressed)
  - restore selinux security context

Chef Run complete in 5.634796214 seconds
Running report handlers
Report handlers complete
Chef Client finished, 1 resources updated
Sending resource update report (run-id: 06edc909-6740-49bb-971d-82657e066236)
```

## Exercise: Check your work

```
chef@node1:~$ cat /etc/motd
```

```
This server is property of Chef
```

## Exercise: Show your test node's pci attribute

```
$ knife node show node1 -a pci
```

```
1 items found

node1:
  pci:
    in_scope: false
```

## Congratulations!

- You now know the 3 most important resources in the history of configuration management
  - Package
  - Template
  - Service

## Review Questions

- What goes in a cookbook's attribute files?
- What are the 3 different levels of precedence?
- When do you need to specify a cookbook dependency?
- What does <%> mean, and where will you encounter it?
- What are the 3 most important resources in configuration management?

## Template Variables, Notifications, and Controlling Idempotency

Refactoring the Apache Cookbook

## Lesson Objectives

- After completing the lesson, you will be able to
  - Use the **execute** resource
  - Control idempotence manually with **not\_if** and **only\_if**
  - Navigate the Resources page on [docs.chef.io](https://docs.chef.io)
  - Describe the Directory resource
  - Implement resource notifications
  - Explain what Template Variables are, and how to use them
  - Use Ruby variables, loops, and string expansion

## The Problem and the Success Criteria

- **The Problem:** We need to deploy multiple custom home pages running on different ports
- **Success Criteria:** Be able to view our custom home page

## Exercise: Change the cookbook's version number in the metadata



OPEN IN EDITOR: cookbooks/apache/metadata.rb

```
maintainer      "YOUR_COMPANY_NAME"
maintainer_email "YOUR_EMAIL"
license         "All rights reserved"
description     "Installs/Configures apache"
long_description IO.read(File.join(File.dirname(__FILE__), 'README.md'))
version        "0.2.0"
```

SAVE FILE!

- Major, Minor, Patch
- Semantic Versioning Policy: <http://semver.org/>

## Exercise: Edit the attribute file default.rb



OPEN IN EDITOR: cookbooks/apache/attributes/default.rb

```
default["apache"]["sites"]["clowns"] = { "port" => 80 }
default["apache"]["sites"]["bears"] = { "port" => 81 }
```

SAVE FILE!

- We add information about the sites we need to deploy
  - One about Clowns, running on port 80
  - One about Bears, running on port 81

## Exercise: Open the default apache recipe in your editor



[OPEN IN EDITOR: cookbooks/apache/recipes/default.rb](#)

```
package "httpd" do
  action :install
end

service "httpd" do
  action [:enable, :start]
end

node.default["apache"]["indexfile"] = "index2.html"
cookbook_file "/var/www/html/index.html" do
  source node["apache"]["indexfile"]
  mode "0644"
end

# Tasks:-
# Disable the default virtual host
# Iterate over the apache sites
# Set the document root
# Add a template for Apache virtual host configuration
# Add a directory resource to create the document_root
# Add a template resource for the virtual host's index.html
```

## Exercise: Use the 'execute' resource to disable the default Apache virtual host



[OPEN IN EDITOR: cookbooks/apache/recipes/default.rb](#)

```
service "httpd" do
  action [ :enable, :start ]
end

# Disable the default virtual host
execute "mv /etc/httpd/conf.d/welcome.conf /etc/httpd/conf.d/welcome.conf.disabled" do
  only_if do
    File.exist?("/etc/httpd/conf.d/welcome.conf")
  end
  notifies :restart, "service[httpd]"
end

cookbook_file "/var/www/html/index.html" do
```

**SAVE FILE!**

- Renames the default site config file, but only if the file exists
- If the action succeeds, restart Apache

## Execute resources are generally not idempotent

- Chef will stop your run if a resource fails
- Most command line utilities are not idempotent and can only be run once - they assume a human being is interacting with, and understands, the state of the system
  - e.g. 'mv /foo/file1 /bar' will work the first time its run, but will fail the second time
- The result is - it's up to you to make execute resources idempotent

## Enter the `not_if` and `only_if` metaparameters

```
only_if do
  File.exist?("/etc/httpd/conf.d/welcome.conf")
end
```

- The `only_if` parameter causes the resources actions to be taken only if its argument returns true
- The `not_if` parameter is the opposite of `only_if` - the actions are taken only if its argument returns false

## Building the resource collection - revisited

- It is important to understand the difference between these two pieces of code

```
execute "start-apache" do
  not_if <code to check if apache is running>
  notifies :start, "service[httpd]"
end
```

This code is placed in the resource collection during the 'compile' phase and executed during the 'execute' phase

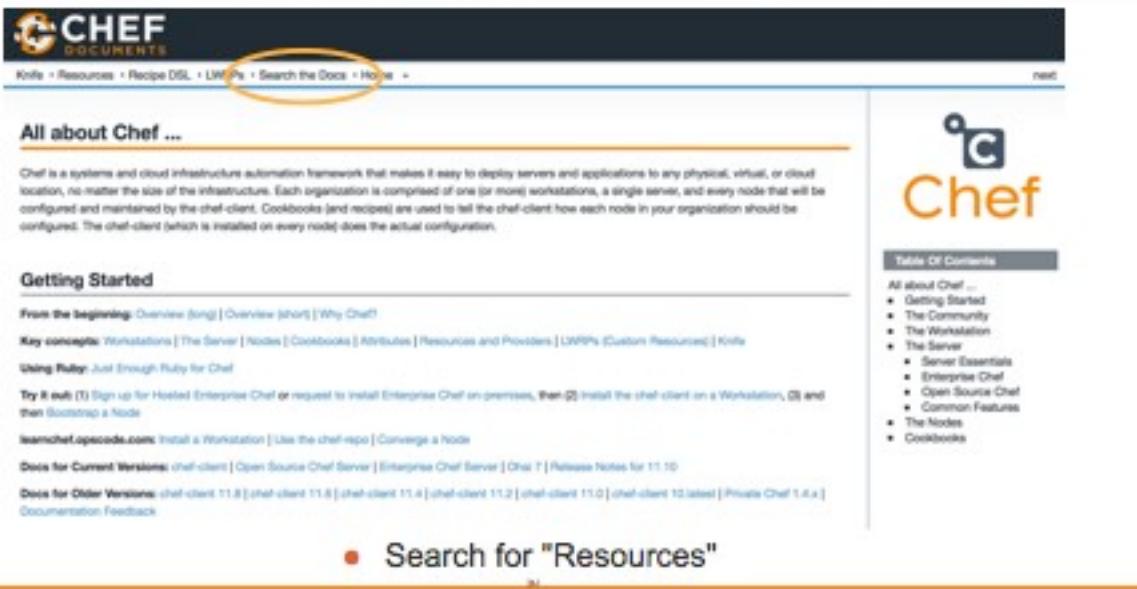
```
if !<code to check if apache is running>
  execute "start-apache" do
    notifies :start, "service[httpd]"
  end
end
```

This code is executed during the 'compile' phase before any Chef DSL  
Executing code during the compile phase could potentially prevent the resource from getting added to the resource collection

## Best Practice: The Chef Docs Site

- The Chef Docs Site is the home for all of the documentation about Chef.
  - It is very comprehensive
  - It has a page on every topic
- <http://docs.chef.io>
- Let's use the docs to learn more about **not\_if** and **only\_if**

## Exercise: Search for more information about Resources

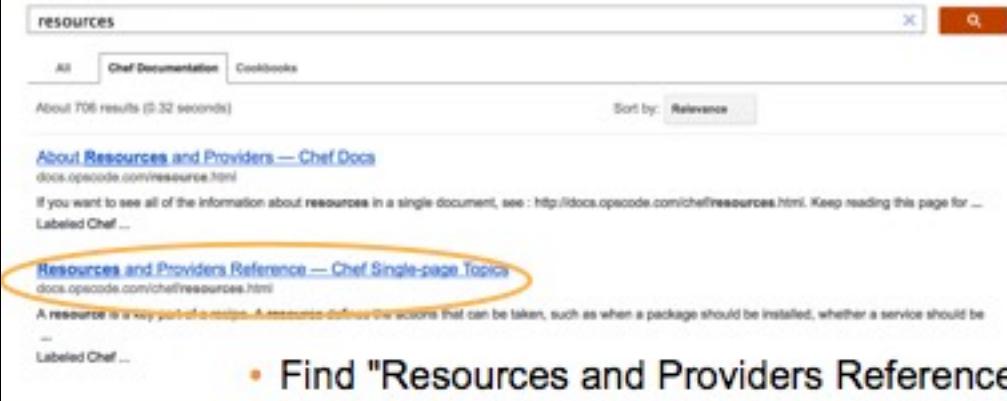


The screenshot shows the official Chef Documentation website. At the top, there's a navigation bar with links to 'Knife', 'Resources', 'Recipe DSL', 'LWRPs', 'Search the Docs', and 'Home'. A yellow circle highlights the 'Search the Docs' link. Below the navigation, there's a section titled 'All about Chef ...' which contains a brief introduction to Chef. Further down, there's a 'Getting Started' section with links to 'From the beginning', 'Key concepts', 'Using Ruby', 'Try it out!', and 'Learnchef'. On the right side, there's a 'Table Of Contents' sidebar with a tree structure: 'All about Chef ...' → 'Getting Started' → 'The Community', 'The Workstation', 'The Server' (which further branches into 'Server Essentials', 'Enterprise Chef', 'Open Source Chef', 'Common Features'), 'The Nodes', and 'Cookbooks'. At the bottom of the page, there's a bullet point: '• Search for "Resources"'.

## Exercise: Search for more information about Resources

### Search the Documentation for Chef

From here you can use a scoped Google search query to search all of the documentation about Chef that is located at docs.opscode.com. (This page requires JavaScript be enabled to view the search box.)



The screenshot shows the search results for the term 'resources' on the Chef Documentation site. The search bar contains 'resources' and has a magnifying glass icon. Below the search bar, there are filters for 'All', 'Chef Documentation', and 'Cookbooks', with 'Chef Documentation' being selected. It also shows 'About 706 results (0.32 seconds)' and a 'Sort by: Relevance' dropdown. The results list includes a link to 'About Resources and Providers — Chef Docs' (docs.opscode.com/resource.html) and a note about viewing resources in a single document. Further down, there's a section titled 'Resources and Providers Reference — Chef Single-page Topics' (docs.opscode.com/chefresources.html), which is highlighted with a yellow circle. A descriptive text follows: 'A resource is a key part of a recipe. A resource defines the actions that can be taken, such as when a package should be installed, whether a service should be'. At the bottom, there's another bullet point: '• Find "Resources and Providers Reference"'.



# The Resources Page

## Resources and Providers Reference

A resource is a key part of a recipe. A resource defines the actions that can be taken, such as when a package should be installed, whether a service should be enabled or restarted, which groups, users, or groups of users should be created, where to put a collection of files, what the name of a new directory should be, and so on. During a chef-client run, each resource is identified and then associated with a provider. The provider then does the work to complete the action defined by the resource. Each resource is processed in the same order as they appear in a recipe. The chef-client ensures that the same actions are taken the same way everywhere and that actions produce the same result every time. A resource is implemented within a recipe using Ruby.

Where a resource represents a piece of the system (and its desired state), a provider defines the steps that are needed to bring that piece of the system from its current state into the desired state. These steps are de-coupled from the request itself. The request is made in a recipe and is defined by a lightweight resource. The steps are then defined by a lightweight provider.

The Chef::Platform class maps providers to platforms (and platform versions). Chef, as part of every chef-client run, verifies the `platform` and `platform_version` attributes on each node. The chef-client then uses those values to identify the correct provider, build an instance of that provider, identify the current state of the resource, do the specified action, and then mark the resource as updated (if changes were made). For example, given the following resource:

```
directory "/tmp/folder" do
  owner "root"
  group "root"
  mode 0755
  action :create
end
```



### Table Of Contents

#### Resources and Providers Reference

- Common Functionality for all Resources
  - Actions
    - Examples
  - Attributes
    - Examples
  - Guards
    - Attributes
    - Arguments
    - `not_if` Examples
    - `only_if` Examples
  - `lifecycle` Attribute Implementation
  - Notifications
    - Notifications Timers
    - Notifications Syntax
      - Examples
    - Subscribed Syntax

# Notifications

```
notifies :restart, "service[httpd]"
```

- Resource Notifications in Chef are used to trigger an action on a resource when the current resources actions are successful.
- "If we delete the site, restart apache"
- The first argument is an action, and the second argument is the string representation of a given resource
- Like `not_if` and `only_if`, **notifies** is a resource metaparameter - any resource can notify any other

## Exercise: Iterate over each apache site

 OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
execute "mv /etc/httpd/conf.d/welcome.conf /etc/httpd/conf.d/welcome.conf.disabled" do
  only_if do
    File.exist?("/etc/httpd/conf.d/welcome.conf")
  end
  notifies :restart, "service[httpd]"
end

node.default["apache"]["indexfile"] = "index2.html"
cookbook_file "/var/www/index.html" do
  source node["apache"]["indexfile"]
  mode "0644"
end
```

- Delete the cookbook\_file resource

## Exercise: Iterate over each apache site and set document\_root

 OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Disable the default virtual host
execute "mv /etc/httpd/conf.d/welcome.conf /etc/httpd/conf.d/welcome.conf.disabled" do
  only_if do
    File.exist?("/etc/httpd/conf.d/welcome.conf")
  end
  notifies :restart, "service[httpd]"
end

# Iterate over the apache sites
node["apache"]["sites"].each do |site_name, site_data|
  # Set the document root
  document_root = "/srv/apache/#{site_name}"
```

**SAVE FILE!**

- node["apache"]["sites"] is a ruby hash, with keys and values

## Exercise: Iterate over each apache site

```
node["apache"]["sites"].each do |site_name, site_data|
  document_root = "/srv/apache/#{site_name}"
```

- Calling .each loops over each site

```
default["apache"]["sites"]["clowns"] = { "port" => 80 }
default["apache"]["sites"]["bears"] = { "port" => 81 }
```

- First pass
  - site\_name = 'clowns'
  - site\_data = { "port" => 80 }
- Second pass
  - site\_name = 'bears'
  - site\_data = { "port" => 81 }

## Exercise: Iterate over each apache site

```
node["apache"]["sites"].each do |site_name, site_data|
  document_root = "/srv/apache/#{site_name}"
```

- Create a variable called document\_root
- #{site\_name} means "insert the value of site\_name here"
- First pass
  - The value is the string "/srv/apache/clowns"
- Second pass
  - The value is the string "/srv/apache/bears"

## Exercise: Add a template for Apache virtual host configuration

```
# Iterate over the apache sites
node["apache"]["sites"].each do |site_name, site_data|
  # Set the document root
  document_root = "/srv/apache/#{site_name}"

  # Add a template for Apache virtual host configuration
  template "/etc/httpd/conf.d/#{site_name}.conf" do
    source "custom.erb"
    mode "0644"
    variables(
      :document_root => document_root,
      :port => site_data["port"]
    )
    notifies :restart, "service[httpd]"
  end
end
```

## Template Variables

- Not all data you might need in a template is necessarily node attributes
- The **variables** parameter lets you pass in custom data for use in a template

## Exercise: Add a directory resource to create the document\_root

- Use a **directory** resource
- The **name** is `document_root`
- The resource has two parameters
  - **mode** is "0755"
  - **recursive** is true
- Use the Resources page on the Docs Site to read more about what **recursive** does.

## Exercise: Add the directory resource

```
# Add a template for Apache virtual host configuration
template "/etc/httpd/conf.d/#{site_name}.conf" do
  source "custom.erb"
  mode "0644"
  variables(
    :document_root => document_root,
    :port => site_data["port"]
  )
  notifies :restart, "service[httpd]"
end

# Add a directory resource to create the document_root
directory document_root do
  mode "0755"
  recursive true
end
```

## Exercise: Add a template resource to supply the index.html for the virtual host

```
# Add a directory resource to create the document_root
directory document_root do
  mode "0755"
  recursive true
end

# Add a template resource for the virtual host's index.html
template "#{document_root}/index.html" do
  source "index.html.erb"
  mode "0644"
  variables(
    :site_name => site_name,
    :port => site_data["port"]
  )
end
end
```

## Don't forget the last "end"

```
# Add a template resource for the virtual host's index.html
template "#{document_root}/index.html" do
  source "index.html.erb"
  mode "0644"
  variables(
    :site_name => site_name,
    :port => site_data["port"]
  )
end
end
```

See the correct, whole file at <https://gist.github.com/8954699>

## Exercise: Add custom.erb to your templates directory



OPEN IN EDITOR: cookbooks/apache/templates/default/custom.erb

- Note the two **template variables** are prefixed with an @ symbol
- If you are feeling hardcore, type it
- <https://gist.github.com/8955103>

```
<# if fport != 80 ->
  Listen <%= fport %>
<# end ->

<VirtualHost *:<%= fport %>>
  ServerAdmin webmaster@localhost

  DocumentRoot <%= @document_root %>
  <Directory />
    Options FollowSymLinks
    AllowOverride None
  </Directory>
  <Directory <%= @document_root %>>
    Options Indexes FollowSymLinks Multiviews
    AllowOverride None
    Order allow,deny
    allow from all
  </Directory>
</VirtualHost>
```

SAVE FILE!

## Exercise: Add index.html.erb to your templates directory



OPEN IN EDITOR: cookbooks/apache/templates/default/index.html.erb

```
<html>
  <body>
    <h1>Welcome to <%= node[ "motd" ][ "company" ] %></h1>
    <h2>We love <%= @site_name %></h2>
    <%= node[ "ipaddress" ] %>:<%= @port %>
  </body>
</html>
```

SAVE FILE!

- Note the two **template variables** are prefixed with an @ symbol
- <https://gist.github.com/8955080>

## Exercise: Upload the Apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [0.2.0]
Uploaded 1 cookbook.
```

## Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
=====
Error executing action 'restart' on resource 'service[httpd]'

Mixlib::ShellOut::ShellCommandFailed
=====
Expected process to
---- begin output of /sbin/service httpd restart
----| Starting httpd...
----| httpd: Syntax error on line 11 of /etc/httpd/conf.d/vhosts.conf:
----|   directive "ServerName" is not allowed here
---- End output of /sbin/service httpd restart
----|
---- Result:
----| exit status 1
----|
---- Standard error:
----| httpd: Could not reliably determine the server's fully qualified domain name, using 127.0.0.1 for ServerName
----|
---- Resource Declaration:
----| # In /var/chef/cache/cookbooks/apache/recipes/default.rb
----|
----| 43: service "httpd" do
----| 44:   action [ :enable, :restart ]
----| 45: end
----| ...
----|
```

**FAIL!**

## Exercise: Fix typo in custom.erb



OPEN IN EDITOR: cookbooks/apache/templates/default/custom.erb

```
<? if $port != 80 -?>
  Listen <?= $port ?>
<? end -?>

<VirtualHost *:<?= $port ?>>
  ServerAdmin webmaster@localhost

  DocumentRoot <?= $document_root ?>
  <Directory />
    Options FollowSymLinks
    AllowOverride None
  </Directory>
  <Directory <?= $document_root ?>>
    Options Indexes FollowSymLinks MultiViews
    AllowOverride None
    Order allow,deny
    allow from all
  </Directory>
</VirtualHost>
```

SAVE FILE!

## Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
...
Error executing action "restart" on resource "service[httpd]"
```

```
mixlib-shellout::ShellOut: sh -c /etc/init.d/httpd restart
=====
Expected process to
---- Begin output of /etc/init.d/httpd restart
stderr: syntax error near unexpected token `}'
---- End output of /etc/init.d/httpd restart
Run /sbin/service httpd restart returned
```

```
Resource Declaration:
```

```
# In /var/chef/cache/cookbooks/apache/recipes/default.rb

43: service "httpd" do
44:   action [ :enable, :restart ]
45: end
...
```

**FAIL!**

## Exercise: Move the service resource to end of recipe

```
# Add a template resource for the virtual host's index.html
template "#{document_root}/index.html" do
  source "index.html.erb"
  mode "0644"
  variables(
    :site_name => site_name,
    :port => site_data["port"]
  )
end
end

service "httpd" do
  action [ :enable, :start ]
end
```

## Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
- start service service[httpd]

Recipe: motd::default
 * template[/etc/motd] action create[2014-02-11T04:34:23-05:00] INFO: Processing template[/etc/motd]
action create (motd::default line 9)
 (up to date)
[2014-02-11T04:34:50-05:00]      0: temp    etc    /com   novas.com   ding   start action to
service[httpd]    yed)
Recipe: apache::ult
 * service[http]  ult  20  -11T04:34:51-05:00  rock   sex   httpd) action
restart (apache:
[2014-02-11T04:34:51-05:00]      0: httpd  e[httpd]  arte
 - restart service[httpd]

[2014-02-11T04:34:56-05:00] INFO: Run completed in 9.... 4215 seconds.
[2014-02-11T04:34:56-05:00] INFO: Running report handlers
[2014-02-11T04:34:56-05:00] INFO: Report handlers complete
Chef Client finished, 5 resources updated
[2014-02-11T04:34:56-05:00] INFO: Sending resource update report (run-id: d1e6c73e-9406-46b8-b022-d7bdf0868432)
```

# Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
[2014-01-07T05:13:43-05:00] INFO: execute[mv /etc/httpd/conf.d/welcome.conf /etc/httpd/conf.d/welcome.conf.disabled] sending restart action to service[httpd] (delayed)
Recipe: apache::default
  * service[httpd] action restart(2014-01-07T05:13:43-05:00) INFO: Processing service[httpd] action
    restart (apache::default line 13)
[2014-01-07T05:13:45-05:00] INFO: service[httpd] restarted
  - restart service service[httpd]

[2014-01-07T05:13:46-05:00] INFO: Chef Run complete in 10.389166164 seconds
[2014-01-07T05:13:46-05:00] INFO: Removing cookbooks/apache/files/default/index2.html from the
cache; it is no longer needed by chef-client.
[2014-01-07T05:13:46-05:00] INFO: Running report handlers
[2014-01-07T05:13:46-05:00] INFO: Report handlers complete
Chef Client finished, 8 resources updated
[2014-01-07T05:13:46-05:00] INFO: Sending resource update report (run-id: b929bc5f-4465-4669-8027-
b7d7b44741d9)
```

## Exercise: Verify the two sites are working!



## Best Practice: Recipes contain the pattern, attributes supply the details

- Recipes contain the pattern for how to do something. ("How we deploy apache virtual hosts")
- Attributes contain the details. ("What virtual hosts should we deploy?")

## Review Questions

- How do you control the idempotence of an Execute resource?
- Where can you learn the details about all the core resources in Chef?
- What is a notification?
- What is a template variable?
- What does `#{foo}` do in a Ruby string?

# Search

A deeper dive....

92.1.8



## Lesson Objectives

- After completing the lesson, you will be able to
  - Describe the query syntax used in search
  - Invoke a search from command line
  - Build a search into your recipe code

## What is search?

- Use search to query data indexed on the chef server
- Syntax is

```
$ knife search INDEX SEARCH_QUERY
```

- Where index can be client, environment, node, role, (or the name of a data bag)
- Search query **runs on** the server and is **invoked from** within a recipe, using knife or via WebUI

## Exercise: Use knife search

```
$ knife search node "*:*"
```

```
Node Name:    node1
Environment:  _default
FQDN:        centos63.example.com
IP:          10.160.201.90
Run List:    recipe[apache], recipe[motd]
Roles:
Recipes:     apache, motd
Platform:   centos 6.4
Tags:
```

## Query Syntax

- Chef search uses the Solr 'key:search\_pattern' syntax

```
$ knife search node "ipaddress:10.20.30.40"
```

- Use an asterisk ("\*") for  $\geq 0$  chars in a wildcard search

```
$ knife search node "ipaddress:10.*"  
$ knife search node "platfo*:ubuntu"
```

- Use a question mark (?) to replace a single character

```
$ knife search node "platform_version:12.0?"
```

## Indexing and search delay

- Search is based on an index, so if you write to a node it may not be immediately available

```
node.default["serveradmin"] = "Jane"  
search("node", "*:*").each do |servers|  
  user servers[ "serveradmin" ]  
end
```

BAD PRACTICE!

- Here, since the attribute value has just been written the search will not pick it up until the index is rebuilt
- Also, node attributes changed during a Chef run aren't posted to the server until the entire Chef run is successful
- But the value is stored in memory - so no need to search!

## Remember what you're searching for!

```
use knife          for nodes  
$ knife search node 'ipaddress:10.1.1.*'  
      to search      with this search criteria
```

- A successful search returns **all objects**, not just the attributes, that satisfy the search criteria
- In other words 'search for nodes containing this attribute', **NOT** 'search for this attribute across nodes'!

### Exercise: Find a node with the specific attribute value

```
$ knife search node "ipaddress:10.*"
```

```
1 items found

Node Name:    node1
Environment:  _default
FQDN:        centos63.example.com
IP:          10.160.201.90
Run List:    recipe[apache], recipe[motd]
Roles:
Recipes:     apache, motd
Platform:   centos 6.4
Tags:
```

## Exercise: Using Knife to find an attribute value

```
$ knife search node "*:*" -a ipaddress
```

```
1 items found
```

```
node1:
```

```
  ipaddress: 10.160.201.90
```

## Exercise: Return just the attribute value

```
$ knife search node "ipaddress:10.*" -a ipaddress
```

```
1 items found
```

```
node1:
```

```
  ipaddress: 10.160.201.90
```

## Exercise: Boolean matching

```
$ knife search node "ipaddress:10* AND  
platform_family:rhel"
```

```
1 items found

Node Name:  node1
Environment: _default
FQDN:      centos63.example.com
IP:        10.160.201.90
Run List:   recipe[apache], recipe[motd]
Roles:
Recipes:    apache, motd
Platform:   centos 6.4
Tags:
```

## Search from within a recipe

- You can invoke a search within a recipe, and use the results to apply further Chef primitives

```
search(INDEX, SEARCH_QUERY).each do |result|
  foo
end
```

- For example

```
search("node", "role:webserver").each do |webserver|
  [register webserver with load balancer]
  or
  [configure firewall so webserver can access database]
end
```

## Exercise: Search for an attribute in a recipe



OPEN IN EDITOR: cookbooks/apache/recipes/ip-logger.rb

```
search("node", "platform:centos").each do |server|
  log "The CentOS servers in your organization have the following
FQDN/IP Addresses:- #{server["fqdn"]}/#{server["ipaddress"]}"
end
```

SAVE FILE!

## Exercise: Upload the Apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [0.2.0]
Uploaded 1 cookbook.
```

## Exercise: Add the recipe to your test node's run list

```
$ knife node run_list add node1 'recipe[apache::ip-logger]'
```

```
node1:  
  run_list:  
    - recipe[apache]  
    - recipe[motd]  
    - recipe[apache::ip-logger]
```

## Recipe Naming

- Recipes are referenced using the notation '**cookbook::recipe**', where 'recipe' is the name of the '.rb' file in the "cookbook/recipe" directory
- The "default.rb" recipe for a given cookbook may be referred to just the cookbook name (e.g. '*mysql*')
- If we added another recipe to this cookbook named 'backup', we would refer to it as '*mysql::backup*'

# Exercise: Re-run the Chef Client

```
chef@node1:~$ sudo chef-client
```

```
Recipe: apache::ip-logger
 * log[The CentOS servers in your organization have the following FQDN/IP
 Addresses:- centos63.example.com/10.160.201.90] action
write[2014-01-07T05:42:21-05:00] INFO: Processing log[The CentOS servers in your
organization have the following FQDN/IP Addresses: Server centos63.example.com/
10.160.201.90] action write (apache::ip-logger line 4)
[2014-01-07T05:42:21-05:00] INFO: The CentOS servers in your organization have the
following FQDN/IP Addresses: Server centos63.example.com/10.160.201.90

[2014-01-07T05:42:22-05:00] INFO: Chef Run complete in 6.376181602 seconds
[2014-01-07T05:42:22-05:00] INFO: Running report handlers
[2014-01-07T05:42:22-05:00] INFO: Report handlers complete
Chef Client finished, 1 resources updated
[2014-01-07T05:42:22-05:00] INFO: Sending resource update report (run-id:
cecf34e9-eab6-4379-a5a6-d6d135ae8cc7)
```

## Exercise: Remove the recipe to your test node's run list

```
$ knife node run_list remove node1 'recipe[apache::ip-logger]'
```

```
node1:
  run_list:
    recipe[apache]
    recipe[motd]
```

## Questions

- What search engine is Chef search built upon?
- What is searchable in Chef?
- What character can you use for a wildcard search in Solr?
- Why should you not set an attribute and then immediately search for it?

## Recipe Inclusion, Data Bags, and Search

Writing a Users cookbook

## Lesson Objectives

- After completing the lesson, you will be able to
  - Explain what Data Bags are, and how they are used
  - Describe the User and Group resources
  - Use the `include_recipe` directive
  - Describe the role Search plays in recipes

## The Problem and the Success Criteria

- **The Problem:** Employees should have local user accounts created on servers, along with custom groups
- **Success Criteria:** We can add new employees and groups to servers dynamically

## Where should we store the user data?

- As we've seen, we could start by storing information about users as Node Attributes
- This is sort of a bummer, because we would be duplicating a lot of information - every user in the company would be stored in every Node object!
- Additionally, it would be very hard to integrate such a solution with another source of truth about users

## Introducing Data Bags

- A data bag is a **container** for **items** that represent information about your infrastructure that is not tied to a single node
- Examples
  - Users
  - Groups
  - Application Release Information

## Make a data\_bags directory

```
$ mkdir -p data_bags/users
```

(No output)

## Exercise: Create a data bag named users

```
$ knife data_bag create users
```

Created data\_bag[users]

## Exercise: Create a user item in the users data bag

 OPEN IN EDITOR: data\_bags/users/bobo.json

```
{  
  "id": "bobo",  
  "comment": "Bobo T. Clown",  
  "uid": 2000,  
  "gid": 0,  
  "home": "/home/bobo",  
  "shell": "/bin/bash"  
}
```

SAVE FILE!

## Exercise: Create the data bag item

```
$ knife data_bag from file users bobo.json
```

```
Updated data_bag_item[users::bobo]
```

## Exercise: Create another user in the users data bag

 OPEN IN EDITOR: data\_bags/users/frank.json

```
{  
  "id": "frank",  
  "comment": "Frank Belson",  
  "uid": 2001,  
  "gid": 0,  
  "home": "/home/frank",  
  "shell": "/bin/bash"  
}
```

SAVE FILE!

## Exercise: Create the data bag item

```
$ knife data_bag from file users frank.json
```

```
Updated data_bag_item[users::frank]
```

## Exercise: Show all the items in users data bag

```
$ knife search users "*;*"
```

```
2 items found

chef_type: data_bag_item
comment: Frank Balsom
data_bag: users
gid: 0
home: /home/frank
id: frank
shell: /bin/bash
uid: 2001

chef_type: data_bag_item
comment: Bobo T. Clown
data_bag: users
gid: 0
home: /home/bobo
id: bobo
shell: /bin/bash
uid: 2000
```

## Exercise: Find Bobo's shell in Chef

```
$ knife search users "id:bobo" -a shell
```

```
1 items found

data_bag_item_users_bobo:
  shell: /bin/bash
```

## Exercise: Create a data bag named groups

```
$ mkdir data_bags/groups  
$ knife data_bag create groups
```

```
Created data_bag[groups]
```

## Exercise: Create a group item in the group data bag



OPEN IN EDITOR: data\_bags/groups/clowns.json

```
{  
  "id": "clowns",  
  "gid": 3000,  
  "members": ["bobo", "frank"]  
}
```

SAVE FILE!

## Exercise: Create the data bag item

```
$ knife data_bag from file groups clowns.json
```

```
Updated data_bag_item[groups::clowns]
```

## Exercise: Show all the groups in Chef

```
$ knife search groups "*:*"
```

```
1 items found

chef_type:  data_bag_item
data_bag:    groups
gid:        3000
id:         clowns
members:
  bobo
  frank
```

## Exercise: Create a cookbook named 'users'

```
$ knife cookbook create users
```

```
** Creating cookbook users
** Creating README for cookbook: users
** Creating CHANGELOG for cookbook: users
** Creating metadata for cookbook: users
```

## Exercise: Open the default recipe in your editor



OPEN IN EDITOR: cookbooks/users/recipes/default.rb

```
#  
# Cookbook Name:: users  
# Recipe:: default  
#  
# Copyright 2013, YOUR_COMPANY_NAME  
#  
# All rights reserved - Do Not Redistribute  
#
```

SAVE FILE!

## Exercise: Open the default recipe in your editor

```
search("users", "*:*").each do |user_data|
  user user_data["id"] do
    comment user_data["comment"]
    uid user_data["uid"]
    gid user_data["gid"]
    home user_data["home"]
    shell user_data["shell"]
  end
end

include_recipe "users::groups"
```

- We use the same search we just tried with Knife in the recipe
- Each item is bound to `user_data`
- Use the Chef Docs for information about the `user` resource

## Exercise: Open the default recipe in your editor

```
search("users", "*:*").each do |user_data|
  user user_data["id"] do
    comment user_data["comment"]
    uid user_data["uid"]
    gid user_data["gid"]
    home user_data["home"]
    shell user_data["shell"]
  end
end

include_recipe "users::groups"
```

- `include_recipe` ensures another recipes resources are complete before we continue
- Chef will only include each recipe once

## Best Practice: Use include\_recipe liberally

- If there is a pre-requisite for your recipe that resides in another recipe (the JVM existing for your Java application, for example)
- Always use `include_recipe` to include it specifically, even if you put it in a run list

### Exercise: Open the `users::groups` recipe in your editor

 OPEN IN EDITOR: `cookbooks/users/recipes/groups.rb`

```
search("groups", "*:*").each do |group_data|
  group group_data["id"] do
    gid group_data["gid"]
    members group_data["members"]
  end
end
```

SAVE FILE!

- This file follows the same pattern as the default `users` recipe
- Use the Chef Docs for information about the **group** resource

## Exercise: Upload the users cookbook

```
$ knife cookbook upload users
```

```
Uploading users [0.1.0]
Uploaded 1 cookbook.
```

## Exercise: Add the users recipe to your test node's run list

```
$ knife node run_list add node1 'recipe[users]'
```

```
node1:
  run_list:
    recipe[apache]
    recipe[motd]
    recipe[users]
```

## Exercise: Re-run the Chef Client

```
chef@node1$ sudo chef-client
```

```
...
Recipe: users::default
 * user[bobo] action create[2014-01-07T06:16:26-05:00] INFO: Processing user[bobo] action create
(users::default line 10)
[2014-01-07T06:16:26-05:00] INFO: user[bobo] created

 - create user user[bobo]

 * user[frank] action create[2014-01-07T06:16:26-05:00] INFO: Processing user[frank] action create
(users::default line 10)
[2014-01-07T06:16:27-05:00] INFO: user[frank] created

 - create user user[frank]

Recipe: users::groups
 * group[circusmonkeys] action create[2014-01-07T06:16:27-05:00] INFO: Processing group[circusmonkeys] action
create (users::groups line 2)
[2014-01-07T06:16:27-05:00] INFO: group[circusmonkeys] created

 - create group[circusmonkeys]
...
...
```

## Exercise: Verify the users and groups exist

```
chef@node1:-$ cat /etc/passwd
```

```
frank:x:2001:0:Frank Belson:/home/frank:/bin/bash
bobo:x:2000:0:Bobo T. Clown:/home/bobo:/bin/bash
```

```
chef@node1:-$ cat /etc/group
```

```
clowns:x:3000:bobo,frank
```

## Let's review real quick..

- We just created a centralized user and group repository, from scratch
  - (That's kind of like what LDAP and Active Directory do, only they are, well, fancier.)
- Between Data Bags and Node Attribute precedence, Chef provides a plethora of ways to inform the patterns you use to configure your infrastructure

## Review Questions

- What are Data Bags?
- How are they used?
- What does the User resource do?
- What is `include_recipe`, and why is it useful?
- How does search work inside a recipe?
- What other applications do you see for search?
- How could we have used Data Bags in the refactored Apache recipe?
- Where would you go to find out more?

# Roles

Role-based Attributes and Merge Order Precedence

92.1.8

48



## Lesson Objectives

- After completing the lesson, you will be able to
- Explain what Roles are, and how they are used to provide clarity
- Discuss the Role Ruby DSL
- Show a Role with Knife
- Explain how merge order affects the precedence hierarchy
- Describe nested Roles

51

## What is a Role?

- So far, we've been just adding recipes directly to a single node
- But that's not how your infrastructure works - think about how you refer to servers
  - "*It's a web server*"
  - "*It's a database server*"
  - "*It's a monitoring server*"

## What is a Role?

- Roles allow you to conveniently encapsulate the run lists and attributes required for a server to "be" what you already think it is
- In practice, Roles make it **easy to configure many nodes identically** without repeating yourself each time

## Best Practice: Roles live in your chef-repo

- Like Data Bags, you have options with how to create a Role
- The best practice is that all of your Roles live in the roles directory of your chef-repo
- They can be created via the API and Knife, but it's nice to be able to see them evolve in your source control history

## Exercise: Create the webserver role



OPEN IN EDITOR: roles/webserver.rb

- A Role has a:
- name
- description
- run\_list

```
name "webserver"
description "Web Server"
run_list "recipe[apache]"
default_attributes({
  "apache" => {
    "sites" => {
      "admin" => {
        "port" => 8000
      }
    }
  }
})
```

SAVE FILE!

## Exercise: Create the webserver role

 OPEN IN EDITOR: roles/webserver.rb

- You can set default node attributes within a role.

```
name "webserver"
description "Web Server"
run_list "recipe[apache]"
default_attributes({
  "apache" => {
    "sites" => {
      "admin" => {
        "port" => 8000
      }
    }
  }
})
```

**SAVE FILE!**

## Exercise: Create the role

```
$ knife role from file webserver.rb
```

```
Updated Role webserver!
```

## Exercise: Show the role with knife

```
$ knife role show webserver
```

```
chef_type:          role
default_attributes:
  apache:
    sites:
      admin:
        port: 8000
description:        Web Server
env_run_lists:
json_class:         Chef::Role
name:               webserver
override_attributes:
run_list:           recipe[apache]
```

## Exercise: Search for roles with recipe[apache] in their run list

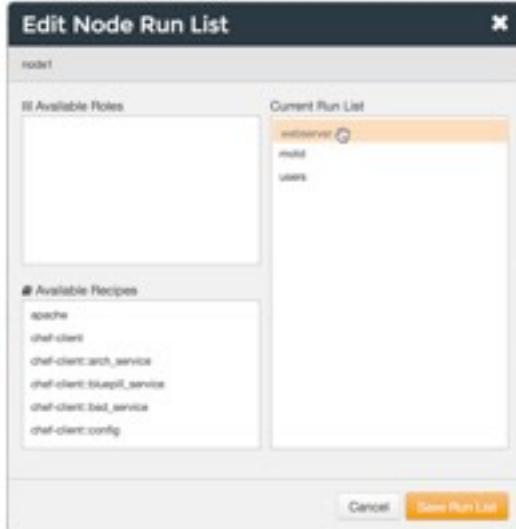
```
$ knife search role "run_list:recipe\[apache\]"
```

```
1 items found

chef_type:          role
default_attributes:
  apache:
    sites:
      admin:
        port: 8000
description:        Web Server
env_run_lists:
json_class:         Chef::Role
name:               webserver
override_attributes:
run_list:           recipe[apache]
```

## Exercise: Replace recipe[apache] with role[webserver] in run list

- Click the 'Nodes' tab then select node 'node1'
- Click 'Edit Run List' from left navigation bar
- Drag 'Apache' over from 'Current Run List' to 'Available Recipes'
- Drag 'webserver' over from 'Available Roles' to the top of 'Current Run List'
- Click 'Save Run List'



## Exercise: Re-run the Chef Client

```
chef@node1$ sudo chef-client
```

```
INFO: *** Chef 11.8.2 ***
INFO: Chef-client pid: 5634
INFO: Run List is [role[webserver], recipe[motd], recipe[users]]
INFO: Run List expands to [apache, motd, users]
```

# Exercise: Re-run the Chef Client

```
* directory[/srv/apache/admin] action create[2014-01-07T06:39:51-05:00] INFO: Processing directory[/srv/apache/admin] action create (apache::default line 53)
[2014-01-07T06:39:51-05:00] INFO: directory[/srv/apache/admin] created directory /srv/apache/admin

- create new directory /srv/apache/admin[2014-01-07T06:39:51-05:00] INFO: directory[/srv/apache/admin] mode changed to 755

- change mode from '' to '0755'
- restore selinux security context

* template[/srv/apache/admin/index.html] action create[2014-01-07T06:39:51-05:00] INFO: Processing template[/srv/apache/admin/index.html] action create (apache::default line 58)
[2014-01-07T06:39:51-05:00] INFO: template[/srv/apache/admin/index.html] created file /srv/apache/admin/index.html

- create new file /srv/apache/admin/index.html[2014-01-07T06:39:51-05:00] INFO: template[/srv/apache/admin/index.html] updated file contents /srv/apache/admin/index.html

- update content in file /srv/apache/admin/index.html from none to $@bbb@3
  --- /srv/apache/admin/index.html 2014-01-07 06:39:51.762119942 -0500
  +++ /tmp/chef-rendered-template20140107-17953-1aqbli 2014-01-07 06:39:51.764120098 -0500
@@ -1 +1 @@
<!DOCTYPE html>
<html>
  <head>
    <title>Welcome to Chef!</title>
  </head>
  <body>
    <h1>We love admin!</h1>
    <p>19.168.201.96:#8000</p>
  </body>
</html>[2014-01-07T06:39:51-05:00] INFO: template[/srv/apache/admin/index.html] mode changed to 644

- change mode from '' to '0644'
```

## Node Attributes that are hashes are merged

- The apache cookbooks attribute file contains:

```
default["apache"]["sites"]["clowns"] = { "port" => 80 }
default["apache"]["sites"]["bears"] = { "port" => 81 }
```

```
default_attributes({
  "apache" => {
    "sites" => {
      "admin" => {
        "port" => 8000
      }
    }
  }
})
```

- While our role has...

Exercise: Display the apache.sites attribute on all nodes with webserver role

```
$ knife search node "role:webserver" -a apache.sites
```

```
1 items found

node1:
apache.sites:
  admin:
    port: 8000
  bears:
    port: 81
  clowns:
    port: 80
```

## Exercise: Edit the webserver role



OPEN IN EDITOR: roles/webserver.rb

- Do not forget the **comma** after the admin site
- Change the value of the bears site to be 8081

```
default_attributes({
  "apache" => {
    "sites" => {
      "admin" => {
        "port" => 8000
      },
      "bears" => {
        "port" => 8081
      }
    }
})
```

SAVE FILE!

## Exercise: Create the role

```
$ knife role from file webserver.rb
```

Updated Role webserver!

## Exercise: Re-run the Chef Client

```
chef@node1$ sudo chef-client
```

```
...
[2014-01-07 06:57:48-05:00] INFO: template[/etc/httpd/conf.d/bears.conf] updated file contents /etc/
httpd/conf.d/bears.conf

- update content in file /etc/httpd/conf.d/bears.conf from 30610b to 92e1e4
  --- /etc/httpd/conf.d/bears.conf      2014-01-07 05:13:42.450113970 -0500
  +++ /tmp/chef-rendered-template20140107-18263-5563t5      2014-01-07 06:57:48.091397610 -0500
@@ -1,6 +1,6 @@
- Listen 81
+ Listen 8081

-<VirtualHost *:81>
+<VirtualHost *:8081>
  ServerAdmin webmaster@localhost
  DocumentRoot /srv/apache/bears
...

```

Exercise: Display the apache.sites attribute on all nodes with the webserver role

```
$ knife search node 'role:webserver' -a apache.sites
```

```
1 items found

node1:
apache.sites:
  admin:
    port: 8000
  bears:
    port: 8081
  clowns:
    port: 80
```

When you combine merge  
order and precedence  
rules, you get this:

# Merge Order and Precedence

	Attribute Files	Node / Recipe	Environment	Role
default	1	2	3	4
force_default	5	6		
normal	7	8		
override	9	10	12	11
force_override	13	14		
automatic		15		

## Best Practice: Roles get default attributes

- While it is awesome that you can use overrides, in practice there is little need
- If you always set **default** node attributes in your cookbook attribute files
- You can almost **always** set default node attributes in your role, and let merge order do the rest

## Best Practice: Have "base" roles

- In addition to obvious roles, such as "webserver", it is a common practice to group any functionality that "goes together" in a role
- The most common example here is a **base** role, where you include all the recipes that should be run on every node

## Exercise: Create the base role

 OPEN IN EDITOR: roles/base.rb

```
name "base"
description "Base Server Role"
run_list "recipe[motd]", "recipe[users]"
```

SAVE FILE!

## Exercise: Create the role

```
$ knife role from file base.rb
```

Updated Role base!

## Exercise: Add the base role to the webserver role's run list

 OPEN IN EDITOR: roles/webserver.rb

```
name "webserver"
description "Web Server"
run_list "role[base]", "recipe[apache]"
default_attributes({
  "apache" => {
```

- Put `role[base]` at the front of the `run_list`

SAVE FILE!

## Exercise: Update the role

```
$ knife role from file webserver.rb
```

```
Updated Role webserver!
```

## Exercise: Re-run the Chef Client

```
chef@node1$ sudo chef-client
```

```
INFO: *** Chef 11.8.2 ***
Chef-client pid: 5634
INFO: Run List is [role[webserver], recipe[motd],
recipe[users]]
INFO: Run List expands to [motd, users, apache]
```

## Best Practice: Be explicit about what you need or expect

- Chef will only execute a recipe the first time it appears in the run list
- So be explicit about your needs and expectations - either by nesting roles or using `include_recipe`

## Exercise: Set the run list to just `role[webserver]`

- Remove all the entries in the run list other than `role[webserver]`

## Review Questions

- What is a Role?
- What makes for a "good" role?
- How do you search for roles with a given recipe in their run list?
- How many times will Chef execute a recipe in the same run?

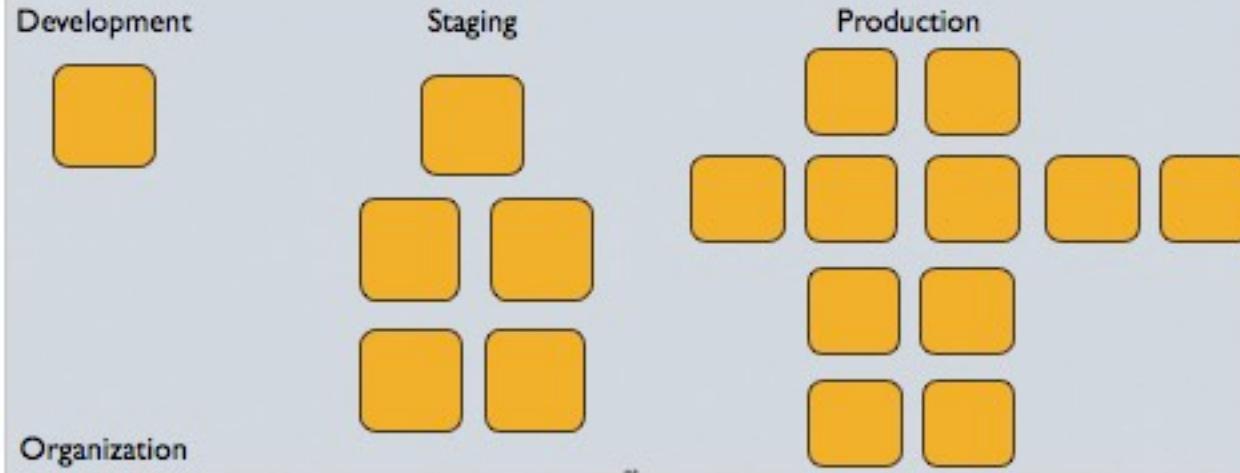
## Environments

Cookbook Version Constraints and Override Attributes

## Lesson Objectives

- After completing the lesson, you will be able to
  - Describe what an Environment is, and how it is different from an Organization
  - Set cookbook version constraints
  - Explain when to set attributes in an environment

## Environments



## Environments

- Every Organization starts with a single environment
- Environments reflect your patterns and workflow
  - Development
  - Test
  - Staging
  - Production
  - etc.

## Environments Define Policy

- Each environment may include attributes necessary for configuring the infrastructure in that environment
  - Production needs certain Yum repos
  - QA needs different Yum repos
  - The version of the Chef cookbooks to be used

## Environment Best Practice

- We cannot share cookbooks between organizations
- Best Practice: If you need to share cookbooks or roles, you likely want an Environment rather than an organization
- Environments allow for isolating resources within a single organization

Exercise: Use knife to show the available cookbook versions

```
$ knife cookbook show apache
```

```
apache    0.2.0  0.1.0
```

## Exercise: List current environments

```
$ knife environment list
```

```
_default
```

- The `_default` environment is read-only, and sets no policy at all

## Make an environments directory

```
$ mkdir environments
```

```
(No output)
```

## Exercise: Create a dev environment



OPEN IN EDITOR: environments/dev.rb

```
name "dev"  
description "For developers!"  
cookbook "apache", "= 0.2.0"
```

SAVE FILE!

- Environments have **names**
- Environments have a **description**
- Environments *can* have one or more **cookbook** constraints

## Cookbook Version Constraints

- `=` Equal to
- There are other options but equality is the recommended practice.
- Learn more at [http://docs.chef.io/chef/essentials\\_cookbook\\_versions.html](http://docs.chef.io/chef/essentials_cookbook_versions.html)

## Exercise: Create the dev environment

```
$ knife environment from file dev.rb
```

```
Updated Environment dev
```

## Exercise: Show your dev environment

```
$ knife environment show dev
```

```
chef_type:          environment
cookbook_versions:
  apache:  = 0.2.0
default_attributes:
description:        For developers!
json_class:         Chef::Environment
name:               dev
override_attributes:
```

## Exercise: Change your node's environment to "dev"

- Click the 'Nodes' tab then select node 'node1'
- Select dev from the 'Environments' drop-down list
- Click 'Save'

The screenshot shows the Chef web interface with the 'Nodes' tab selected. A single node, 'node1', is listed. The 'Environment' dropdown menu is open, showing 'dev' as the selected option. A yellow oval highlights this selection. Below the dropdown, a message states 'Node environment changed from devtest to dev'. At the bottom of the modal, there is a 'Save' button.

## Exercise: Re-run the Chef Client

```
chef@node1$ sudo chef-client
```

```
INFO: Chef Run complete in 1.587776095 seconds
INFO: Running report handlers
INFO: Report handlers complete
```

## Exercise: Create a production environment



OPEN IN EDITOR: environments/production.rb

- Make sure the apache cookbook is set to version 0.1.0
- Set an override attribute for being in scope - no matter what, you are in scope

```
name "production"
description "For Prods!"
cookbook "apache", "= 0.1.0"
override_attributes({
  "pci" => {
    "in_scope" => true
  }
})
```

SAVE FILE!

## Exercise: Create the production environment

```
$ knife environment from file production.rb
```

Updated Environment production

## Exercise: Change your node's environment to "production"

- Click the 'Nodes' tab then select node 'node1'
- Select production from the 'Environments' drop-down list
- Click 'Save'

The screenshot shows the Chef web interface. On the left, there is a sidebar with options like 'Delete', 'Manage Tags', 'Reset Key', 'Edit Run List', and 'Edit Attributes'. The main area has tabs for 'Nodes', 'Reports', 'Policy', and 'Administration'. The 'Nodes' tab is selected, showing a table with one row for 'node1'. The node details are shown in a modal window. In the 'Environment' dropdown, 'staging' is currently selected, but an orange oval highlights the 'production' option, which is described as 'Node environment changed from staging to production'. A 'Save' button is visible at the bottom of this section.

## Exercise: Re-run the Chef Client

```
chef@node1$ sudo chef-client
```

```
INFO: Loading cookbooks (apache, motd, poi, users)
Synchronizing Cookbooks:
...
  Recipe: motd::default
    * template[/etc/motd] action create[2014-01-07T08:40:00-05:00] INFO: Processing template[/etc/motd] action create
      (motd::default line 9)
    [2014-01-07T08:40:00-05:00] INFO: template[/etc/motd] backed up to /var/chef/backup/etc/motd.chef-20140107084000.670961
    [2014-01-07T08:40:00-05:00] INFO: template[/etc/motd] updated file contents /etc/motd
      - update contents in file /etc/motd from d36elf to 42eb89
      (current file is binary, diff output suppressed)

  * cookbook_file[/var/www/index.html] action create[2014-01-07T08:40:05-05:00] INFO: Processing cookbook_file[/var/www/index.html] action create (apache::default line 18)
    (up to date)
  [2014-01-07T08:40:06-05:00] INFO: Chef Run complete in 8.048087322 seconds
  [2014-01-07T08:40:06-05:00] INFO: removing cookbooks/apache/templates/default/index.html.erb from the cache; it is no longer needed by chef-client.
  [2014-01-07T08:40:06-05:00] INFO: Removing cookbooks/apache/templates/default/custom.erb from the cache; it is no longer needed by chef-client.
```

## Rollbacks and Desired State Best Practice

- Chef is not magic - it manages state for **declared** resources
- We just rolled back to an earlier version of the apache cookbook
- While the recipe applied fine, investigating the system will reveal Apache is still configured as it was in the 0.2.0 cookbook
- A better way to ensure a smooth rollback: write contra-resources to clean up, and have a new version of the cookbook.

## Review Questions

- What is an Environment?
- How is it different from an Organization?
- What kind of node attributes do you typically set from an Environment?

# Chef Supermarket

Open Source: Saving you time!



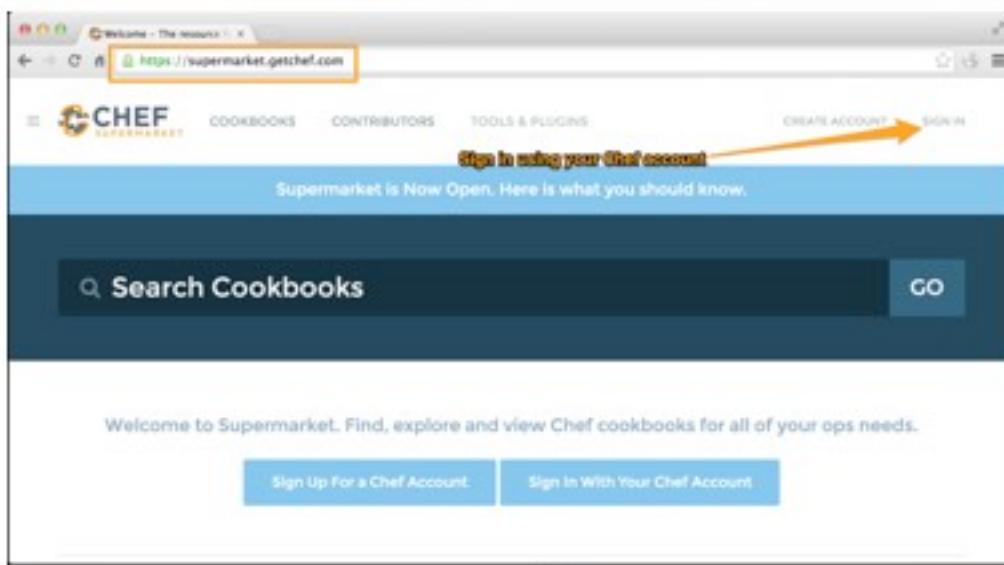
## Lesson Objectives

- After completing the lesson, you will be able to
  - Find, preview and download cookbooks from the Chef Supermarket community site
  - Use knife to work with the Chef Supermarket site API
  - Download, extract, examine and implement cookbooks from the Chef Supermarket

## The easy way...

- We've been writing some cookbooks so far...
- Hundreds already exist for a large number of use cases and purposes. Many (but only a fraction) are maintained by Chef.
- Think of it like RubyGems.org, CPAN.org, or other focused plugin-style distribution sites.

## Exercise: Log into Supermarket



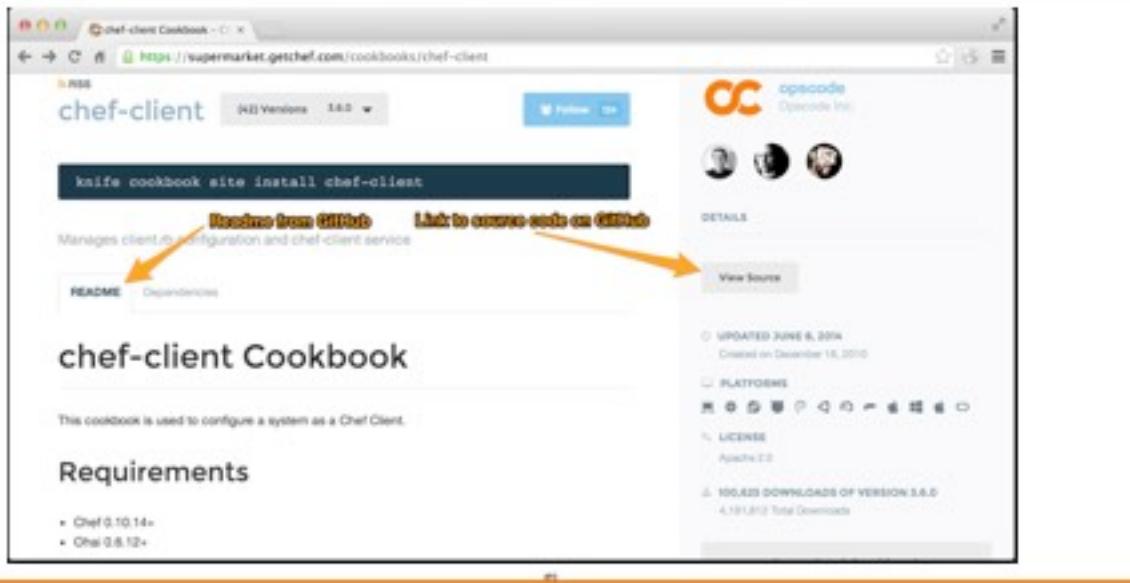
# Exercise: Search for chef-client

A screenshot of a web browser displaying the Chef Supermarket dashboard at <https://supermarket.getchef.com/dashboard>. The search bar at the top contains the query "chef-client". Below the search bar, there are sections for "Followed Cookbooks Activity" and "Your Cookbooks". A message says "Supermarket is Now Open. Here is what you should know." and a "GO" button is visible.

## Search results

A screenshot of the Chef Supermarket search results page for "chef-client" at <https://supermarket.getchef.com/cookbooks?utf8=%E2%9C%93&q=chef-client&order=&button=>. The results show 167 cookbooks. An orange arrow points from the "Click chef-client" link above the results to the "chef-client" entry in the list. The "chef-client" entry shows a version of 3.6.0 and was updated on June 6, 2014. It is described as managing client.rb configuration and the chef-client service. A snippet of the cookbook's code is shown: "cookbook 'chef-client', '=> 3.6.0'".

# Viewing a cookbook



You can download cookbooks directly from the site...

- You can download cookbooks directly from the Supermarket site, but:
  - It doesn't put them in your Chef Repository
  - It isn't fast if you know what you're looking for (click, click...)
  - It isn't necessarily fast if you **don't** know what you're looking for.
  - You're already using `knife` for managing cookbooks and other things in your Chef Repository.

## Introducing Knife Cookbook Site plugin

- Knife includes a "cookbook site" plugin with some sub-commands:
  - search
  - show
  - download
  - ... and more!

Download and use  
chef-client cookbook

## Exercise: Use knife to search the Chef Supermarket

```
$ knife cookbook site search chef-client
```

```
chef:
  cookbook: http://cookbooks.opscode.com/api/v1/cookbooks/chef
  cookbook_description: Installs and configures Chef for chef-client and chef-server
  cookbook_maintainer: opscode
  cookbook_name: chef
chef-client:
  cookbook: http://cookbooks.opscode.com/api/v1/cookbooks/chef-client
  cookbook_description: Manages client.rb configuration and chef-client service
  cookbook_maintainer: opscode
  cookbook_name: chef-client
chef-client-cron:
  cookbook: http://cookbooks.opscode.com/api/v1/cookbooks/chef-client-cron
  cookbook_description: Manages aspects of only chef-client
  cookbook_maintainer: bryansb
  cookbook_name: chef-client-cron
```

## Exercise: Use knife to show the chef-client cookbook on the Chef Supermarket

```
$ knife cookbook site show chef-client
```

```
average_rating: 4.85714
category: Utilities
created_at: 2010-12-16T23:00:45Z
description: Manages client.rb configuration and chef-client service
external_url: github.com/opscode-cookbooks/chef-client
latest_version: http://cookbooks.opscode.com/api/v1/cookbooks/chef-client/
versions/3_2_0
maintainer: opscode
name: chef-client
updated_at: 2013-12-19T19:02:44Z
versions:
  http://cookbooks.opscode.com/api/v1/cookbooks/chef-client/versions/3_2_0
  http://cookbooks.opscode.com/api/v1/cookbooks/chef-client/versions/3_1_2
  http://cookbooks.opscode.com/api/v1/cookbooks/chef-client/versions/3_1_0
```

## Exercise: Download the chef-client cookbook

```
$ knife cookbook site download chef-client
```

```
Downloading chef-client from the cookbooks
site at version 3.2.0 to /Users/
johnfitzpatrick/cheftraining/fundamentals2.0/
chef-repo/chef-client-3.2.0.tar.gz
```

```
Cookbook saved: /Users/YOU/chef-repo/chef-
client-3.2.0.tar.gz
```

## Exercise: Extract chef-client cookbook tarball

```
$ tar -zxvf chef-client*.tar.gz -C cookbooks/
```

```
x chef-client/
x chef-client/attributes/
x chef-client/CHANGELOG.md
x chef-client/CONTRIBUTING
x chef-client/LICENSE
x chef-client/metadata.json
x chef-client/metadata.rb
x chef-client/README.md
x chef-client/recipes/
x chef-client/templates/
x chef-client/templates/arch/
x chef-client/templates/default/
x chef-client/templates/windows/
x chef-client/templates/default/debian/
x chef-client/templates/default/redhat/
x chef-client/templates/default/solaris/
x chef-client/templates/arch/conf.d/
x chef-client/templates/arch/rc.d/
x chef-client/recipes/config.rb
x chef-client/recipes/cron.rb
x chef-client/recipes/default.rb
x chef-client/recipes/delete_validation.rb
```

## What we just did...

- Cookbooks are distributed as a versioned .tar.gz archive.
- The latest version is downloaded by default (you can specify the version).
- Extract the cookbook into the "cookbooks" directory with tar.
- Next, let's examine the contents.

### Best practice: well written cookbooks have a README!

- Documentation for cookbooks doesn't need to be extensive, but a README should describe some important aspects of a cookbook:
  - Expectations (cookbooks, platform, data)
  - Recipes and their purpose
  - LWRPs, Libraries, etc.
  - Usage notes
- Read the README first!

## Best Practice: This runs as root!

- So, you just downloaded source code from the internet.
- As root.
- To load in the magic machine that:
  - **Makes your computers run code**
- Read the *entire* cookbook first!

## Examining the chef-client cookbook

- We're going to use two recipes on the node from the chef-client cookbook.
  - `delete_validation`
  - `service` (via default)

## Exercise: View the chef-client::delete\_validation recipe

 OPEN IN EDITOR: cookbooks/chef-client/recipes/delete\_validation.rb

```
unless chef server?
  file Chef::Config[:validation_key] do
    action :delete
    backup false
    only_if { ::File.exists?(Chef::Config[:client_key]) }
  end
end
```

SAVE FILE!

## Exercise: Add chef-client::delete\_validation to your base role

 OPEN IN EDITOR: roles/base.rb

```
name "base"
description "Base Server Role"
run_list "recipe[chef-client::delete_validation]", "recipe[motd]", "recipe[users]"
```

SAVE FILE!

- Add the **delete\_validation** recipe

## Best Practice: Delete the validation certificate when it isn't required

- Once Chef enters the actual run, synchronizing cookbooks, it has registered its own API client with the validation certificate
- That certificate is no longer required. We do this first because in case the run fails for another reason, we know at least the validation certificate is gone

## Exercise: View the chef-client::default recipe



OPEN IN EDITOR: cookbooks/chef-client/recipes/default.rb

```
#  
# Unless required by applicable law or agreed to in writing, software  
# distributed under the License is distributed on an "AS IS" BASIS,  
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.  
# See the License for the specific language governing permissions and  
# limitations under the License.  
#  
include_recipe "chef-client::service"
```

**SAVE FILE!**

## Best Practice: Sane defaults do "pretty much" what you expect

- The main point of the "chef-client" cookbook is managing the "chef-client" program. It is designed that it can run as a daemonized service.
- The least surprising thing for most users is that the default recipe starts the service.
- You can manage the service in a number of ways, see the cookbook's README.md.

## Exercise: View the chef-client::service recipe



[OPEN IN EDITOR: cookbooks/chef-client/recipes/service.rb](#)

- The recipe supports a number of **service** providers and styles.
- It works on a lot of **platforms**.
- Everything is controllable through **attributes**.

```
supported_init_styles = [
  'arch',
  'bluepill',
  'bsd',
  'daemontools',
  'init',
  'launchd',
  'runit',
  'smf',
  'upstart',
  'winsw'
]

init_style = node['chef_client']['init_style']

# Services moved to recipes
if supported_init_styles.include? init_style
  include_recipe "chef-client::#{init_style}_service"
else
  log "Could not determine service init style, manual intervention required
  to start up the chef-client service."
end
```

## Best Practice: Well-written cookbooks change behavior based on attributes

- Ideally, you don't have to modify the contents of a cookbook to use it for your specific use case.
- Look at the attributes directory for things you can override through roles to affect behavior of the cookbook.
- Of course, well written cookbooks have sane defaults, and a README to describe all this.

## Exercise: Upload the chef-client cookbook

```
$ knife cookbook upload chef-client
```

```
Uploading chef-client (1.2.0)
ERROR: Cookbook 'chef-client' is not currently being loaded and has no files on the server.
ERROR: which version '>= 1.2.0',
```



## Exercise: Download the cron cookbook

```
$ knife cookbook site download cron
```

```
Downloading cron from the cookbooks site at version  
1.2.8 to /Users/YOU/chef-repo/cron-1.2.8.tar.gz  
Cookbook saved: /Users/YOU/chef-repo/  
cron-1.2.8.tar.gz
```

## Exercise: Extract cron cookbook tarball

```
$ tar -zxvf cron*.tar.gz -C cookbooks/
```

```
x cron/  
x cron/CHANGELOG.md  
x cron/README.md  
x cron/metadata.json  
x cron/metadata.rb  
x cron/providers  
x cron/providers/d.rb  
x cron/recipes  
x cron/recipes/default.rb  
x cron/recipes/test.rb  
x cron/resources  
x cron/resources/d.rb  
x cron/templates  
x cron/templates/default  
x cron/templates/default/cron.d.erb
```

## Exercise: Upload the cron cookbook

```
$ knife cookbook upload cron
```

```
Uploading cron [1.2.8]  
Uploaded 1 cookbook.
```

## Exercise: Upload the chef-client cookbook

```
$ knife cookbook upload chef-client
```

**FAIL!**  
Uploading chef-client [1.1.0]  
ERROR: Cookbook chef-client depends on  
cookbook 'logrotate' version '>= 1.2.0',  
ERROR: which is not currently being  
uploaded and cannot be found on the  
server.

## Exercise: Download the logrotate cookbook

```
$ knife cookbook site download logrotate
```

```
Downloading logrotate from the cookbooks site at
version 1.4.0 to /Users/johnfitzpatrick/
cheftraining/chef-repo/logrotate-1.4.0.tar.gz
Cookbook saved: /Users/YOU/chef-repo/
logrotate-1.4.0.tar.gz
```

## Exercise: Extract logrotate cookbook tarball

```
$ tar -zxvf logrotate*.tar.gz -C cookbooks/
```

```
x logrotate/
x logrotate/CHANGELOG.ad
x logrotate/README.md
x logrotate/attributes
x logrotate/attributes/default.rb
x logrotate/definitions
x logrotate/definitions/logrotate_app.rb
x logrotate/libraries
x logrotate/libraries/logrotate_config.rb
x logrotate/metadata.json
x logrotate/metadata.rb
x logrotate/recipes
x logrotate/recipes/default.rb
x logrotate/recipes/global.rb
x logrotate/templates
x logrotate/templates/default
x logrotate/templates/default/logrotate-global.erb
x logrotate/templates/default/logrotate.erb
```

## Exercise: Upload the logrotate cookbook

```
$ knife cookbook upload logrotate
```

```
Uploading logrotate [1.4.0]  
Uploaded 1 cookbook.
```

## Exercise: Upload the chef-client cookbook

```
$ knife cookbook upload chef-client
```

```
Uploading chef-client [3.0.0]  
Uploaded 1 cookbook!
```

A large, bold, green text "WIN!" is displayed prominently in the center of the terminal window, indicating success.

## Exercise: Add chef-client recipe to base role

 OPEN IN EDITOR: roles/base.rb

```
name "base"
description "Base Server Role"
run_list "recipe[chef-client::delete_validation]", "recipe[chef-client]",
"recipe[motd]", "recipe[users]"
```

SAVE FILE!

## Exercise: Upload the base role

```
$ knife role from file base.rb
```

Updated Role base!

# Exercise: Re-run the Chef Client

```
***  
Recipe: chef-client::delete_validation  
* file[/etc/chef/validation.pem] action delete[2014-01-07T09:05:43-05:00] INFO: Processing file[/etc/chef/validation.pem] action  
delete (chef-client::delete_validation line 25)  
[2014-01-07T09:05:43-05:00] INFO: file[/etc/chef/validation.pem] deleted file at /etc/chef/validation.pem  
- delete file /etc/chef/validation.pem  
***  
* service[chef-client] action enable[2014-01-07T09:05:44-05:00] INFO: Processing service[chef-client] action enable (chef-  
client::init_service line 32)  
[2014-01-07T09:05:47-05:00] INFO: service[chef-client] enabled  
- enable service service[chef-client]  
* service[chef-client] action start[2014-01-07T09:05:47-05:00] INFO: Processing service[chef-client] action start (chef-  
client::init_service line 32)  
[2014-01-07T09:05:48-05:00] INFO: service[chef-client] started  
- start service service[chef-client]  
***  
[2014-01-07T09:05:55-05:00] INFO: template[/etc/init.d/chef-client] sending restart action to service[chef-client] (delayed)  
Recipe: chef-client::init_service  
* service[chef-client] action restart[2014-01-07T09:05:55-05:00] INFO: Processing service[chef-client] action restart (chef-  
client::init_service line 32)  
[2014-01-07T09:06:01-05:00] INFO: service[chef-client] restarted  
- restart service service[chef-client]  
[2014-01-07T09:06:01-05:00] INFO: Chef Run complete in 29.341653545 seconds
```

# Exercise: Verify chef-client is running

```
chef@node1$ ps awux | grep chef-client
```

```
root      8933  0.3  2.2 130400 37816 ?          S1   03:19  
0:01 /opt/chef/embedded/bin/ruby /usr/bin/chef-client -d -c /  
etc/chef/client.rb -L /var/log/chef/client.log -P /var/run/  
chef/client.pid -i 1800 -s 300
```

## Convergent infrastructure

- Our node is now running chef-client as a daemon, and it will converge itself over time on a (by default) 30 minute interval.
- The amount of resources converged may vary with longer intervals, depending on configuration drift on the system.
- Because Chef resources are idempotent, it will only configure what it needs to each run.

Download and use ntp cookbook

# NTP Cookbook

- Network time protocol - keeps system clocks in sync
- Chef Server authentication is time sensitive!

## Exercise: Download the ntp cookbook

```
$ knife cookbook site download ntp
```

```
Downloading ntp from the cookbooks  
site at version 1.5.4 to /Users/YOU/  
chef-repo/ntp-1.5.4.tar.gz
```

```
Cookbook saved: /Users/YOU/chef-repo/  
ntp-1.5.4.tar.gz
```

## Exercise: Extract the ntp cookbook

```
$ tar -zxvf ntp*.tar.gz -C cookbooks/
```

```
x ntp/
x ntp/CHANGELOG.md
x ntp/README.md
x ntp/attributes
x ntp/attributes/default.rb
x ntp/files
x ntp/files/default
x ntp/files/default/ntp.ini
x ntp/files/default/ntp.leapseconds
x ntp/files/default/tests
x ntp/files/default/tests/minitest
x ntp/files/default/tests/minitest/default_test.rb
x ntp/files/default/tests/minitest/support
x ntp/files/default/tests/minitest/support/helpers.rb
x ntp/files/default/tests/minitest/undo_test.rb
x ntp/files/default/usr.sbin.ntpd.apparmor
x ntp/metadata.json
x ntp/metadata.rb
x ntp/recipes
```

## Examining the ntp cookbook

- The cookbook is quite flexible, but for this exercise we're just interested in the most basic use, an NTP client.
  - default recipe

## Exercise: View the `ntp::default` recipe

```
node['ntp']['packages'].each do |ntppkg|
  package ntppkg
end

template node['ntp']['conffile'] do
  source  'ntp.conf.erb'
  owner   node['ntp']['conf_owner']
  group   node['ntp']['conf_group']
  mode    '0644'
  notifies :restart,
  "service[#{node['ntp']['service']}]"
end

service node['ntp']['service'] do
  supports :status => true, :restart => true
  action  [:enable, :start]
end
```

- All **packages** are installed
- The **service** is enabled and started
- The **template** notifies the service

## Exercise: Upload the `ntp` cookbook

```
$ knife cookbook upload ntp
```

```
Uploading ntp
Uploaded 1 cookbook.
```

```
[1.5.4]
```

## Exercise: Add the ntp recipe to the base role



**OPEN IN EDITOR:** roles/base.rb

```
name "base"
description "Base Server Role"
run_list "recipe[chef-client::delete_validation]", "recipe[chef-client]",
"recipe[ntp]", "recipe[motd]", "recipe[users]"
```

**SAVE FILE!**

## Exercise: Upload the base role

```
$ knife role from file base.rb
```

Updated Role base!

# Exercise: Re-run the Chef Client

```
Recipe: ntp::default
 * package[ntp] action install[2014-01-07T09:27:01-05:00] INFO: Processing package(ntp) action install
 (ntp::default line 25)
 (up to date)
 * package[ntpdate] action install[2014-01-07T09:27:08-05:00] INFO: Processing package(ntpdate) action
 install (ntp::default line 25)
 (up to date)
 ***
 * template[/etc/ntp.conf] action create[2014-01-07T09:27:09-05:00] INFO: Processing template[/etc/
 ntp.conf] action create (ntp::default line 71)
 [2014-01-07T09:27:09-05:00] INFO: template[/etc/ntp.conf] backed up to /var/chef/backup/etc/
 ntp.conf.chef-20140107092709.254951
 [2014-01-07T09:27:09-05:00] INFO: template[/etc/ntp.conf] updated file contents /etc/ntp.conf

 - update content in file /etc/ntp.conf from ba8f03 to c381bb
   --- /etc/ntp.conf 2013-10-27 09:07:05.541644862 -0400
   +++ /tmp/chef-rendered-template20140107-20557-ttvgaw 2014-01-07 09:27:09.248975286 -0500
 @@ -1,58 +1,30 @@
 ...
 [2014-01-07T09:27:10-05:00] INFO: template[/etc/ntp.conf] sending restart action to service[stpd] (delayed)
 Recipe: ntp::default
 * service[ntpd] action restart[2014-01-07T09:27:10-05:00] INFO: Processing service[stpd] action restart
 (ntp::default line 100)
 [2014-01-07T09:27:11-05:00] INFO: service[ntpd] restarted
 - restart service service[ntpd]
```

## Review Questions

- What is the Chef Supermarket site URL?
- What are two ways to download cookbooks from the Chef Supermarket?
- What is the first thing you should read when downloading a cookbook?
- Who vets the cookbooks on the Chef Supermarket?

## Further Resources

92.1.8

50



## Working with Chef Software, Inc.

51

## Chef Status

- Status for Chef related systems can be found by browsing to
  - <http://status.chef.io>, or
  - [https://twitter.com/opscode\\_status](https://twitter.com/opscode_status)

## How do I interact with Chef support?

- There are two ways to raise a ticket with Chef Support
  - Via the Web UI at <http://www.chef.io/support/>
  - By sending an email to [support@chef.io](mailto:support@chef.io)
- The Web UI allows you to submit any severity level ticket, however emailed tickets default to 'severity 3'
- The Web UI allows you to view previously submitted tickets for an org

## To re-open a closed ticket

- To re-open a closed ticket either
  - Click on the closed ticket in the web interface, and reply in the box provided, or
  - Reply to the email thread
- This will restart the ticket and move it back into the active queue

## Dealing with Support

- When an org signs up for a support contract, the email domain of the primary user is used as a key, and anyone with an email in that domain can submit tickets
- E.g. `user1@domain.com` starts a support contract for an org, then `user2@domain.com` and `user3@domain.com` can also submit tickets on behalf of "domain.com" for that org
- Any user email address be added explicitly to the Support Tool for your org by emailing [sales@chef.io](mailto:sales@chef.io) or submitting a ticket

## New releases

- Release Notes can be found here
  - <http://docs.chef.io/>

## Further Resources

## Chef Fundamentals - Q&A Forum

- Chef Fundamentals Google Group Q&A Forum
- <http://bit.ly/ChefFundamentalsForum>
- Join the group and post questions

## Further Resources: Cookbooks and Plugins

- Useful cookbooks
  - DNS: djbdns, pdns, dnsimple, dynect, route53
  - Monitoring: nagios, munin, zenoss, zabbix
  - Package repos: yum, apt, freebsd
  - Security: ossec, snort, cis\_benchmark
  - Logging: rsyslog, syslog-ng, logstash, logwatch
- Application cookbooks:
  - application, database
  - python, java, php, ruby
- Plugins
  - Cloud: knife-ec2, knife-rackspace, knife-openstack, knife-hp
  - Windows: knife-windows
- More listed on [docs.chef.io](http://docs.chef.io)

## Further Resources

- <http://chef.io/>
- <http://supermarket.chef.io/>
- <http://docs.chef.io/>
- <http://learn.chef.io>
- <http://lists.opscode.com>
- <http://youtube.com/user/getchef>
- irc.freenode.net #chef, #chef-hacking, #learnchef
- Twitter @chef #getchef

## Food Fight Show

- <http://foodfightshow.org>
- The Podcast Where DevOps Chef Do Battle
- Regular updates about new Cookbooks, Knife-plugins, and more
- Best Practices for working with Chef



# CHEFCONF2015

- March 31 - April 2 in Santa Clara Convention Center
- Register now
- <https://www.chef.io/chefconf/>
- Now accepting presentation proposals

## Appendix: Just Enough Ruby for Chef

A quick & dirty crash course

# Lesson Objectives

- After completing the lesson, you will be able to
  - Explain what irb is, and how to use it.
  - Describe the function of:
    - Variable Assignment
    - Basic Arithmetic
    - Strings
    - Truthiness
    - Operators
    - Hashes
    - Regular Expressions
    - Conditionals
    - Method Declaration
    - Classes
    - Objects

## irb - the interactive ruby shell

- irb is the interactive ruby shell
- It is a REPL for Ruby
  - Read-Eval-Print-Loop
  - LISP, Python, Erlang, Clojure, etc.
- An interactive programming environment
- Super-handy for trying things out

## Exercise: Start irb on your ‘target’ node

```
chef@node1$ /opt/chef/embedded/bin/irb
```

```
irb(main):001:0>
```

## Variable assignment

```
chef@node1$ /opt/chef/embedded/bin/irb
```

```
irb(main):001:0> x = "hello"  
=> "hello"
```

```
irb(main):002:0> puts x  
hello  
=> nil
```

## Arithmetic

```
irb(main):003:0> 1 + 2  
=> 3
```

## Arithmetic

```
irb(main):004:0> 18 - 5  
=> 13
```

## Arithmetic

```
irb(main):005:0> 2 * 7  
=> 14
```

## Arithmetic

```
irb(main):006:0> 5 / 2  
=> 2
```

## Arithmetic

```
irb(main):007:0> 5 / 2.0  
=> 2.5
```

## Arithmetic

```
irb(main):008:0> 5.class  
=> Fixnum  
  
irb(main):009:0> 5.0.class  
=> Float
```

## Arithmetic

```
irb(main):010:0> 1 + (2 * 3)
=> 7
```

## Strings

```
irb(main):011:0> 'jungle'
=> "jungle"
```

## Strings

```
irb(main):012:0> 'it\'s alive'  
=> "it's alive"
```

## Strings

```
irb(main):013:0> "animal"  
=> "animal"
```

## Strings

```
irb(main):014:0> "\\"so easy\\"
=> "\\so easy\\"

irb(main):015:0> puts "\\"so easy\\"
"so easy"
=> nil
```

## Strings

```
irb(main):016:0> x = "pretty"
=> "pretty"

irb(main):017:0> "#{x} nice"
=> "pretty nice"

irb(main):018:0> '#{x} nice'
=> "\\#{x} nice"
```

# Truthiness

```
irb(main):019:0> true  
=> true

irb(main):020:0> false  
=> false

irb(main):021:0> nil  
=> nil

irb(main):022:0> !!nil  
=> false

irb(main):023:0> !!0  
=> true

irb(main):024:0> !!x  
=> true
```

# Operators

```
irb(main):025:0> 1 == 1  
=> true

irb(main):026:0> 1 == true  
=> false

irb(main):027:0> 1 != true  
=> true

irb(main):028:0> !!1 == true  
=> true
```

# Operators

```
irb(main):029:0> 2 < 1  
=> false  
  
irb(main):030:0> 2 > 1  
=> true  
  
irb(main):031:0> 4 >= 3  
=> true  
  
irb(main):032:0> 4 >= 4  
=> true  
  
irb(main):033:0> 4 <= 5  
=> true  
  
irb(main):034:0> 4 <= 3  
=> false
```

# Operators

```
irb(main):035:0> 5 <=> 5  
=> 0  
  
irb(main):036:0> 5 <=> 6  
=> -1  
  
irb(main):037:0> 5 <=> 4  
=> 1
```

# Arrays

```
irb(main):038:0> x = ["a", "b", "c"]
=> ["a", "b", "c"]
```

```
irb(main):039:0> x[0]
=> "a"
```

```
irb(main):040:0> x.first
=> "a"
```

```
irb(main):041:0> x.last
=> "c"
```

# Arrays

```
irb(main):042:0> x + ["d"]
=> ["a", "b", "c", "d"]
```

```
irb(main):043:0> x
=> ["a", "b", "c"]
```

```
irb(main):044:0> x = x + ["d"]
=> ["a", "b", "c", "d"]
```

```
irb(main):045:0> x
=> ["a", "b", "c", "d"]
```

## Arrays

```
irb(main):046:0> x << "e"
=> ["a", "b", "c", "d", "e"]
```

```
irb(main):047:0> x
=> ["a", "b", "c", "d", "e"]
```

## Arrays

```
irb(main):048:0> x.map { |i| "the letter #{i}" }
=> ["the letter a", "the letter b", "the letter
c", "the letter d", "the letter e"]
```

```
irb(main):049:0> x
=> ["a", "b", "c", "d", "e"]
```

## Arrays

```
irb(main):050:0> x.map! { |i| "the letter #{i}" }
=> ["the letter a", "the letter b", "the letter
c", "the letter d", "the letter e"]

irb(main):051:0> x
=> ["the letter a", "the letter b", "the letter
c", "the letter d", "the letter e"]
```

## Hashes

```
irb(main):052:0> h = {
irb(main):053:1* "first_name" => "Gary",
irb(main):054:1* "last_name" => "Gygax"
irb(main):055:1> }
=> {"first_name"=>"Gary",
"last_name"=>"Gygax"}
```

## Hashes

```
irb(main):056:0> h.keys  
=> ["first_name", "last_name"]  
  
irb(main):057:0> h["first_name"]  
=> "Gary"  
  
irb(main):058:0> h["age"] = 33  
=> 33  
  
irb(main):059:0> h.keys  
=> ["first_name", "last_name", "age"]
```

## Hashes

```
irb(main):060:0> h.values  
=> ["Gary", "Gygax", 33]
```

## Hashes

```
irb(main):061:0> h.each { |k, v| puts "#{k}: #{v}" }
first_name: Gary
last_name: Gygax
age: 33
=> {"first_name"=>"Gary", "last_name"=>"Gygax",
"age"=>33}
```

## Regular Expressions

```
irb(main):062:0> x = "I want to believe"
=> "I want to believe"

irb(main):063:0> x =~ /I/
=> 0

irb(main):064:0> x =~ /lie/
=> 12

irb(main):065:0> x =~ /smile/
=> nil

irb(main):066:0> x !~ /smile/
=> true
```

# Regular Expressions

```
irb(main):067:0> x.sub(/t/, "T")
=> "I wanT to believe"
```

```
irb(main):068:0> puts x
I want to believe
=> nil
```

```
irb(main):069:0> x.gsub!(/t/, "T")
=> "I wanT To believe"
```

```
irb(main):070:0> puts x
I want To believe
=> nil
```

# Conditionals

```
irb(main):071:0> x = "happy"
=> "happy"
```

```
irb(main):072:0> if x == "happy"
irb(main):073:1>   puts "Sure am!"
irb(main):074:1> elsif x == "sad"
irb(main):075:1>   puts "Boo!"
irb(main):076:1> else
irb(main):077:1*>   puts "Therapy?"
irb(main):078:1> end
Sure am!
=> nil
```

# Conditionals

```
irb(main):079:0> case x
irb(main):080:1> when "happy"
irb(main):081:1>   puts "Sure Am!"
irb(main):082:1>   1
irb(main):083:1> when "sad"
irb(main):085:1>   puts "Boo!"
irb(main):086:1>   2
irb(main):087:1> else
irb(main):088:1>   puts "Therapy?"
irb(main):089:1>   3
irb(main):090:1> end
Sure Am!
=> 1
```

# Method Definition

```
irb(main):091:0> def metal(str)
irb(main):092:1>   puts "!!#{str} is metal!!"
irb(main):093:1> end
=> nil

irb(main):094:0> metal("ozzy")
!!ozzy is metal!!
=> nil
```

## Classes

```
irb(main):095:0> class Person
irb(main):096:1>   attr_accessor :name, :is_metal
irb(main):097:1>
irb(main):098:1>   def metal
irb(main):099:2>     if @is_metal
irb(main):100:3>       puts "!!#{@name} is metal!!"
irb(main):101:3>     end
irb(main):102:2>   end
irb(main):103:1> end
=> nil
```

## Classes

```
irb(main):104:0> p = Person.new
=> #<Person:0x891ab4c>

irb(main):105:0> p.name = "Adam Jacob"
=> "Adam Jacob"

irb(main):106:0> p.is_metal = true
=> true

irb(main):107:0> p.metal
!!Adam Jacob is metal!!
=> nil

irb(main):108:0> p.is_metal = false
=> false

irb(main):109:0> p.metal
=> nil
```

# Review Questions

- What is irb?
- What is true in ruby? What is false?
- How do you press for the truth?
- What does >= do?
- How do you define a method?
- What is a class?
- What is an object?
- The book you want: "Programming Ruby 1.9" <http://pragprog.com/book/ruby3/programming-ruby-1-9>